

ECAI-2025

October 25-30, 2025

Bologna, Italy



UNIVERSITÀ
DI TORINO

Experimenting with Federated Learning: implementing state-of-the-art methods with **fluke**

Bruno Casella, Samuele Fonio, Mirko Polato

Department of Computer Science, University of Turin, Italy

name.surname@unito.it

28th European Conference on Artificial Intelligence - ECAI2025, Bologna



Outline



Introduction to Federated Learning:
challenges, applications, algorithms



Introduction to fluke: motivation, goals,
design & functionalities



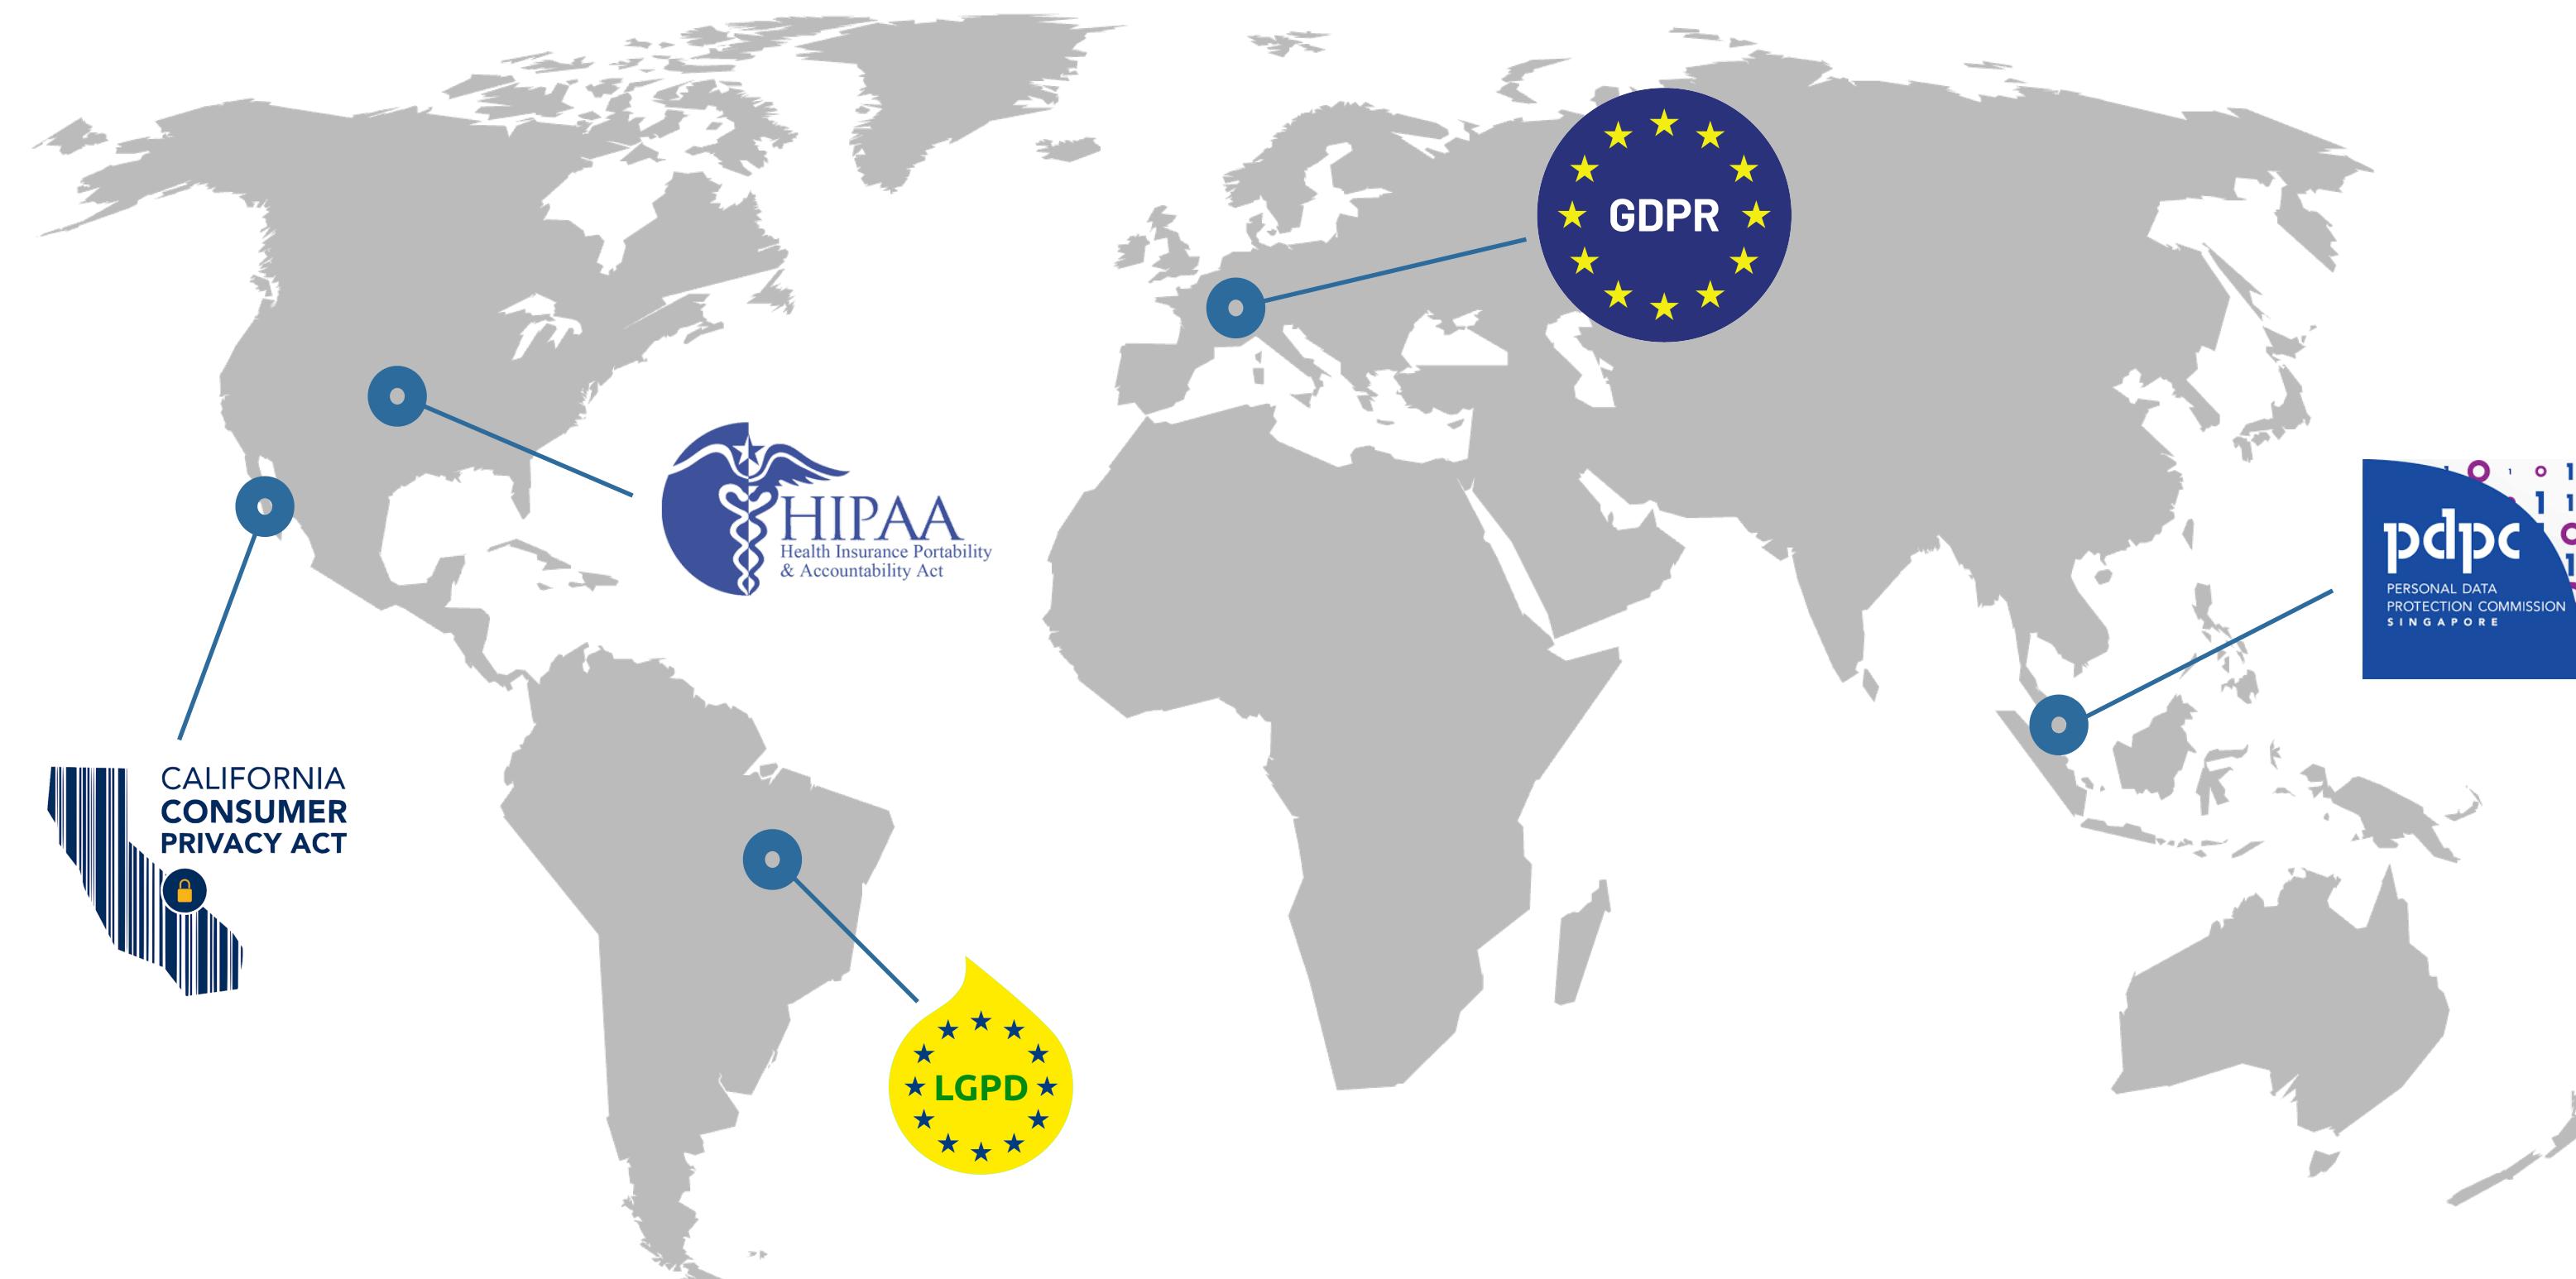
Hands-on-session: setup, run & use cases

Introduction to **Federated Learning**

Federated Learning

Motivation: Data Sovereignty and Compliance

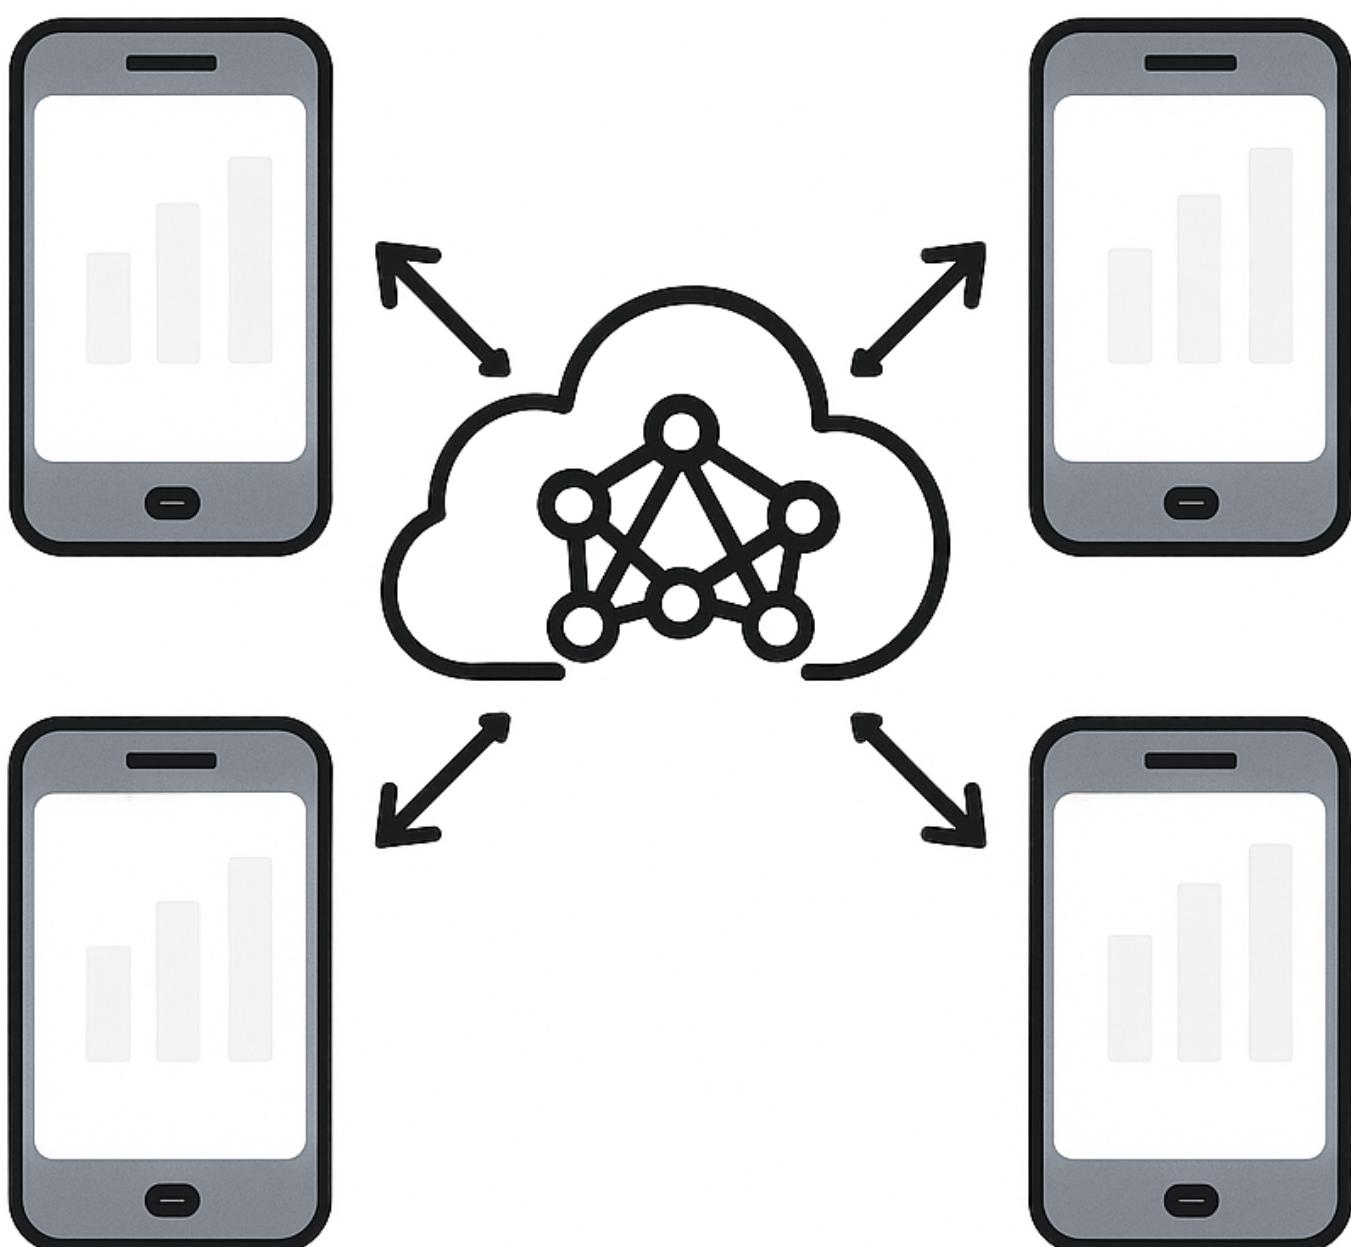
Recent **privacy regulations** (e.g., GDPR) restrict how data can move across borders or be stored.



Federated Learning

Motivation: Security, privacy & leveraging distributed data

- In standard ML, the need of centralizing the data **may expose sensitive data** (e.g., e.g., personal messages, medical records, financial data).
- Many real-world datasets are **naturally distributed** (e.g., smartphones, IoT sensors, hospitals).



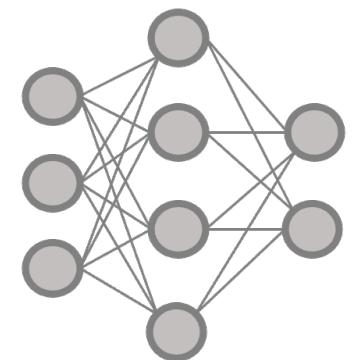
Federated Learning

Privacy-preserving distributed machine learning approach

Federated Learning (FL) is a machine learning setting where **multiple entities (clients) collaborate** in solving a machine learning problem, under the coordination of a central server. **Each client's raw data is stored locally and not exchanged or transferred**; instead, focused updates intended for immediate aggregation are used to achieve the learning objective.

Federated Learning

Parties involved and general idea



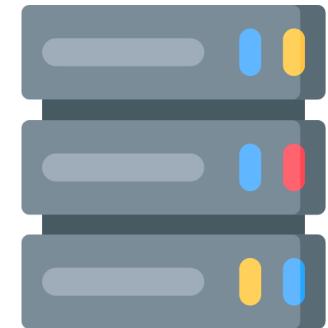
The model is what is exchanged

The model is what is transferred between the parties

Iteratively updated

The learning procedure is divided into “rounds” where a new version of the model is computed

Machine Learning model



Initialisation

The Server initialise the learning process

Orchestration

The Server orchestrates the learning process and acts as a centralisation point

Aggregation

The Server updates the model aggregating the clients' models

Server (aggregator)



Privacy

The clients' data is only used by the owner and never transferred to anyone

Local training

The local data is used to perform part of the learning locally

Model sharing

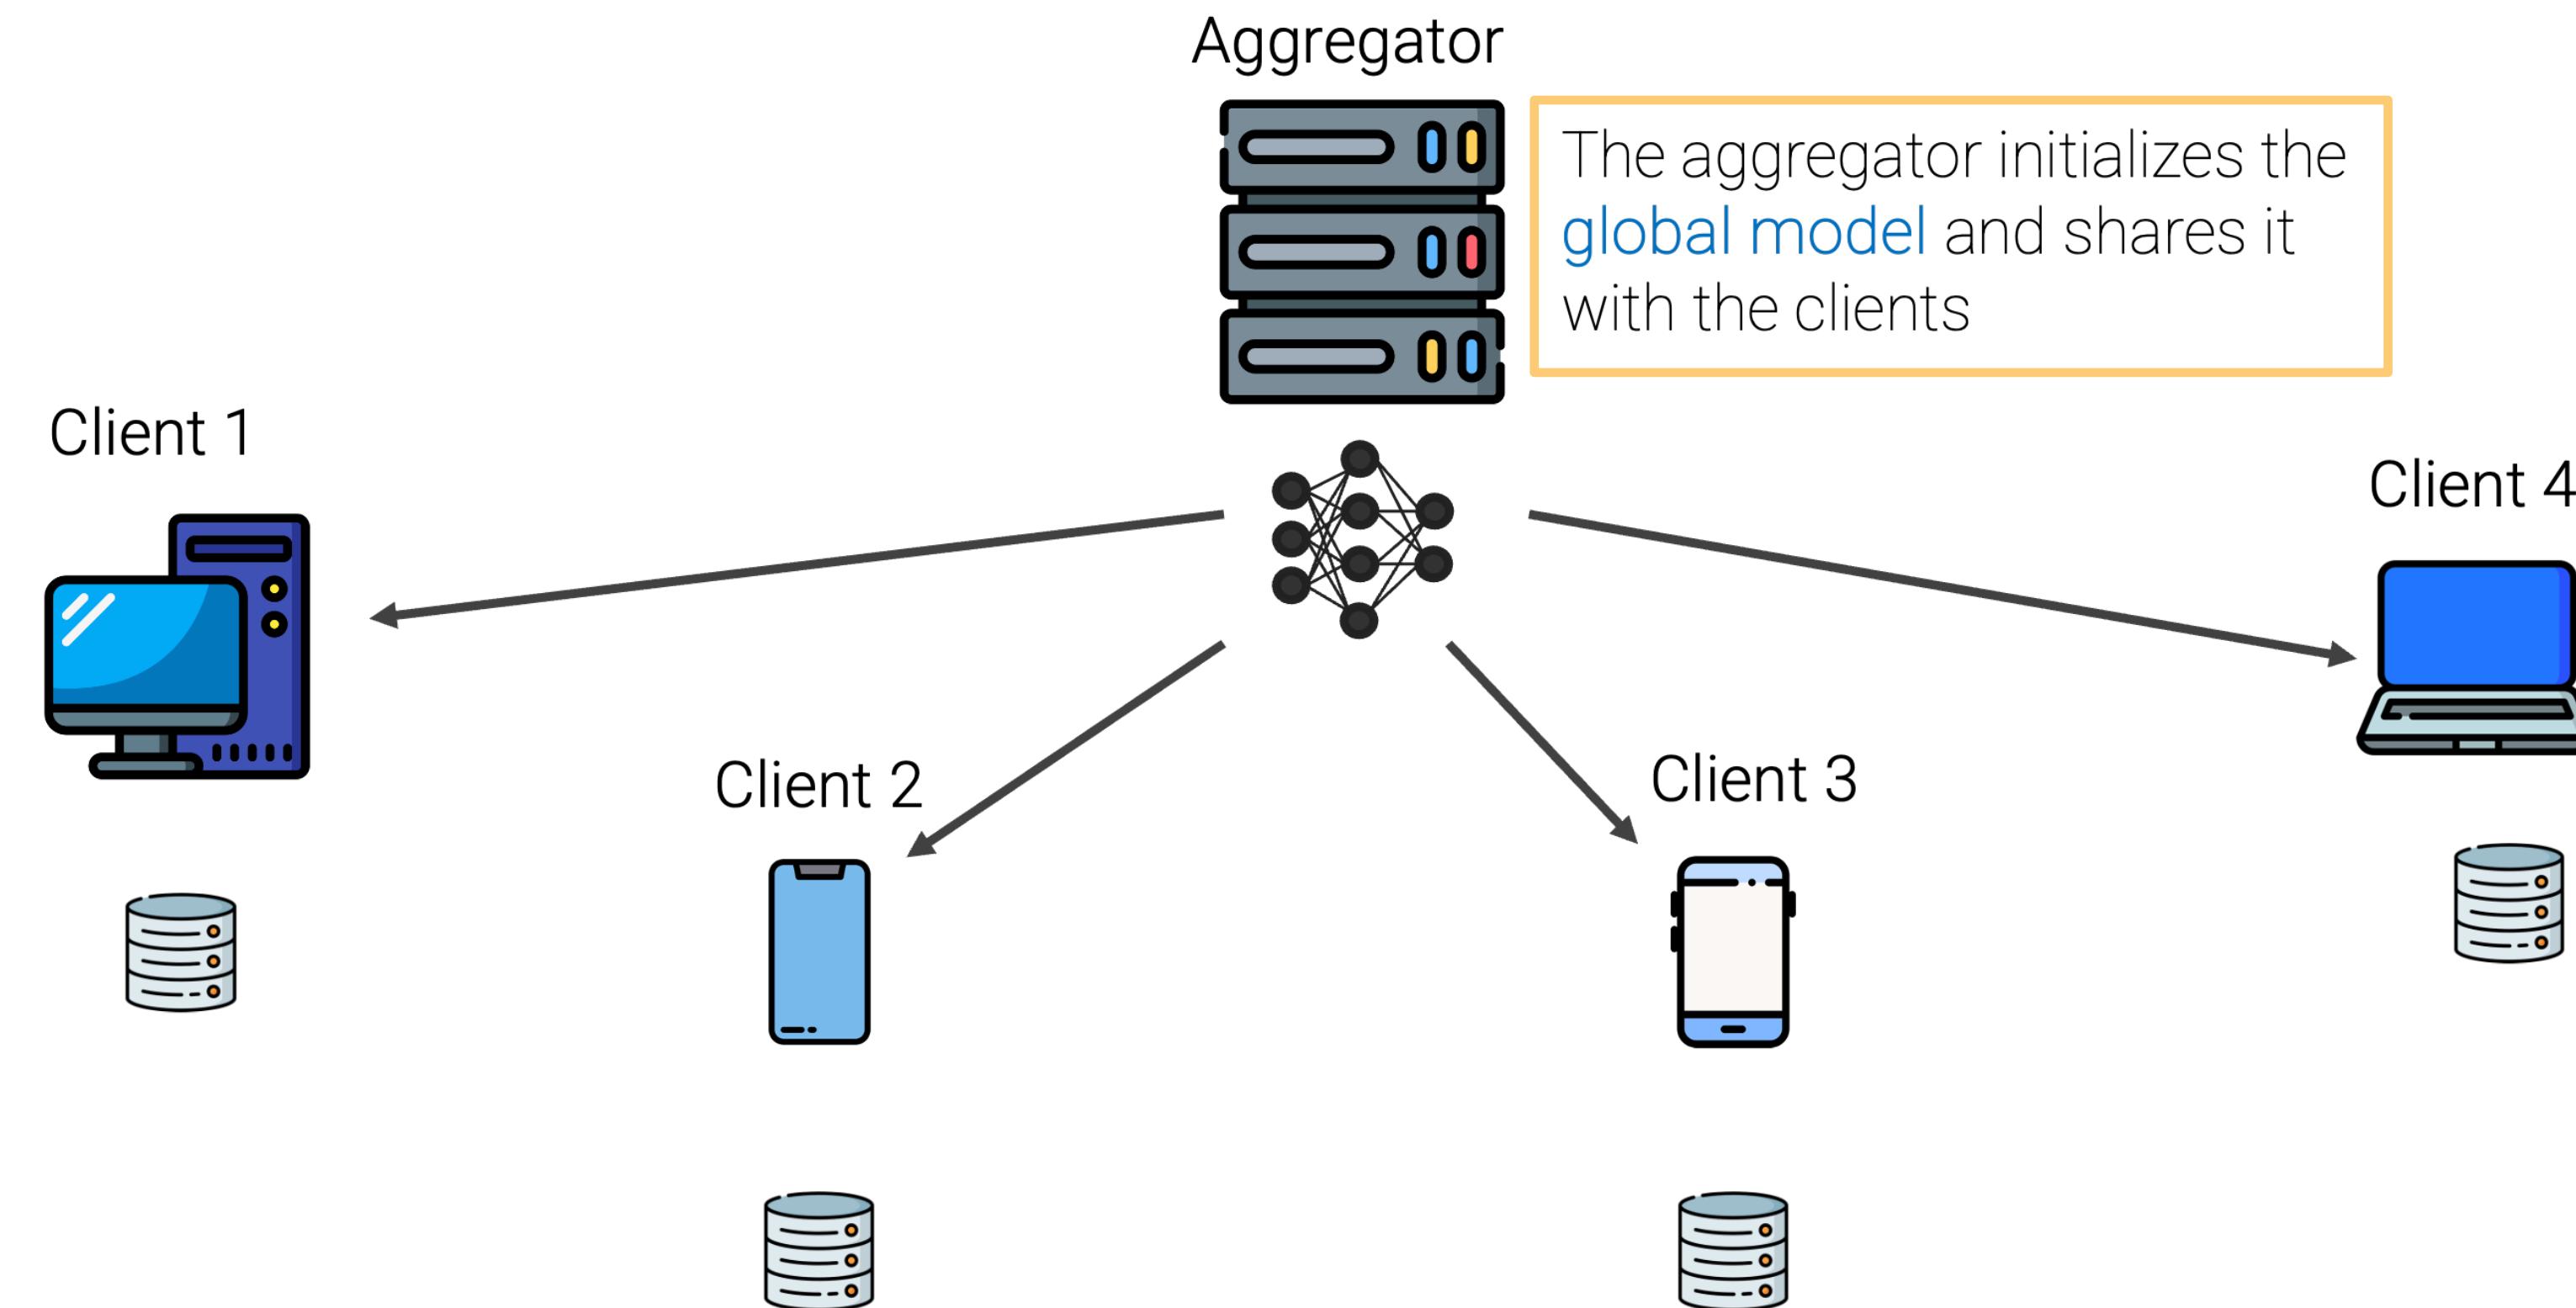
After the local training, the client send its updated model to the Server

Clients (Collaborators)

Federated Learning

Training procedure

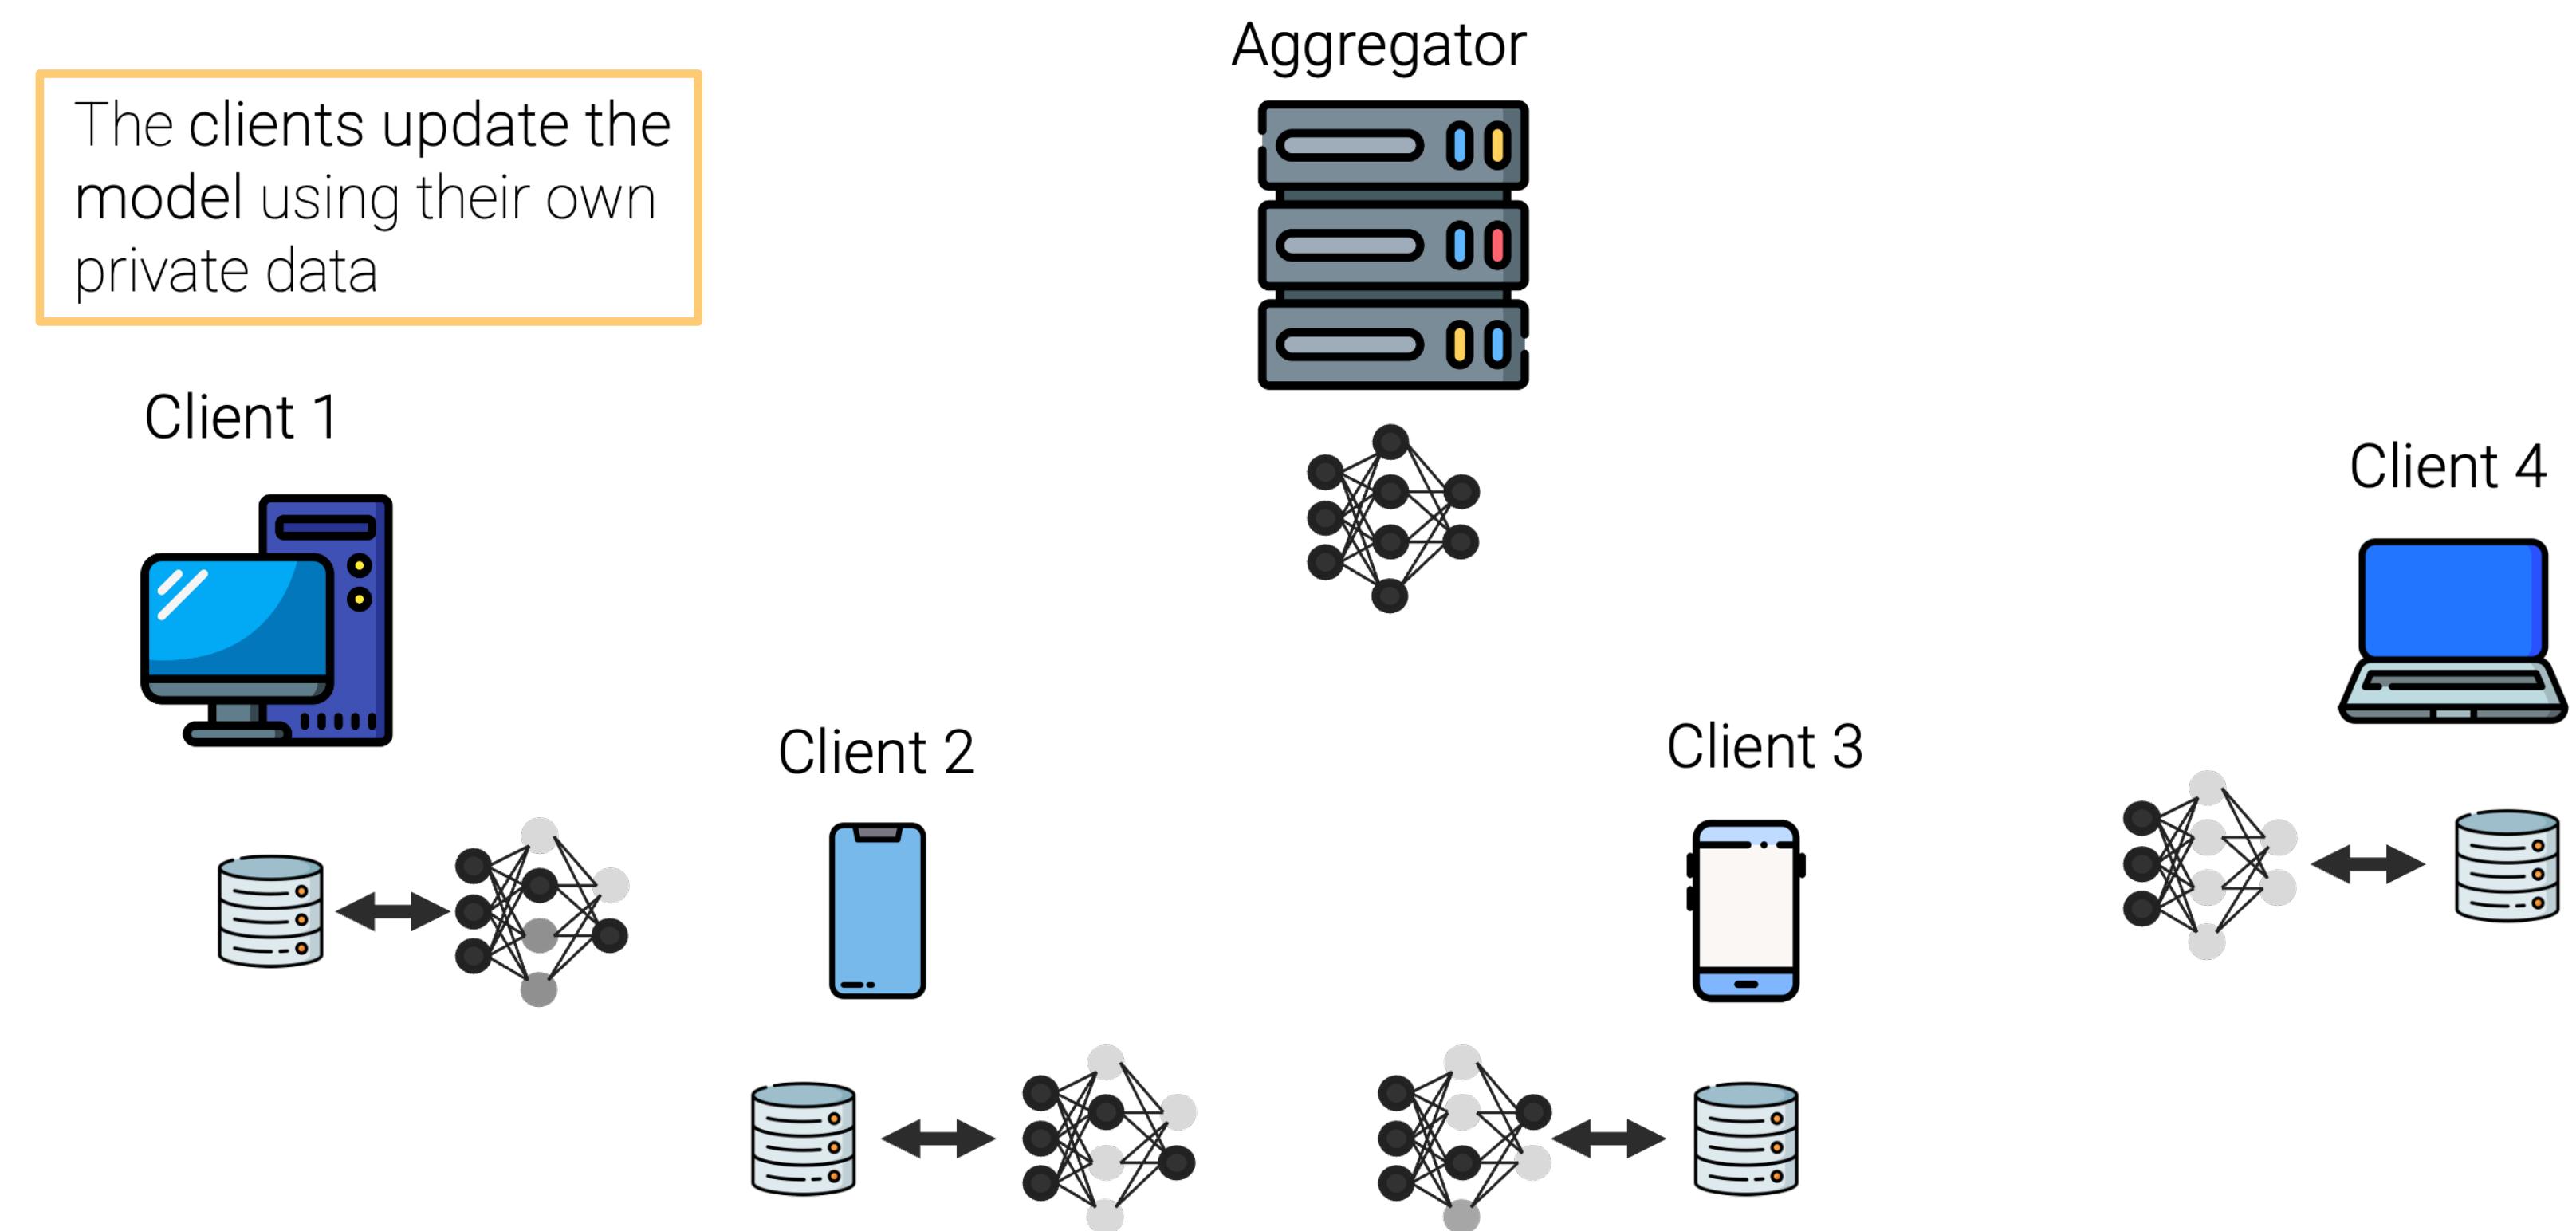
ROUND 1



Federated Learning

Training procedure

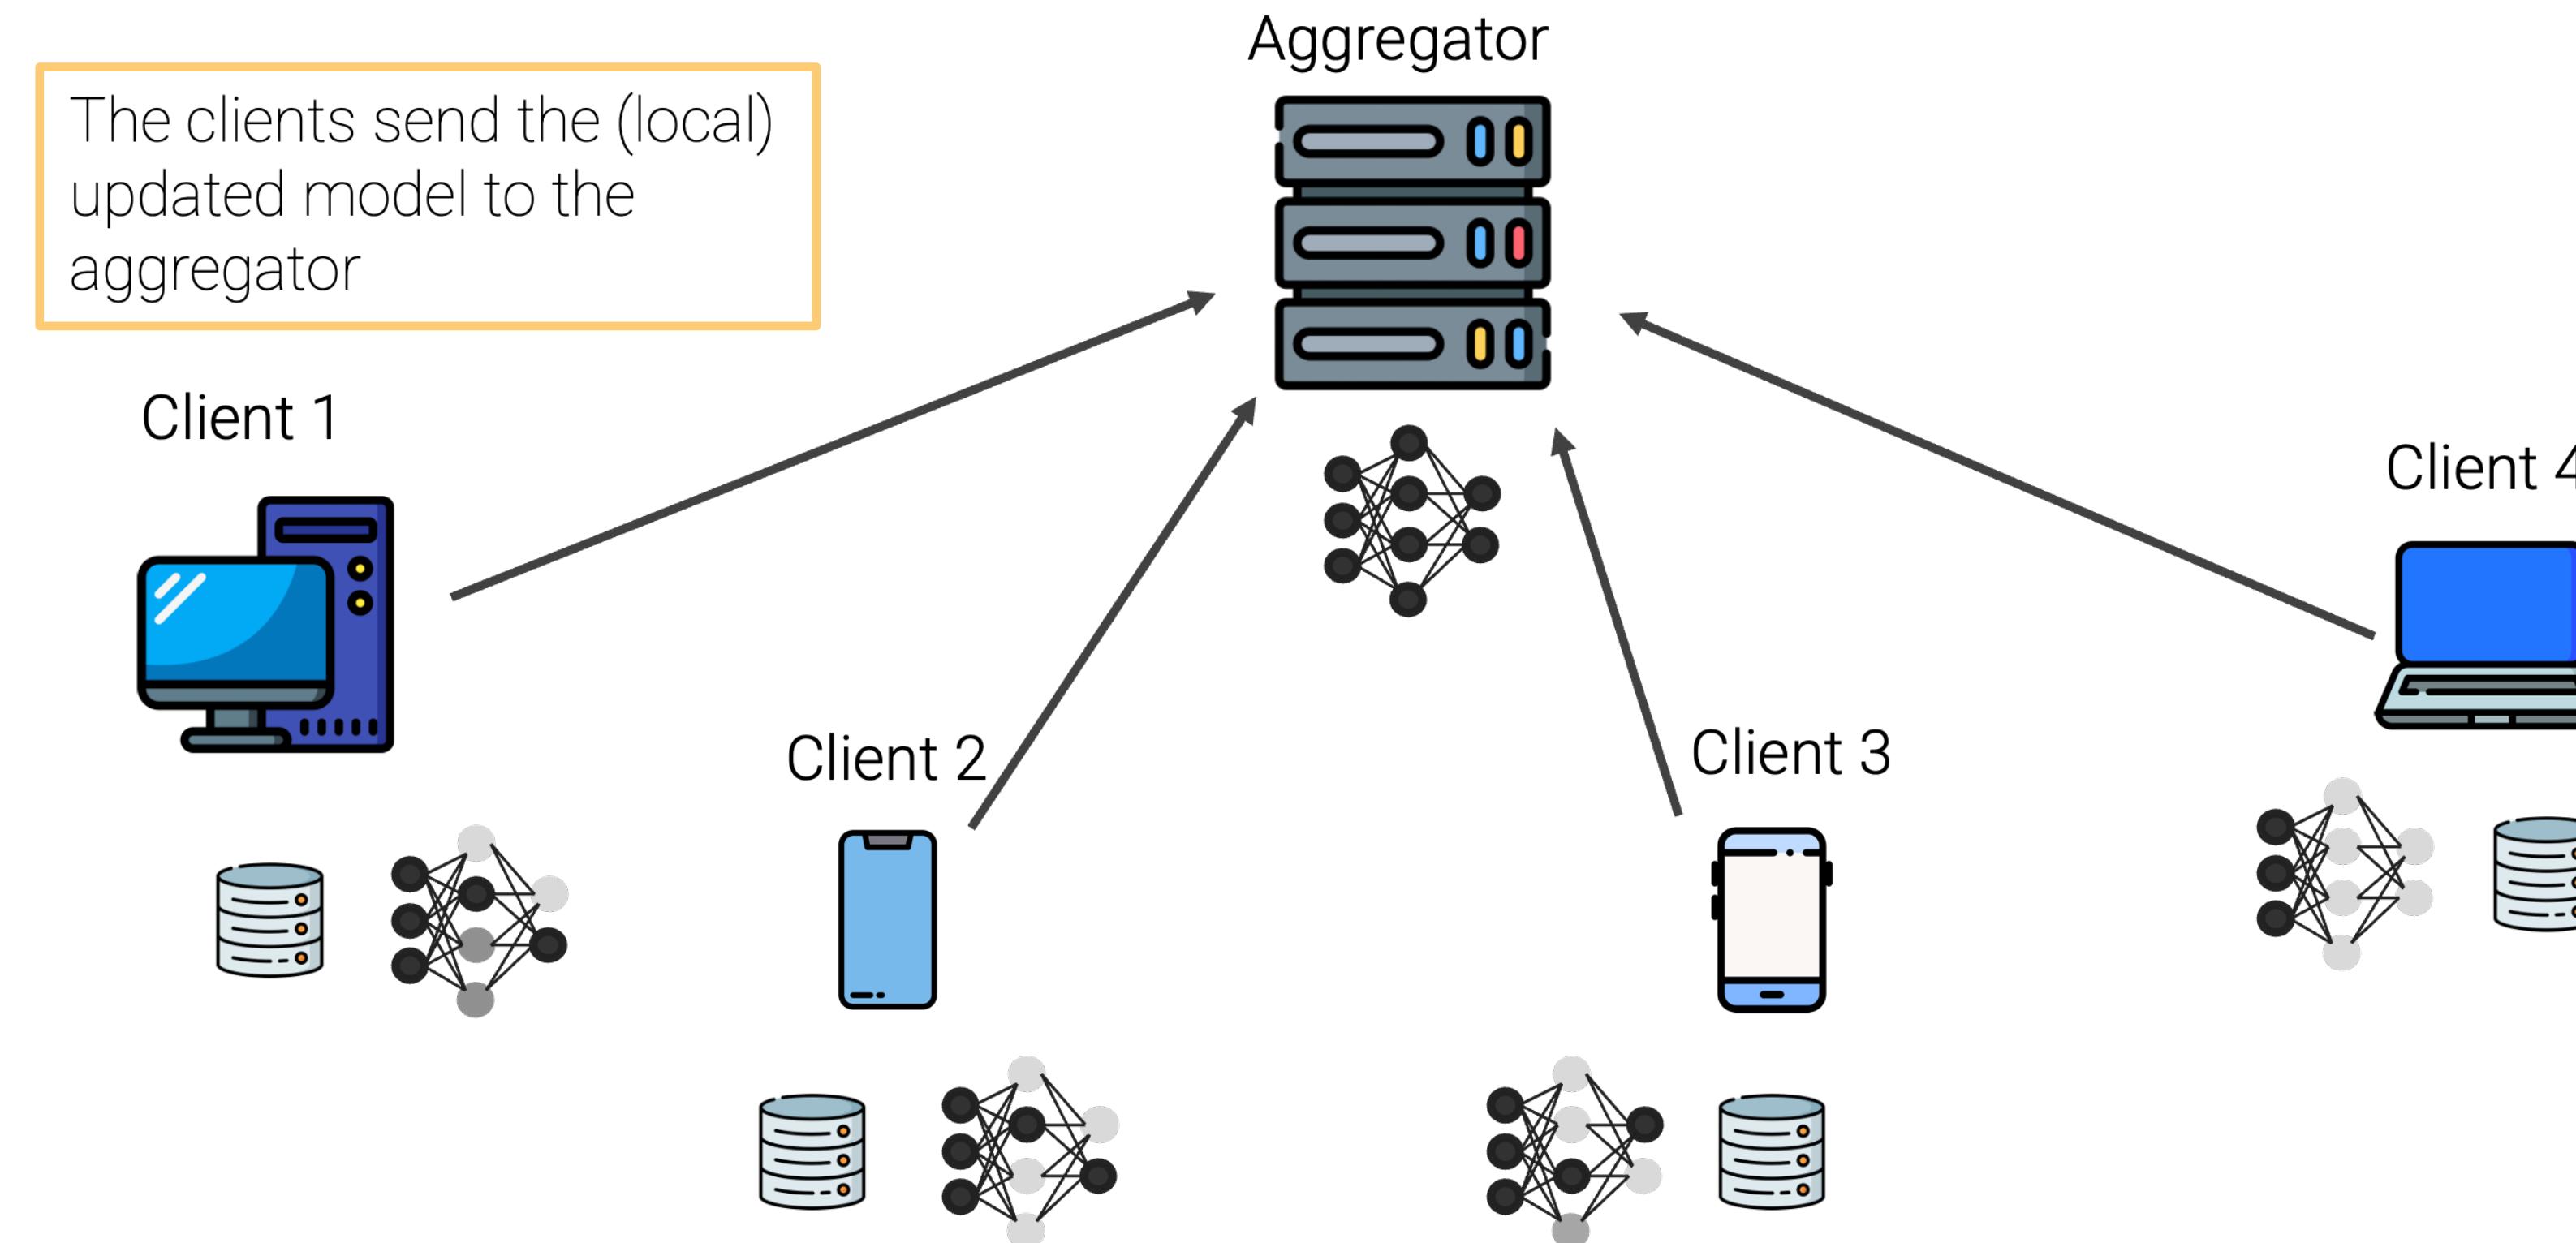
ROUND 1



Federated Learning

Training procedure

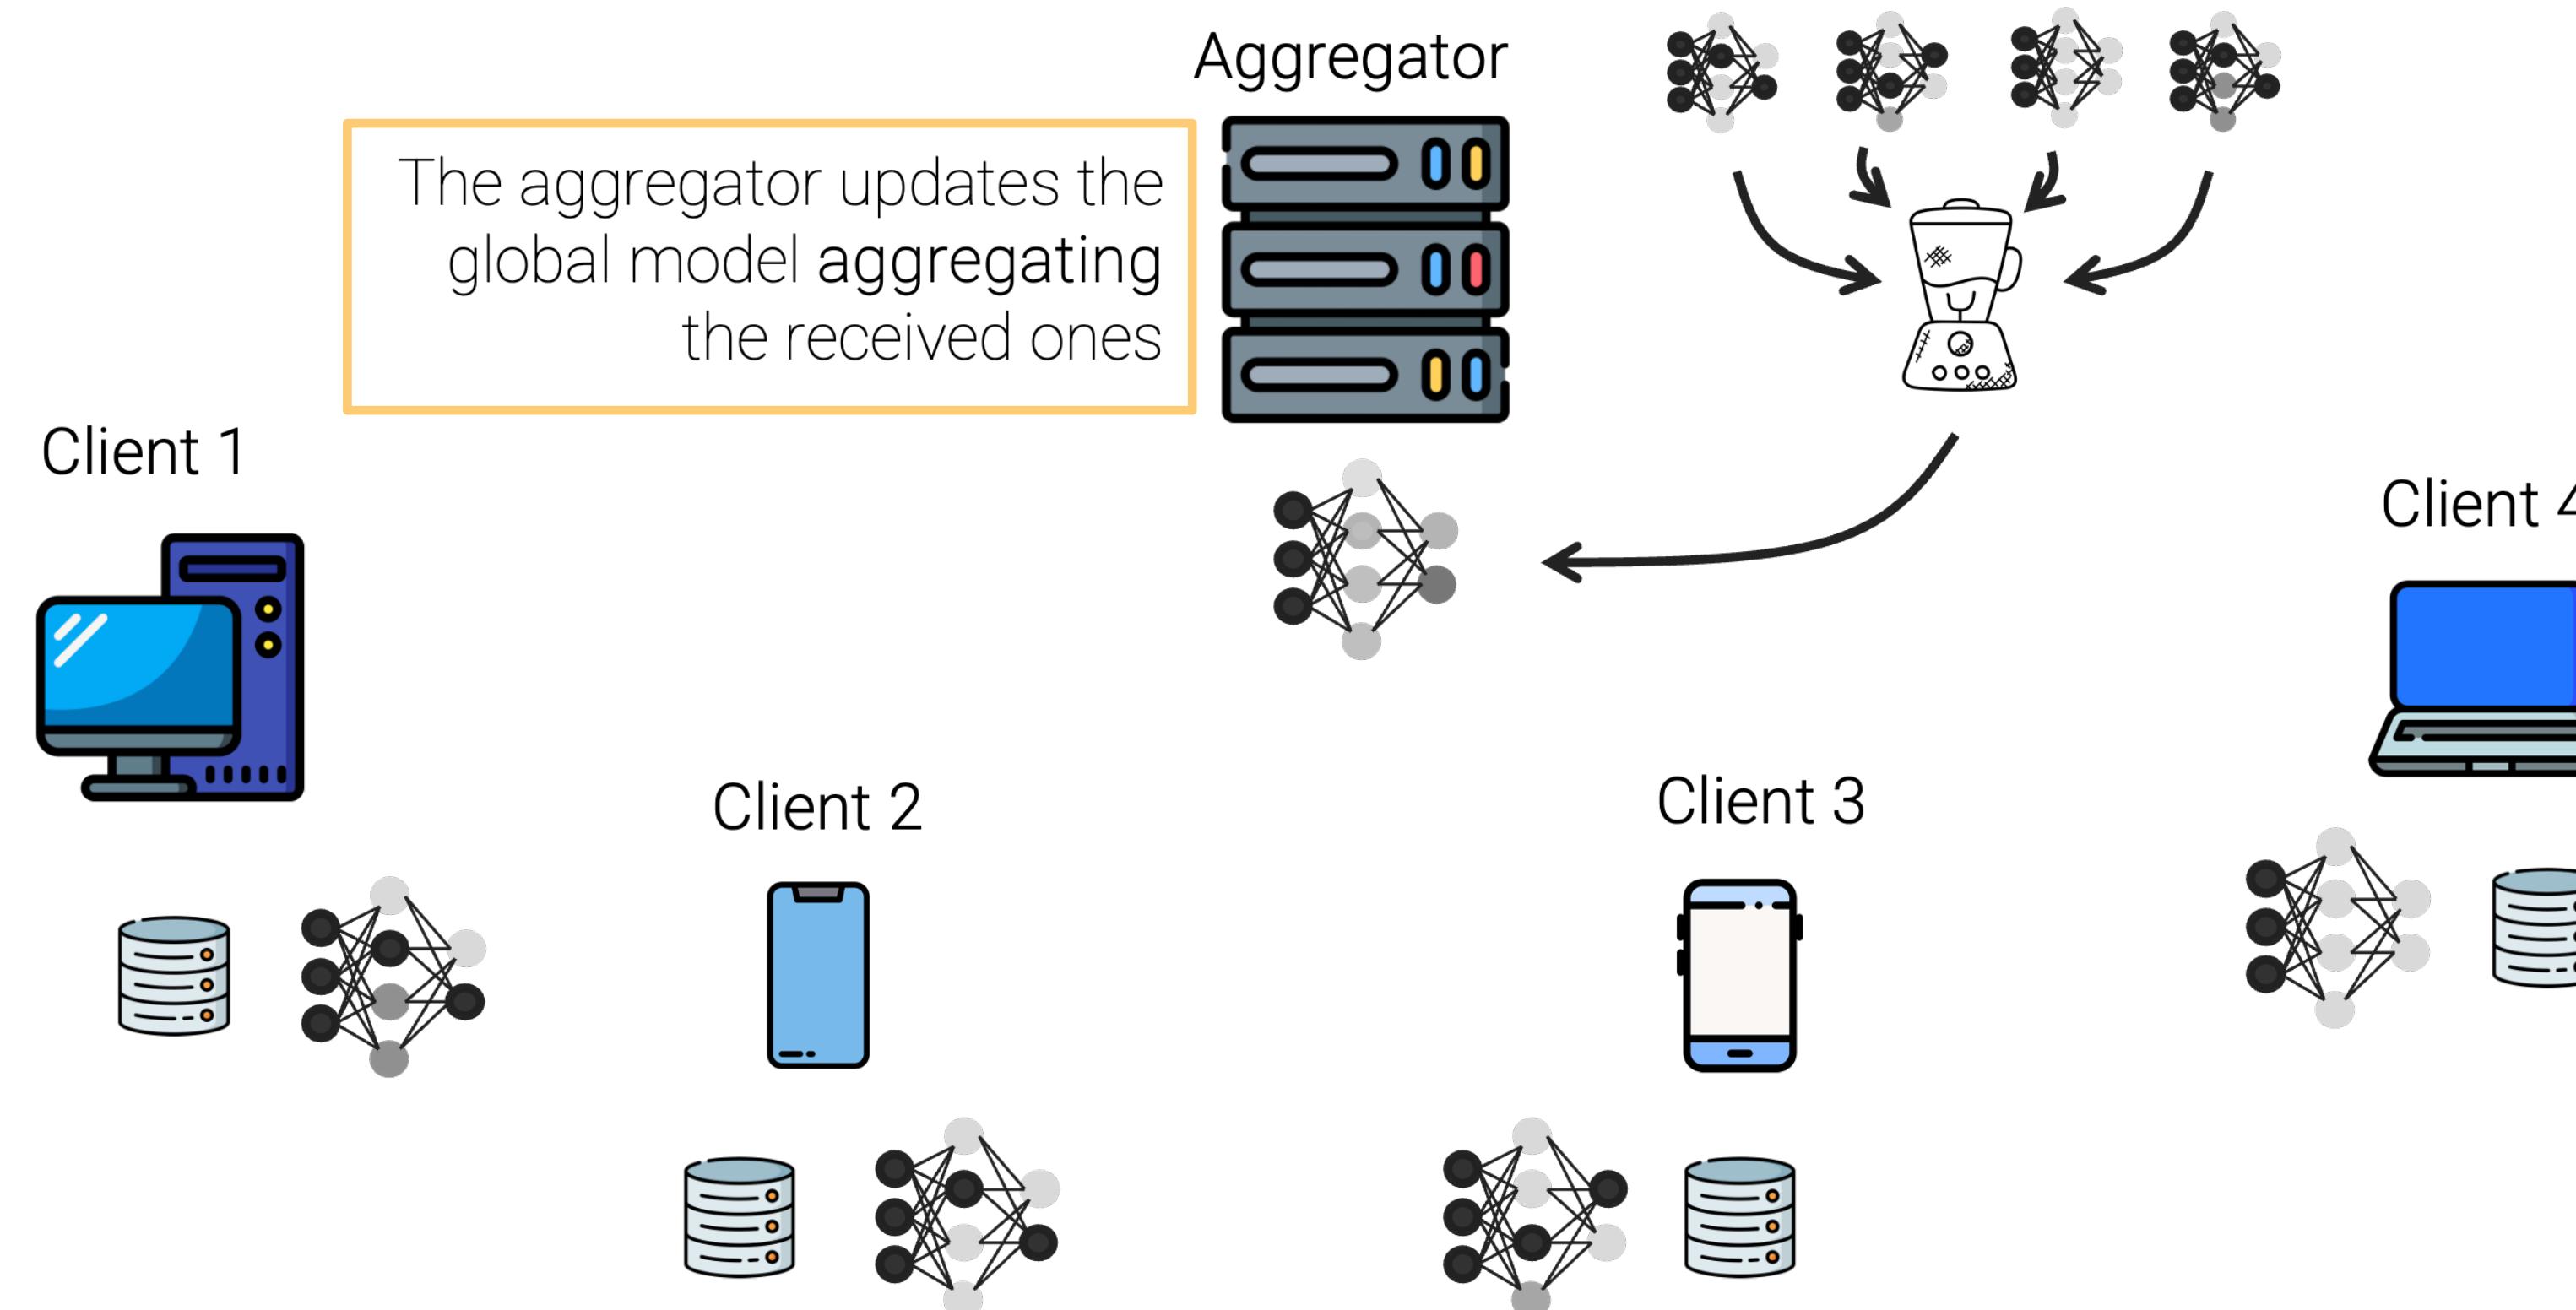
ROUND 1



Federated Learning

Training procedure

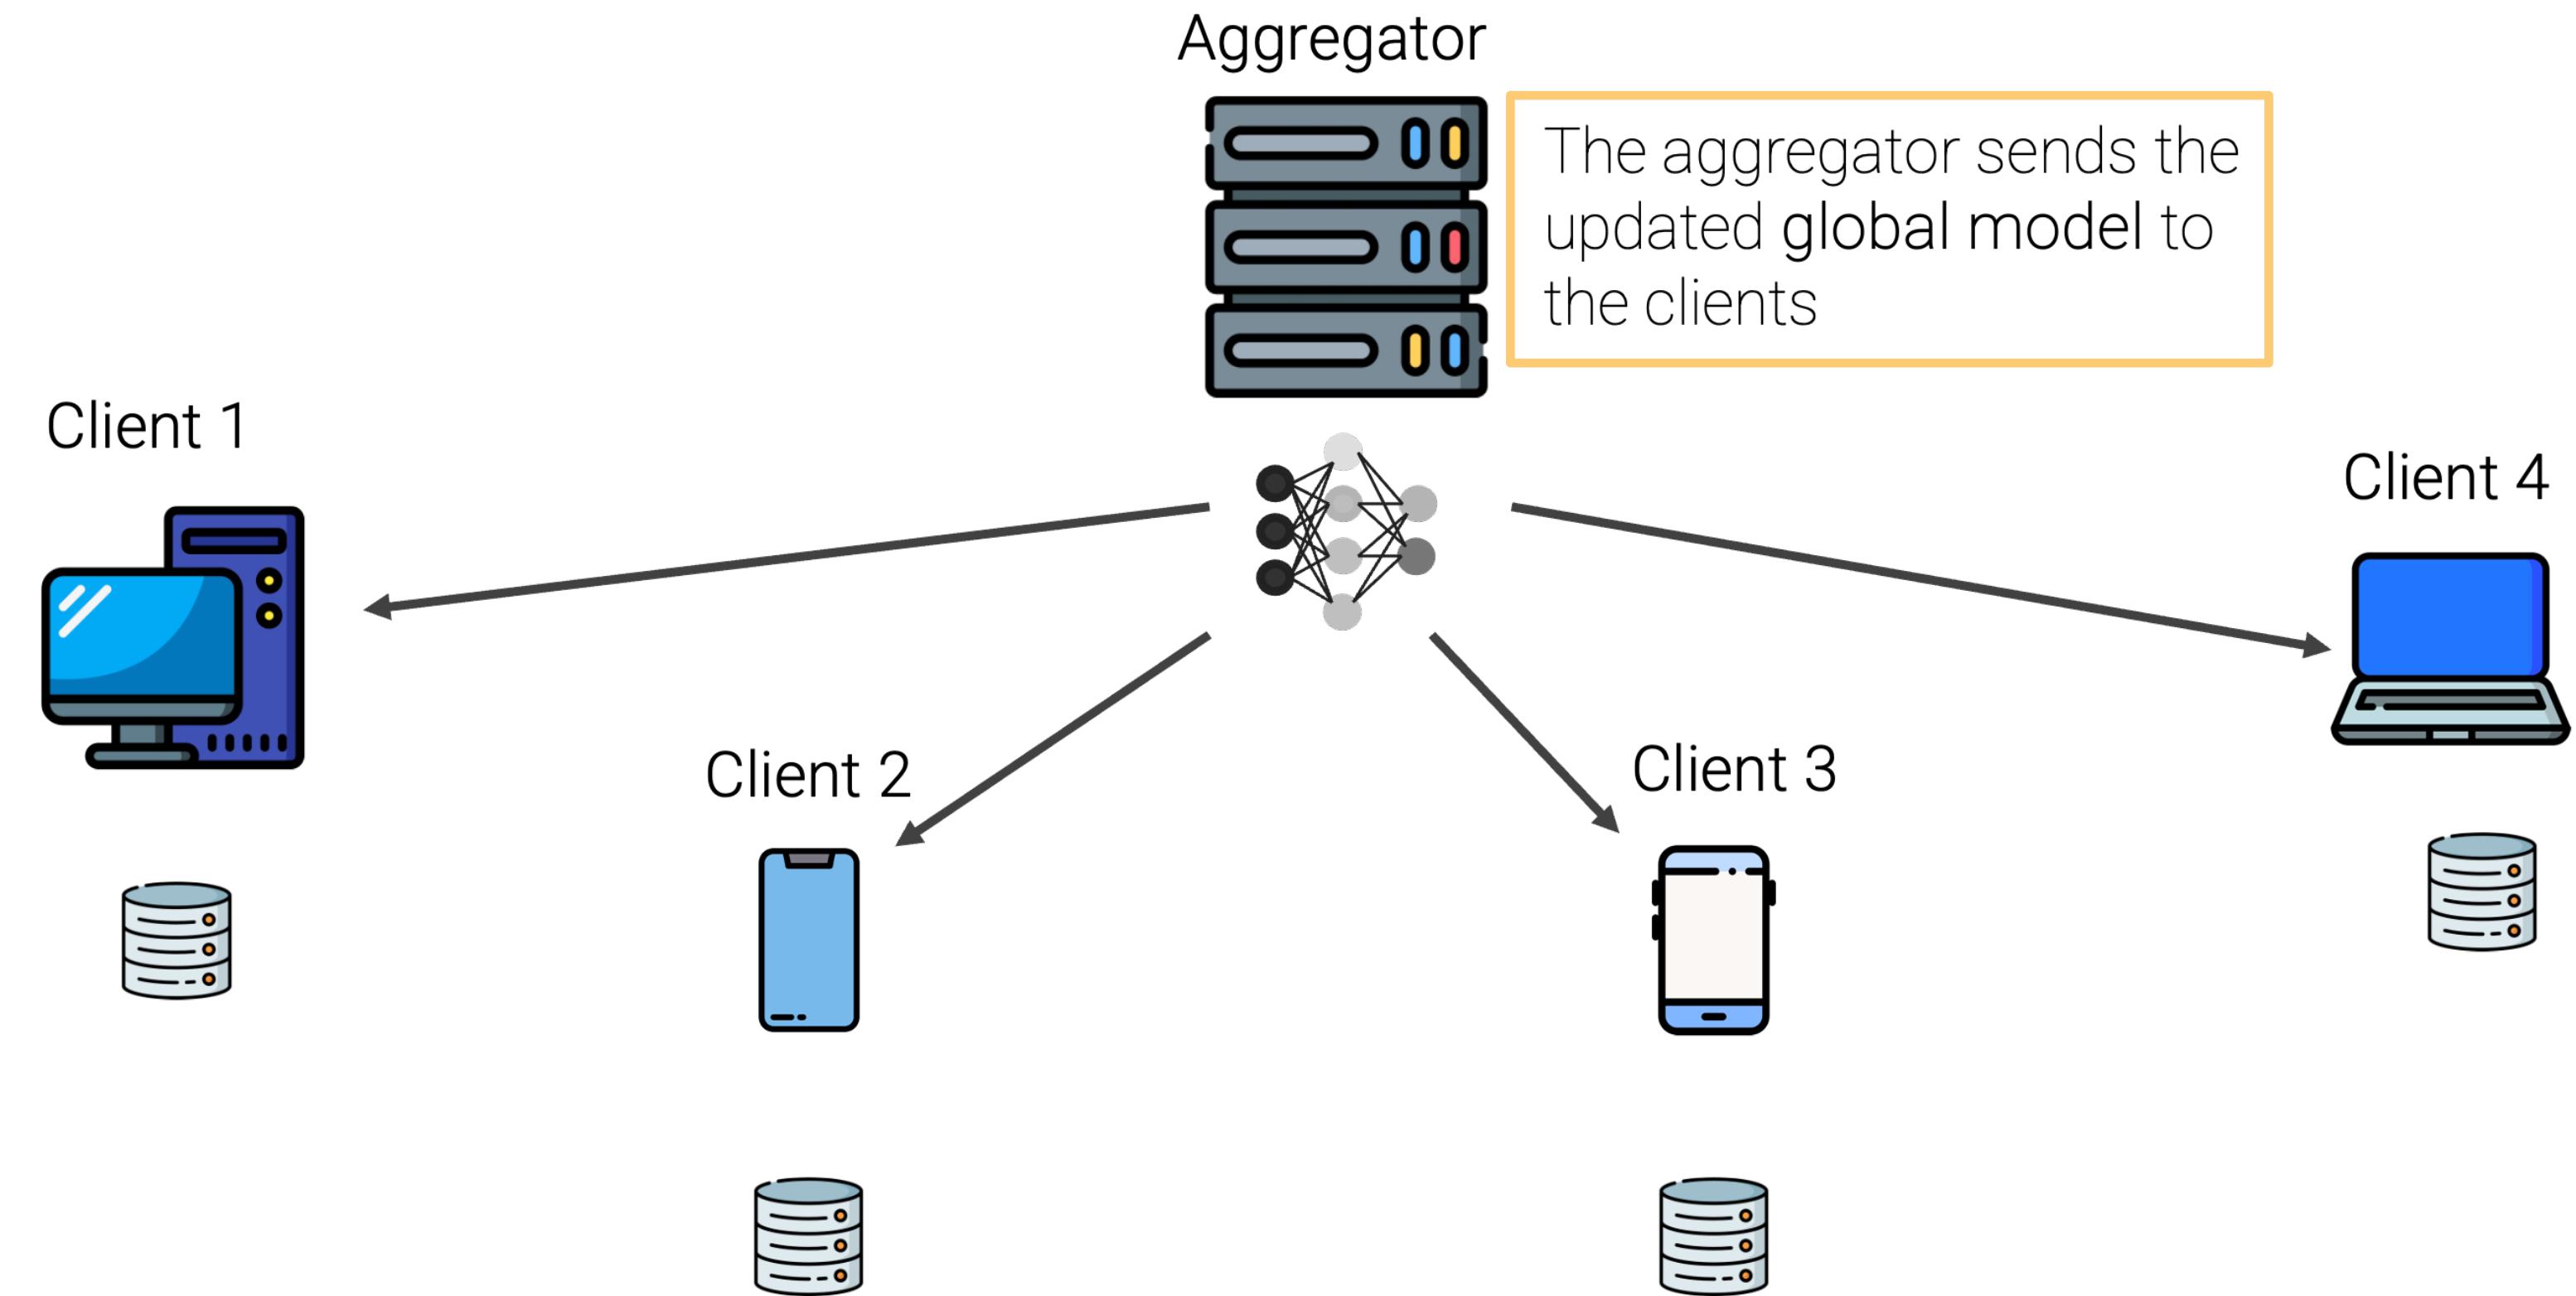
ROUND 1



Federated Learning

Training procedure

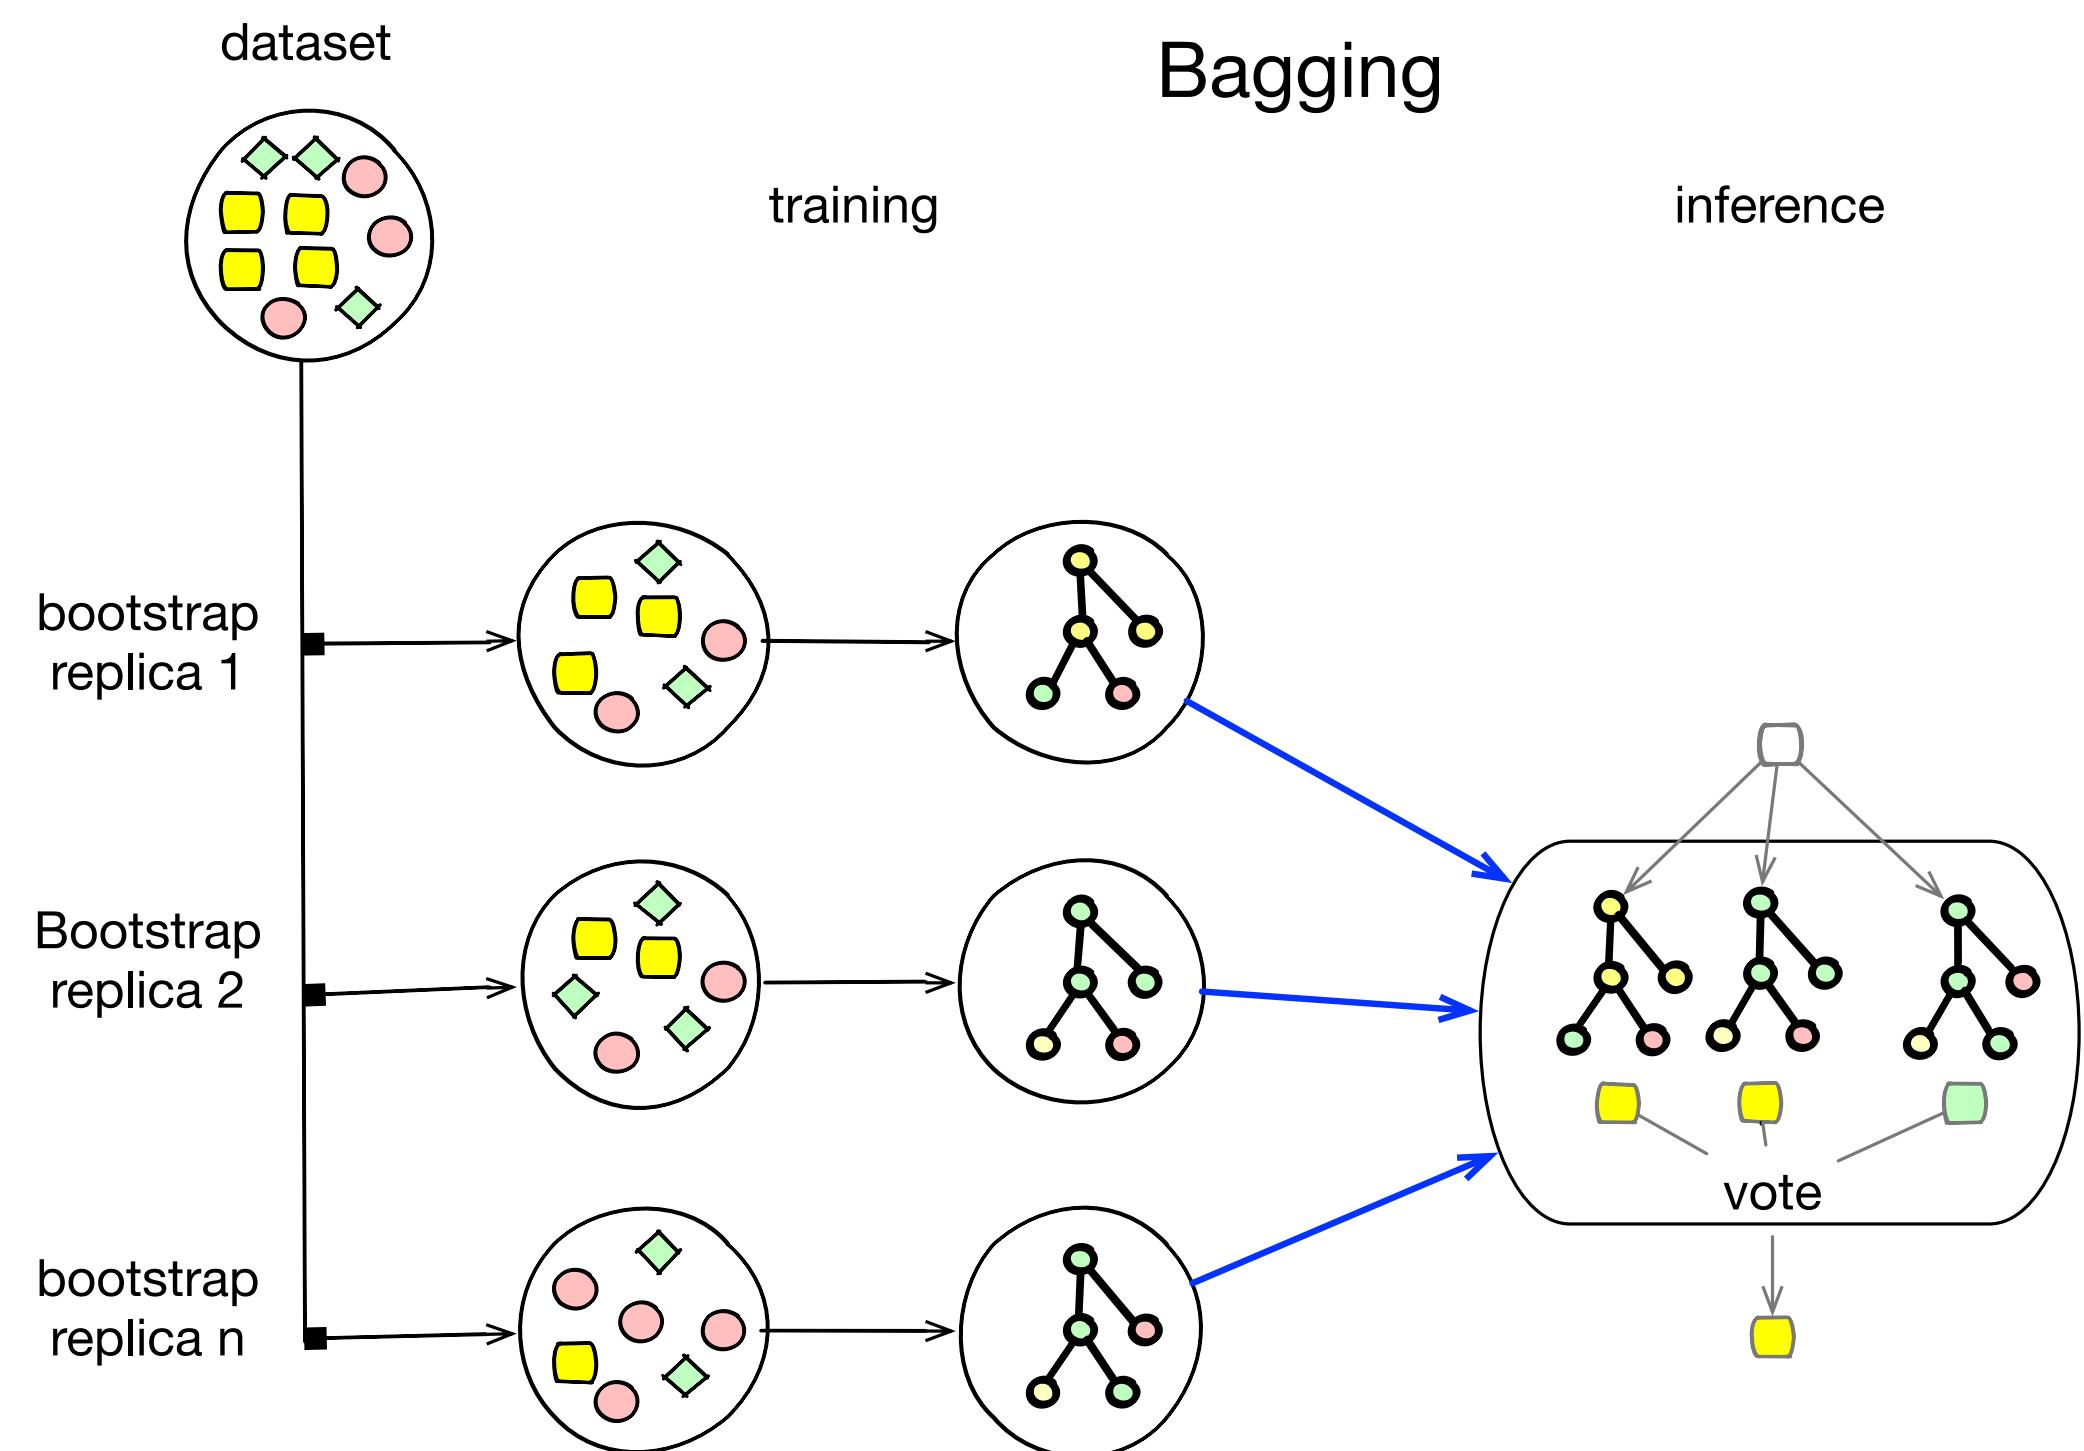
ROUND 2...



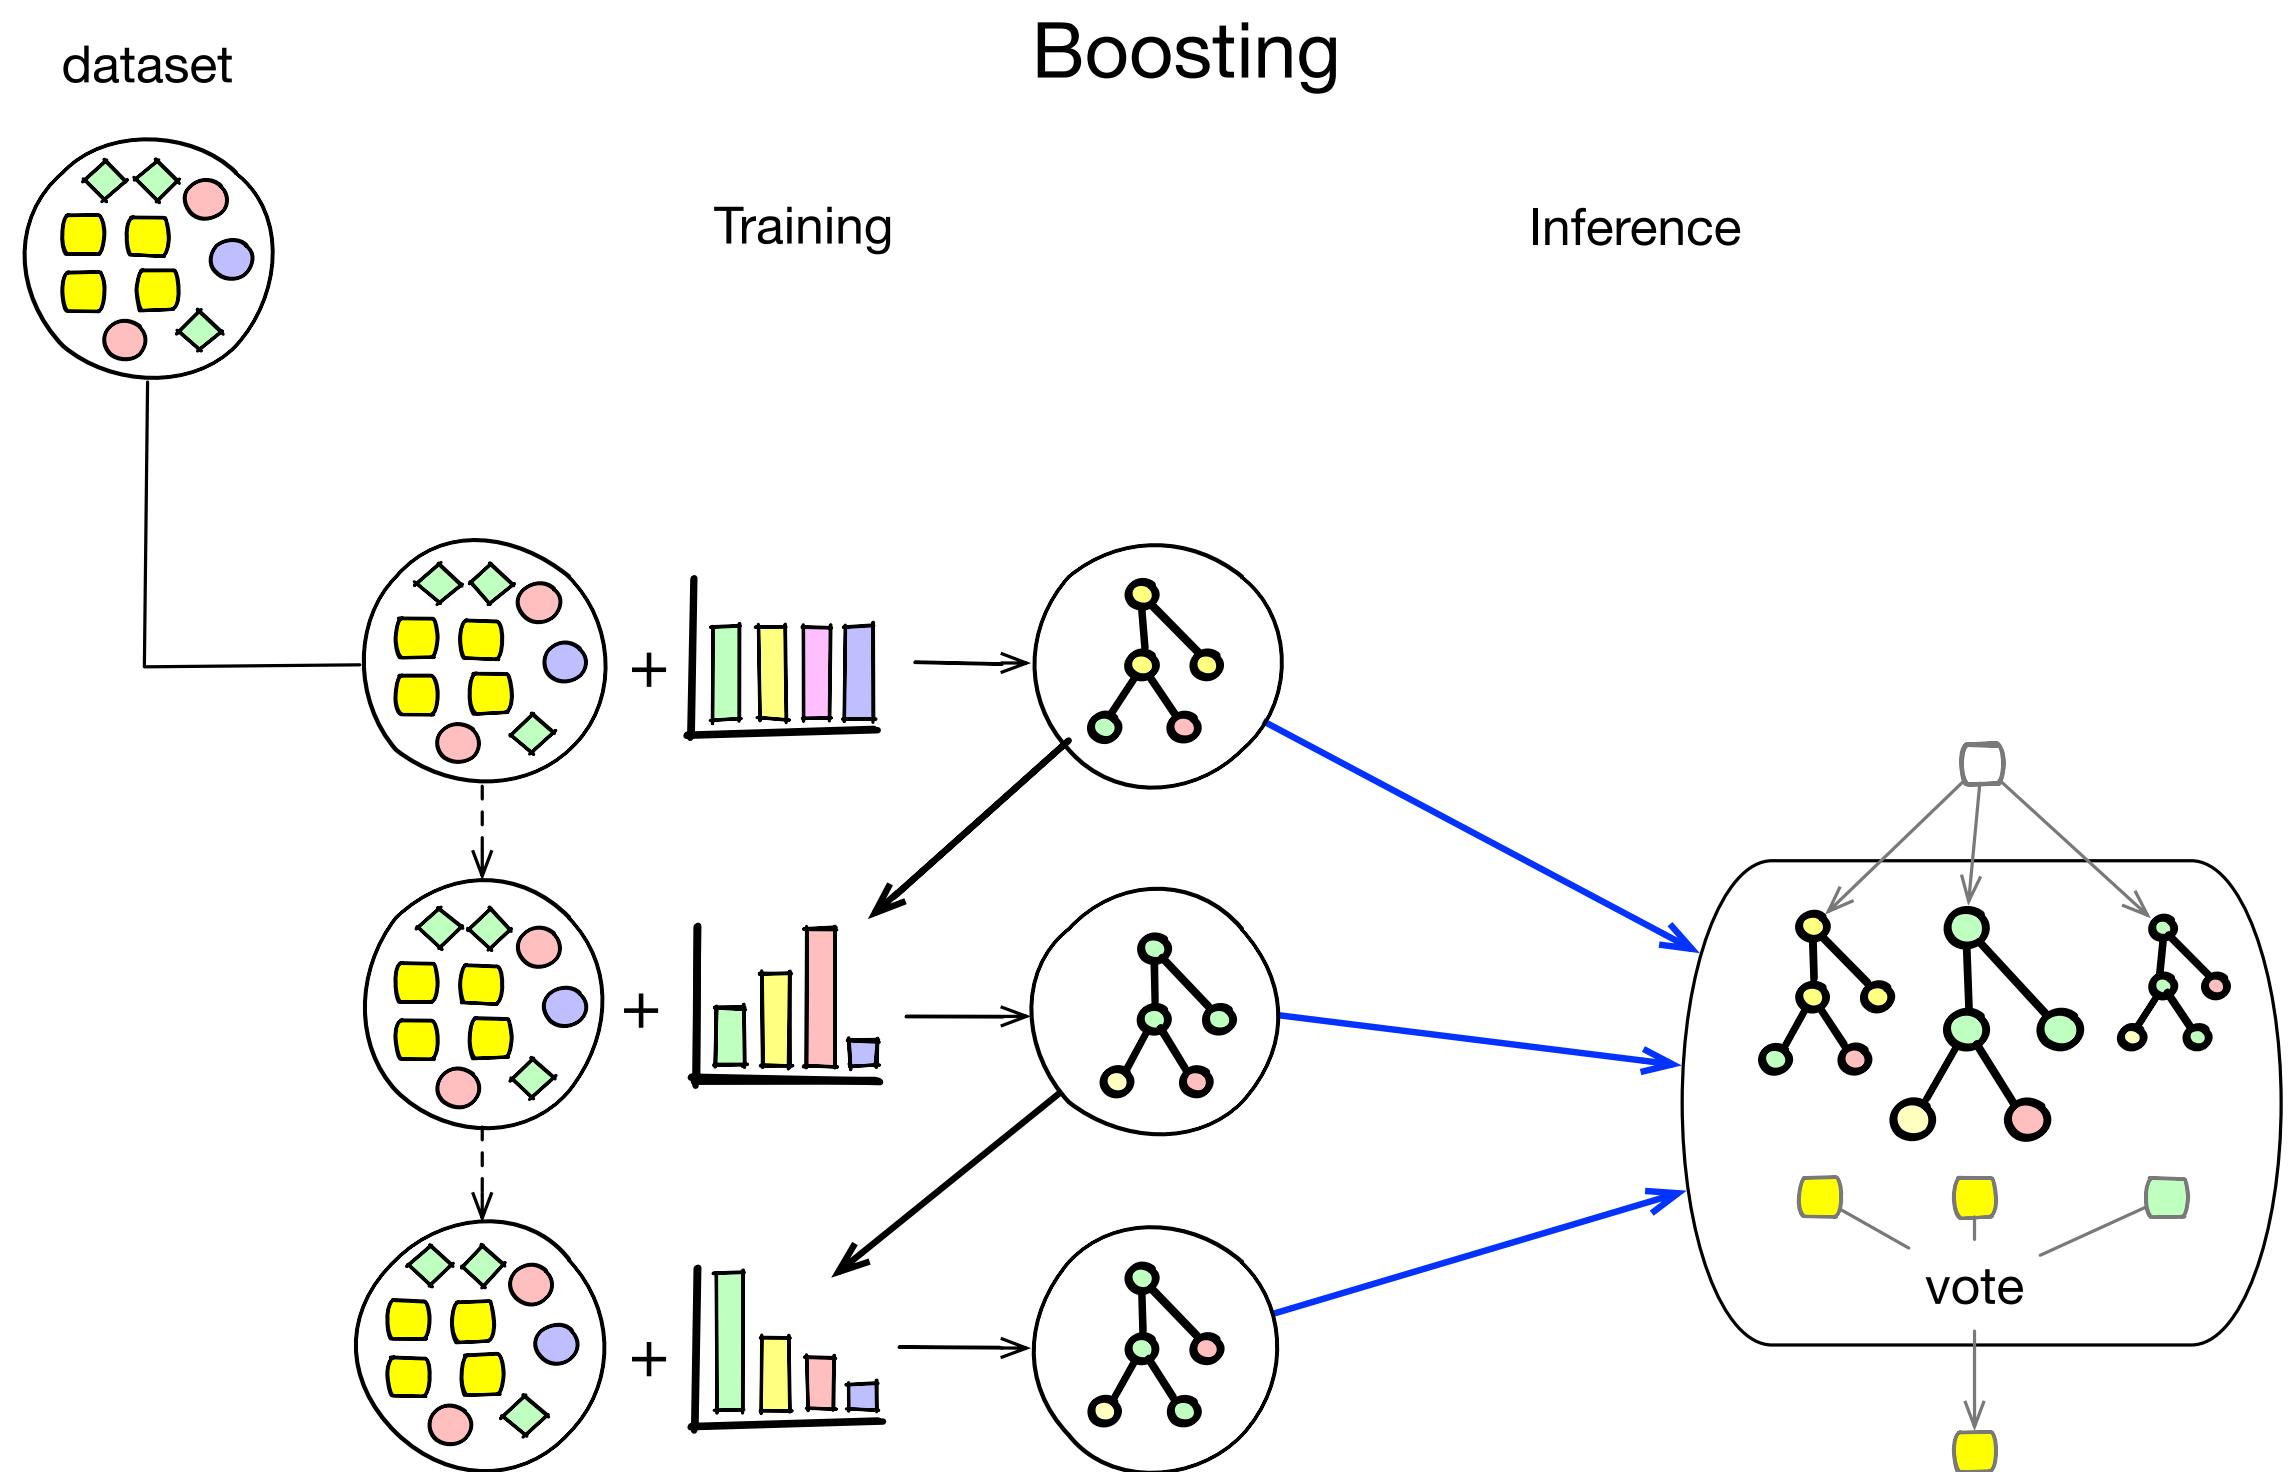
Federated Learning

Not only neural networks

Parallel



Sequential

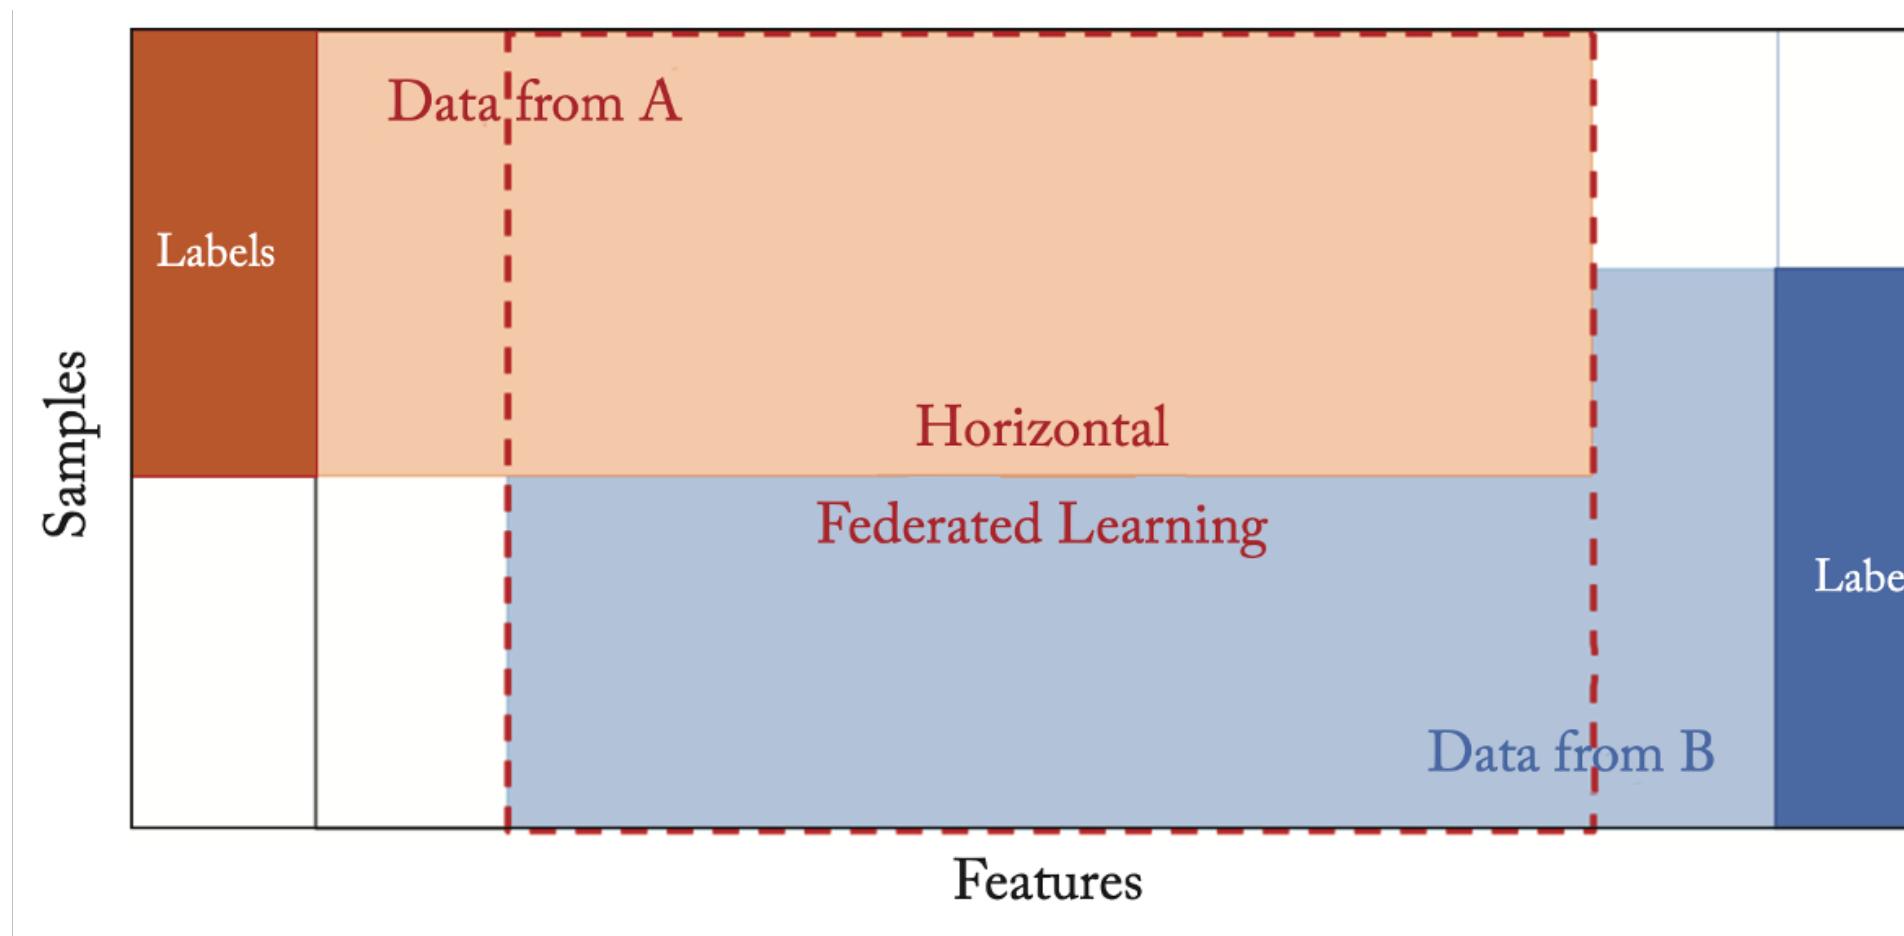


Federated Learning

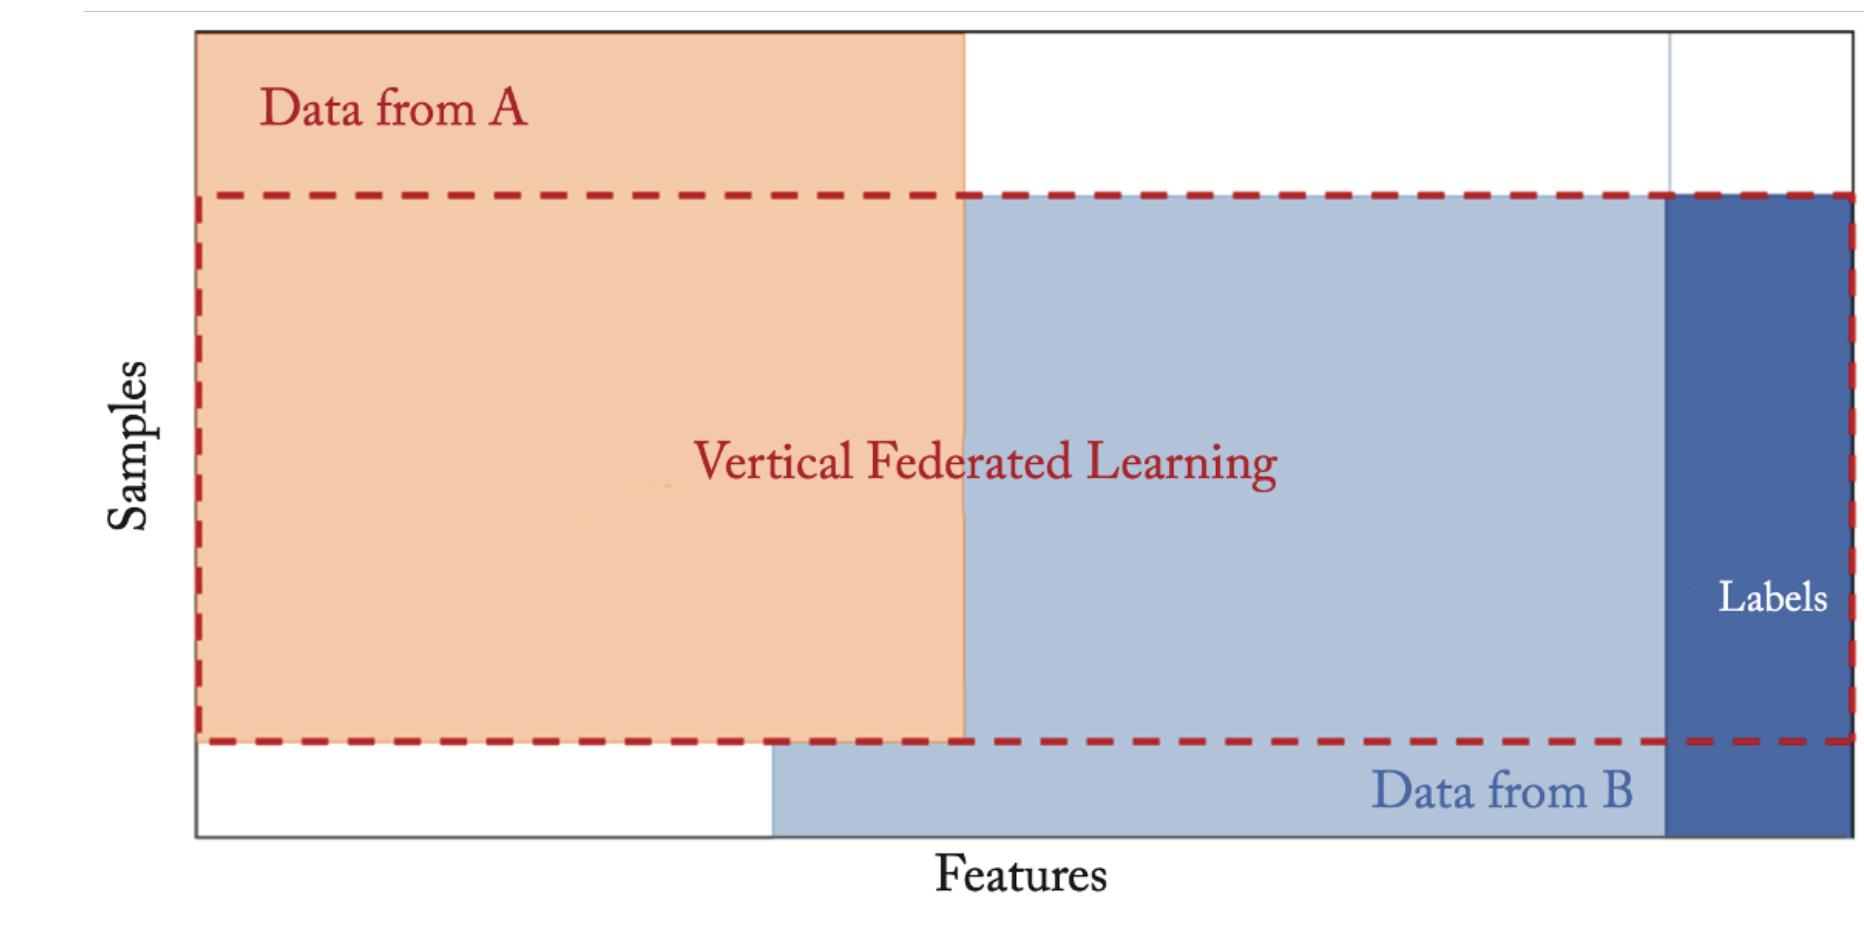
Data Categories

- In **Horizontal FL**, clients share the same feature space, while they differ in the sample space.
- In **Vertical FL**, clients share the same sample space, while they differ in the feature space.

Horizontal



Vertical

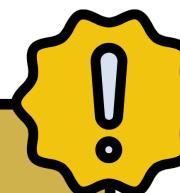


Federated Learning

Deployment scenarios

- FL can be divided into cross-silo and cross-device.
 - **Cross-silo** FL involves a small number of organizations (silos) like corporations or medical institutions. Few clients with large local datasets.
 - **Cross-device** FL involves large scale and often low-power devices like smartphones. Many clients with relatively small local datasets

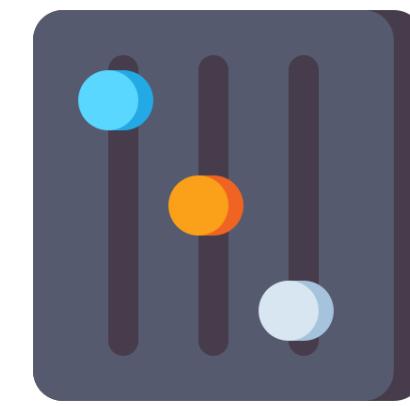
fluke framework can work under all the possible described cases
(paradigms, categories, deployments, traditional ML models, ...)



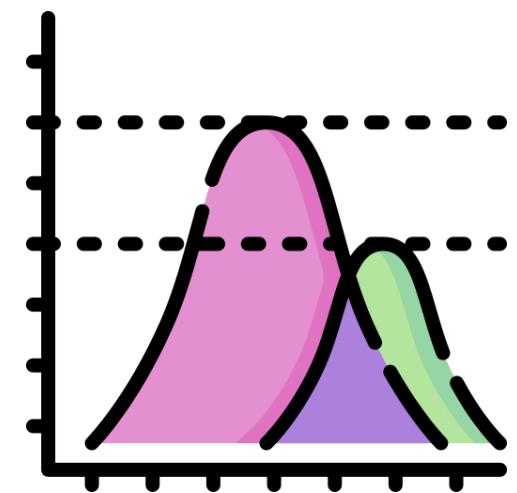
Federated Learning Challenges



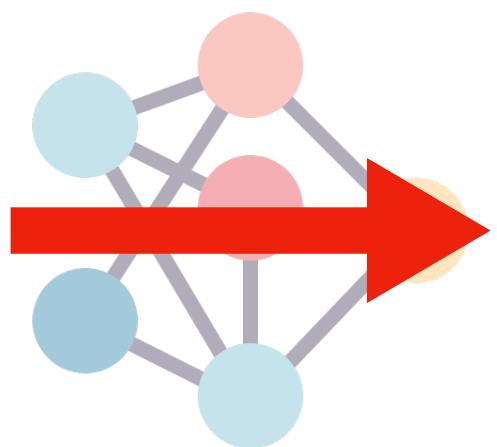
Security & Privacy



Hyper-parameters
tuning



Non-iidness



Beyond
neural networks



Coordination

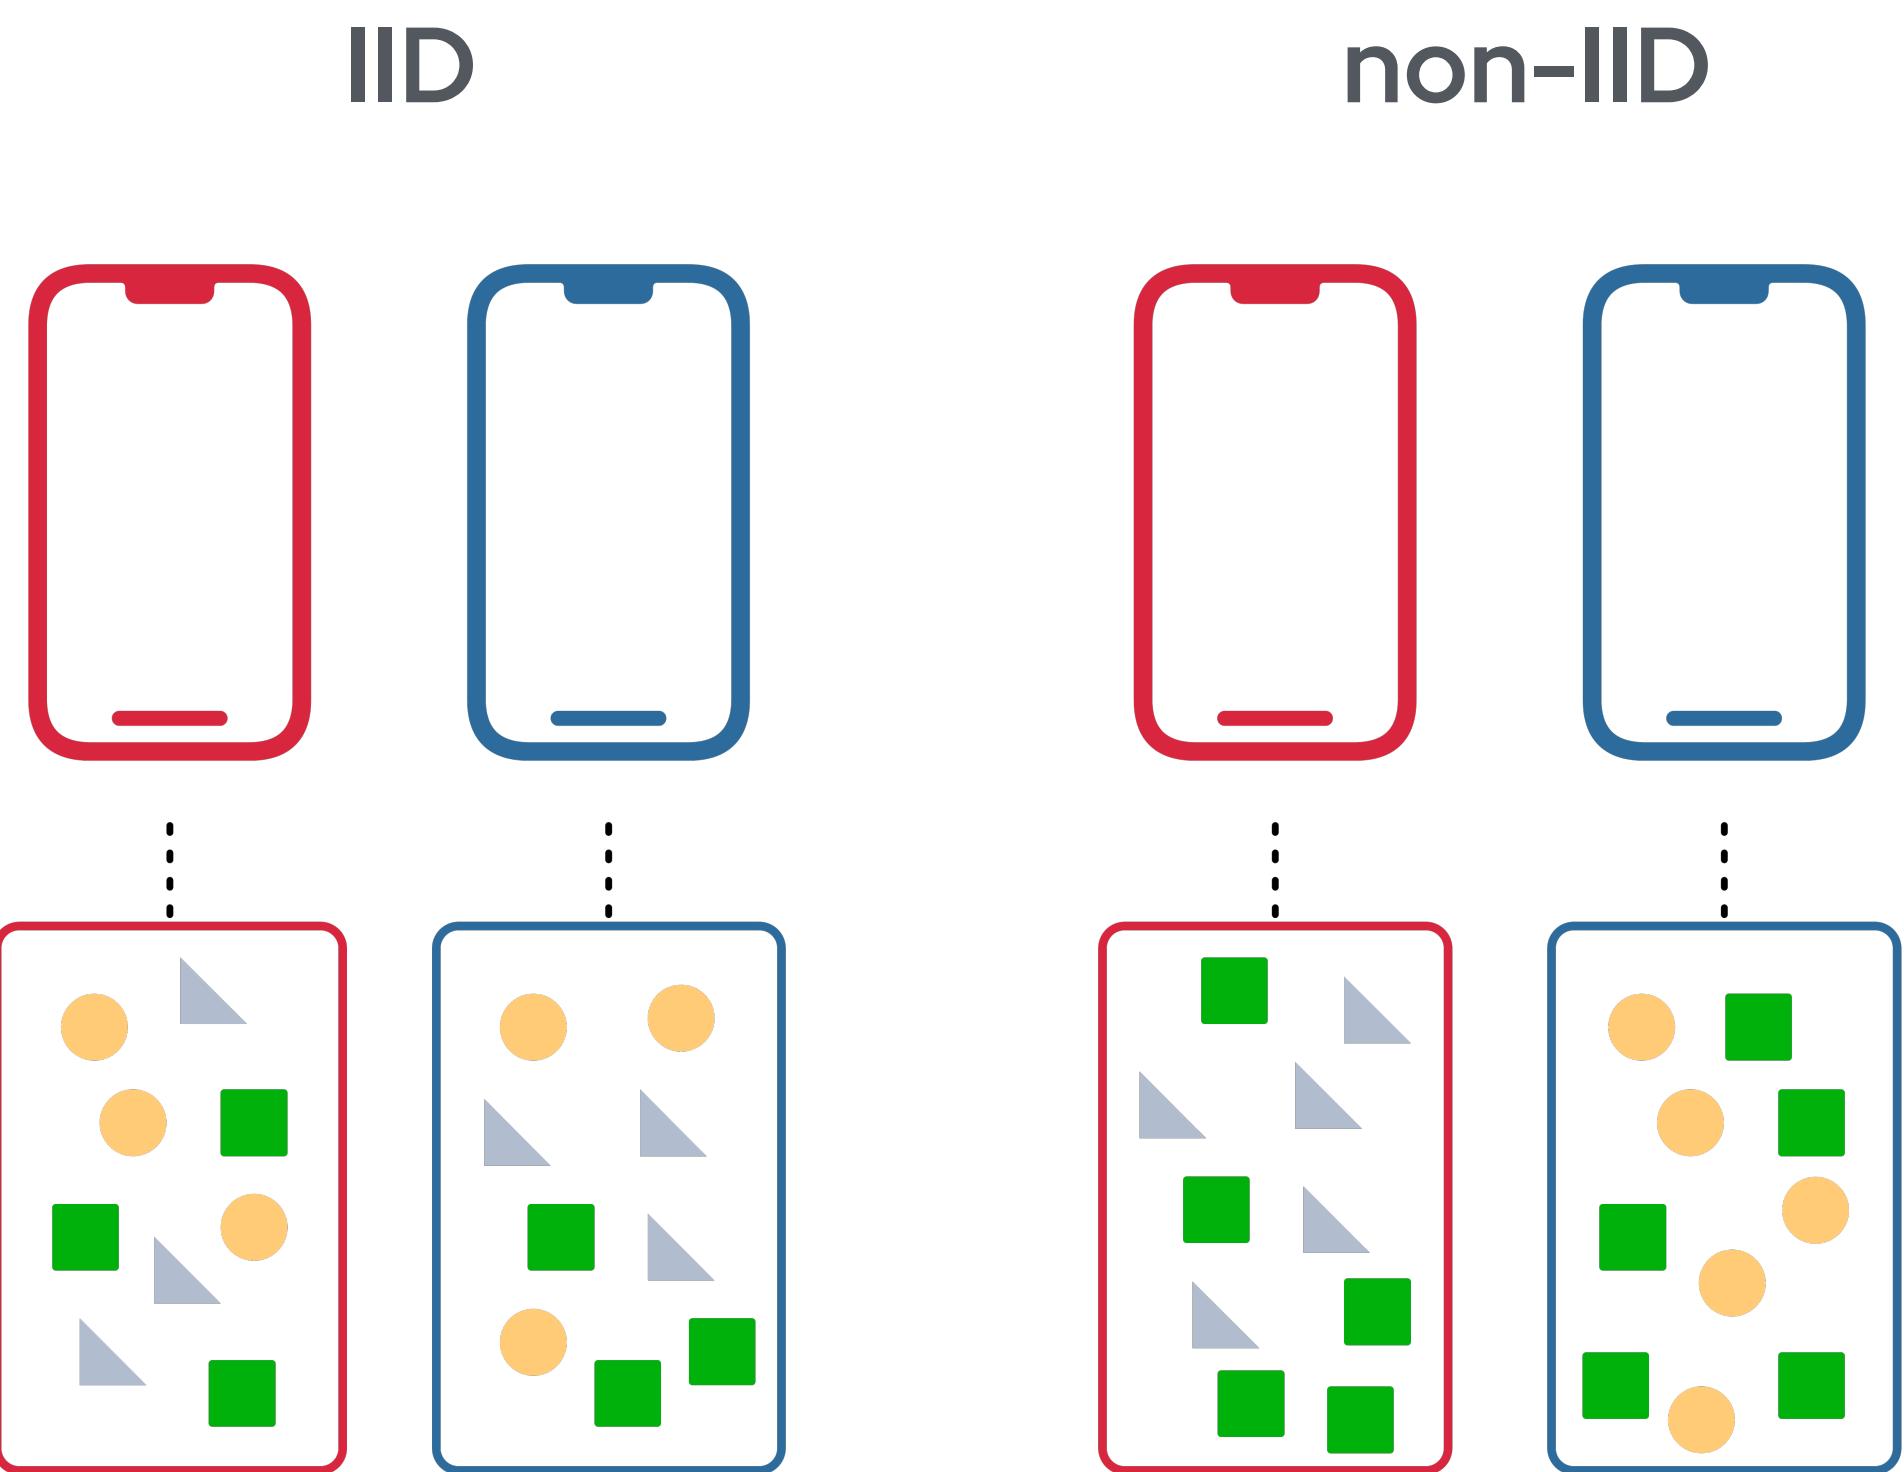


Fairness

Federated Learning

Data Distribution

- Data is considered **IID** when each data point is drawn **identically and independently** from the same probability distribution.
- Each sample is statistically independent of the others.
- All samples follow the same underlying distribution.
- However, in **real-world scenarios, the IID assumption does not hold.**
- **Non-IID** data come from different probability distributions (i.e., one per client).

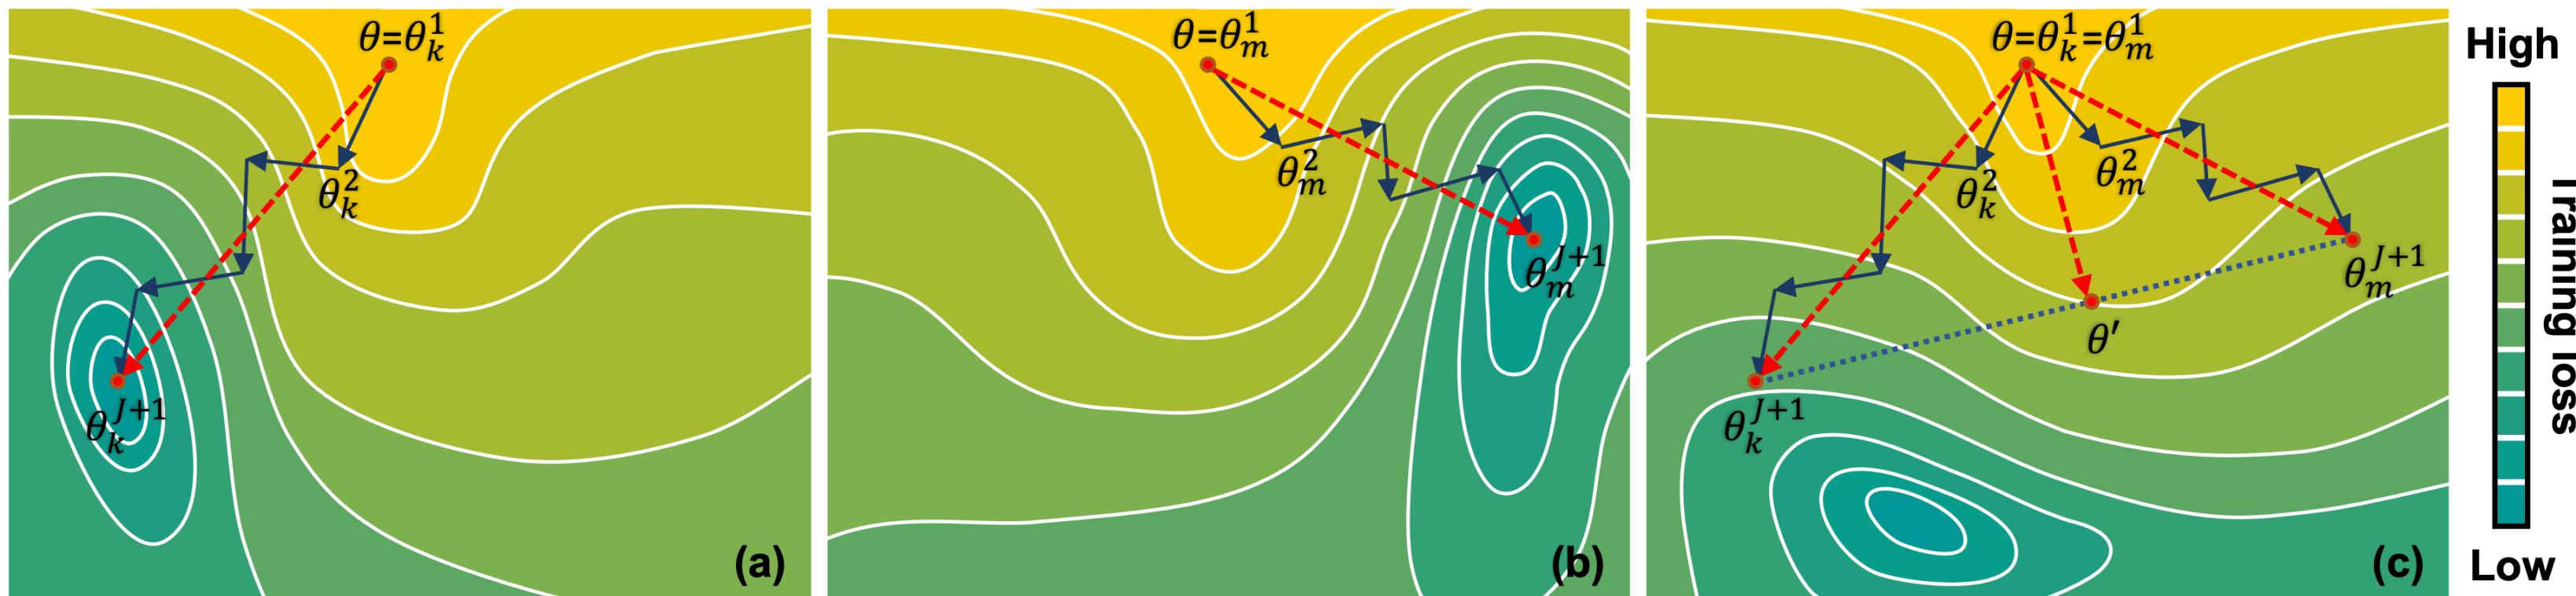


Federated Learning

Data Distribution

- Local data may be really different in terms of distribution, features, and classes.
- Non-IID data cause unstable and **inefficient training**.

- **Quantity skew**: different clients hold different amounts of data.
- **Prior shift**: the distribution of labels varies across clients.
- **Covariate shift**: the input features distribution changes across clients.

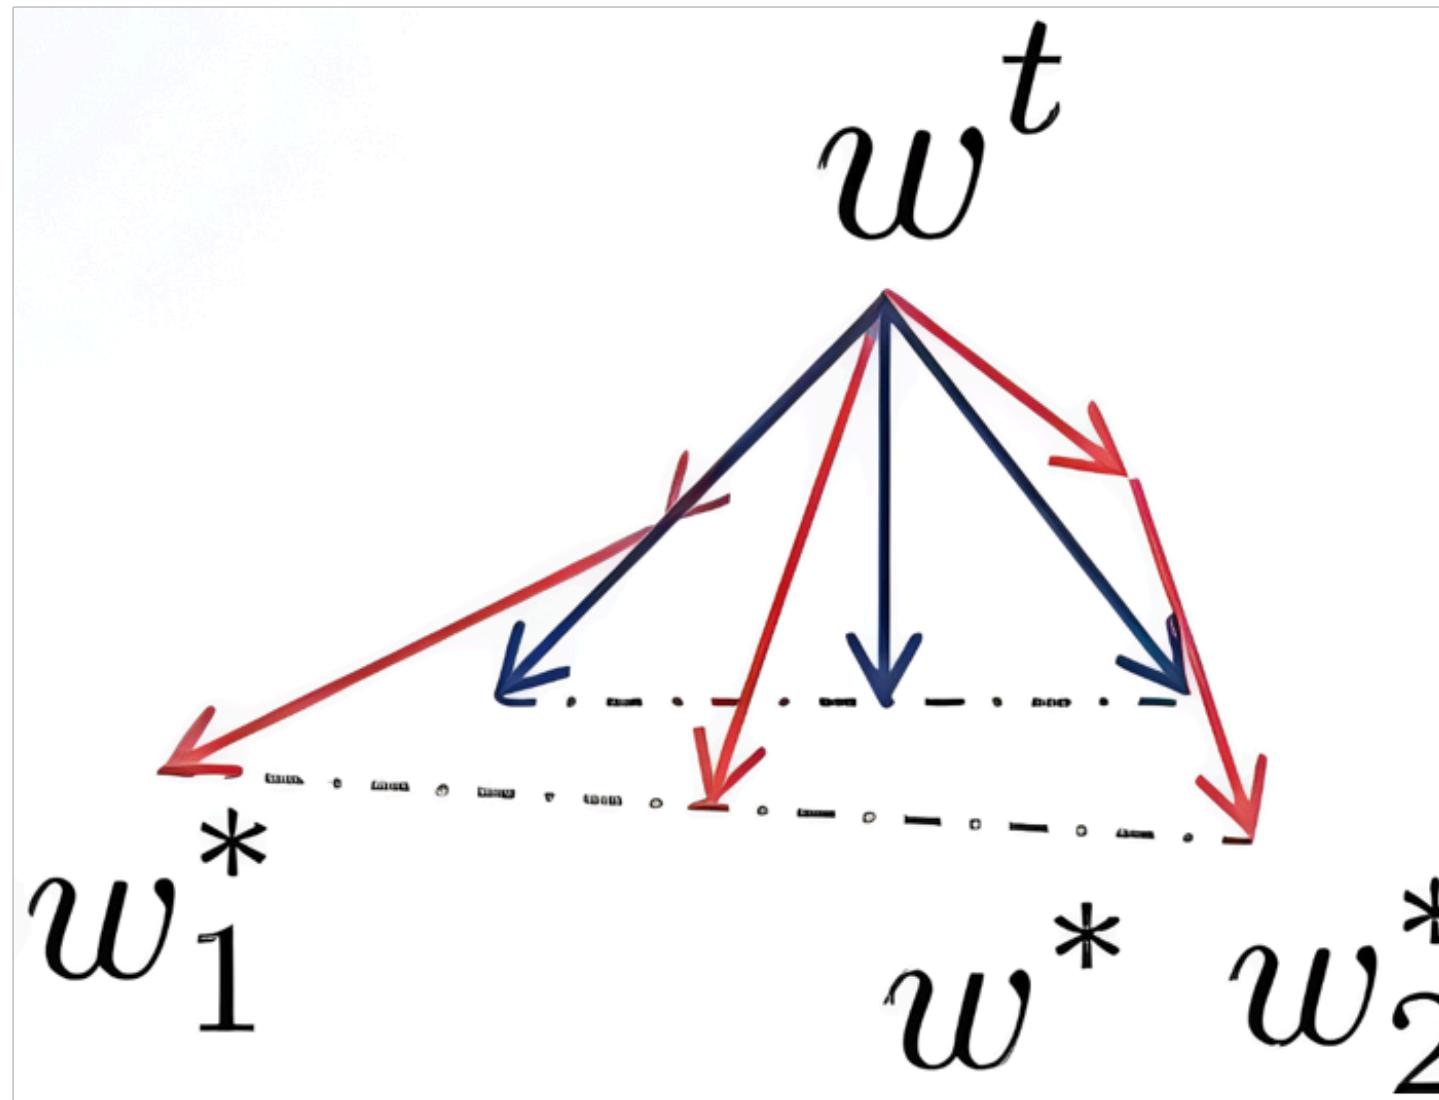


Dealing with non-IID distributions

FedProx

FedProx is a FedAvg variant for heterogeneous clients

$$\min_{w_t^k} \mathcal{L}(w_t^k; \mathcal{D}_k) + \frac{\mu}{2} \|w_t^k - w_{t-1}\|_2^2$$



- **Proximal term** limits the impact of heterogeneous local updates
- Allows to avoid the "device drop" by incorporating partial updates from slow devices
- Theoretical guarantees on convergence

Federated Learning

Privacy risks

Server

Honest-but-curious: a party that adheres strictly to the prescribed FL protocol without any malicious alterations, but harboring intentions to infer private information.

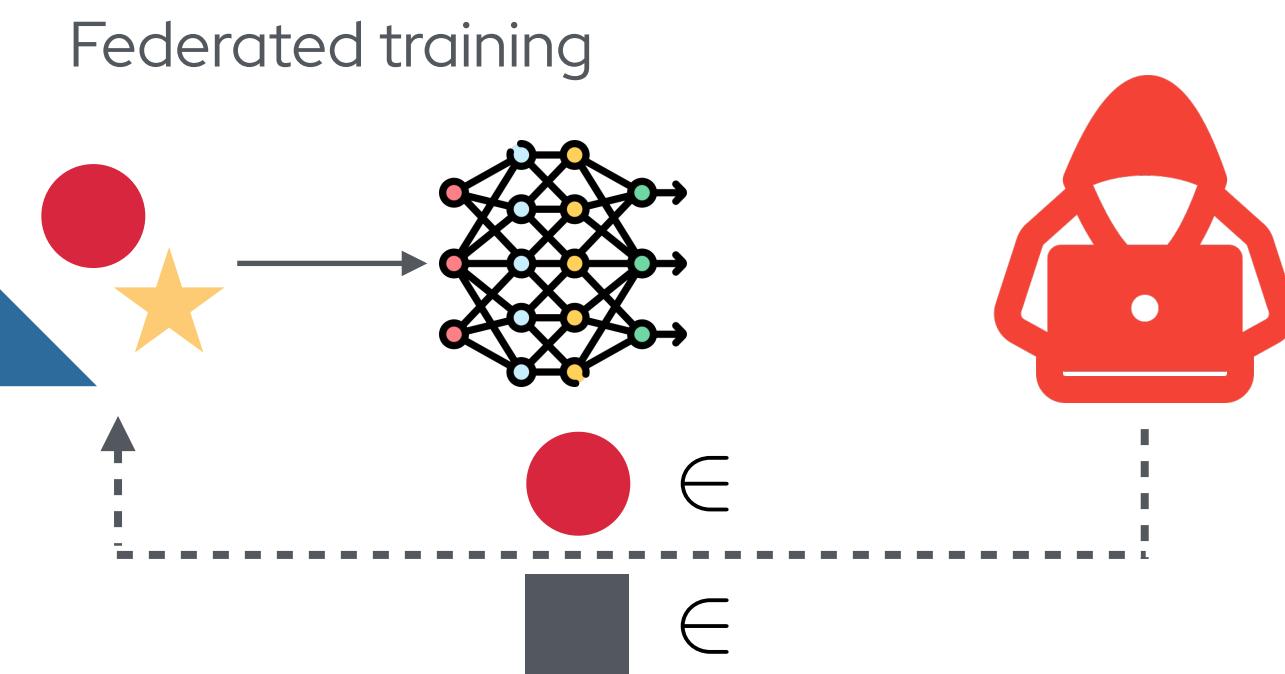
Clients

Malicious: actively seeks to undermine the privacy and integrity of the FL process.

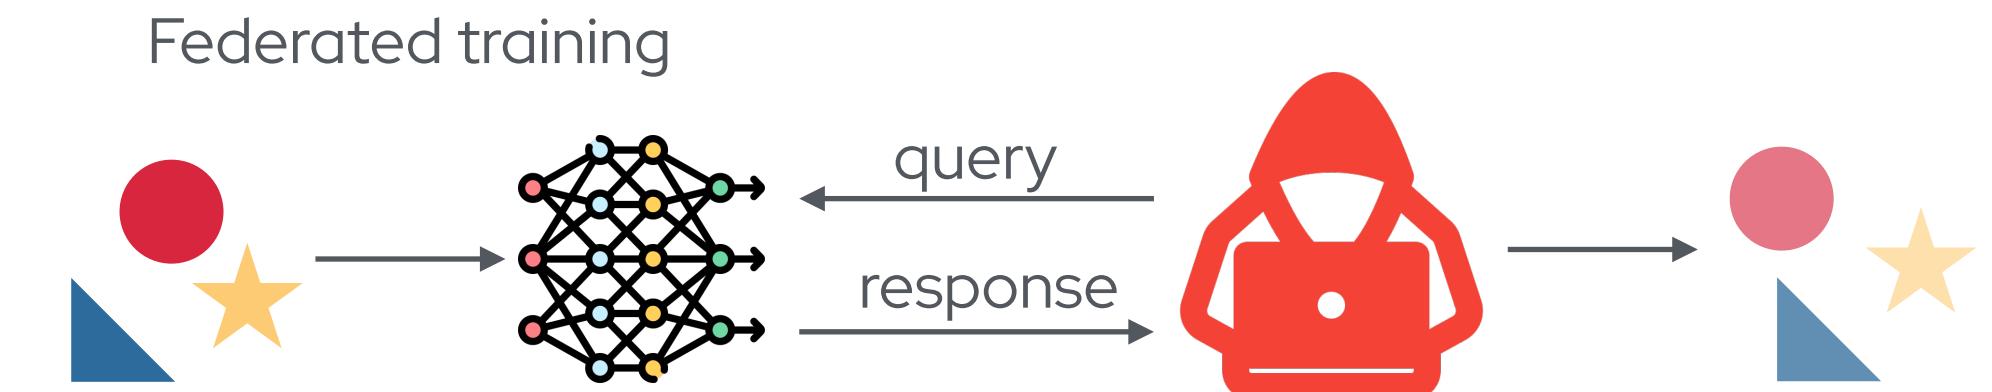
Federated Learning

Security & Privacy - Attacks

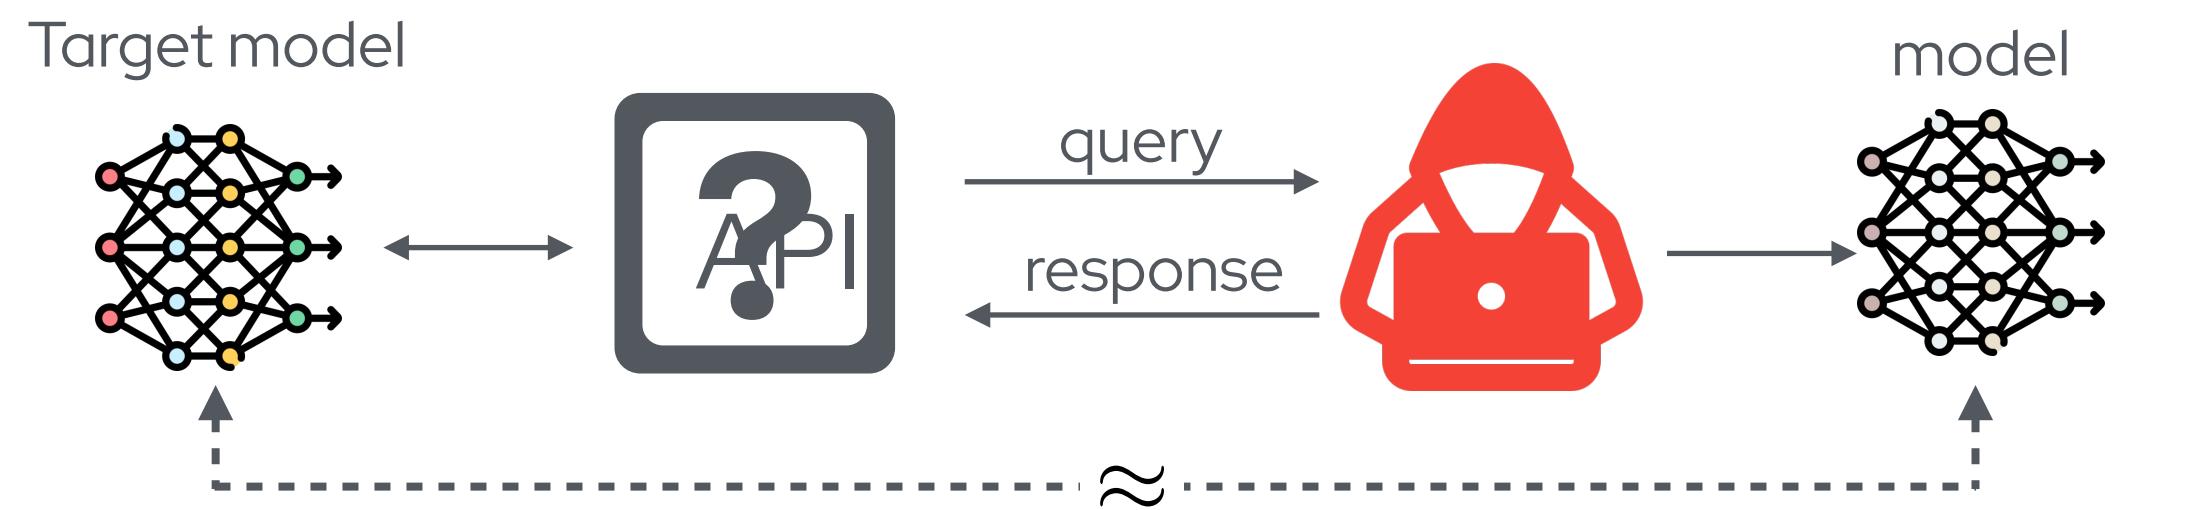
Membership inference



Model Inversion



Model Extraction

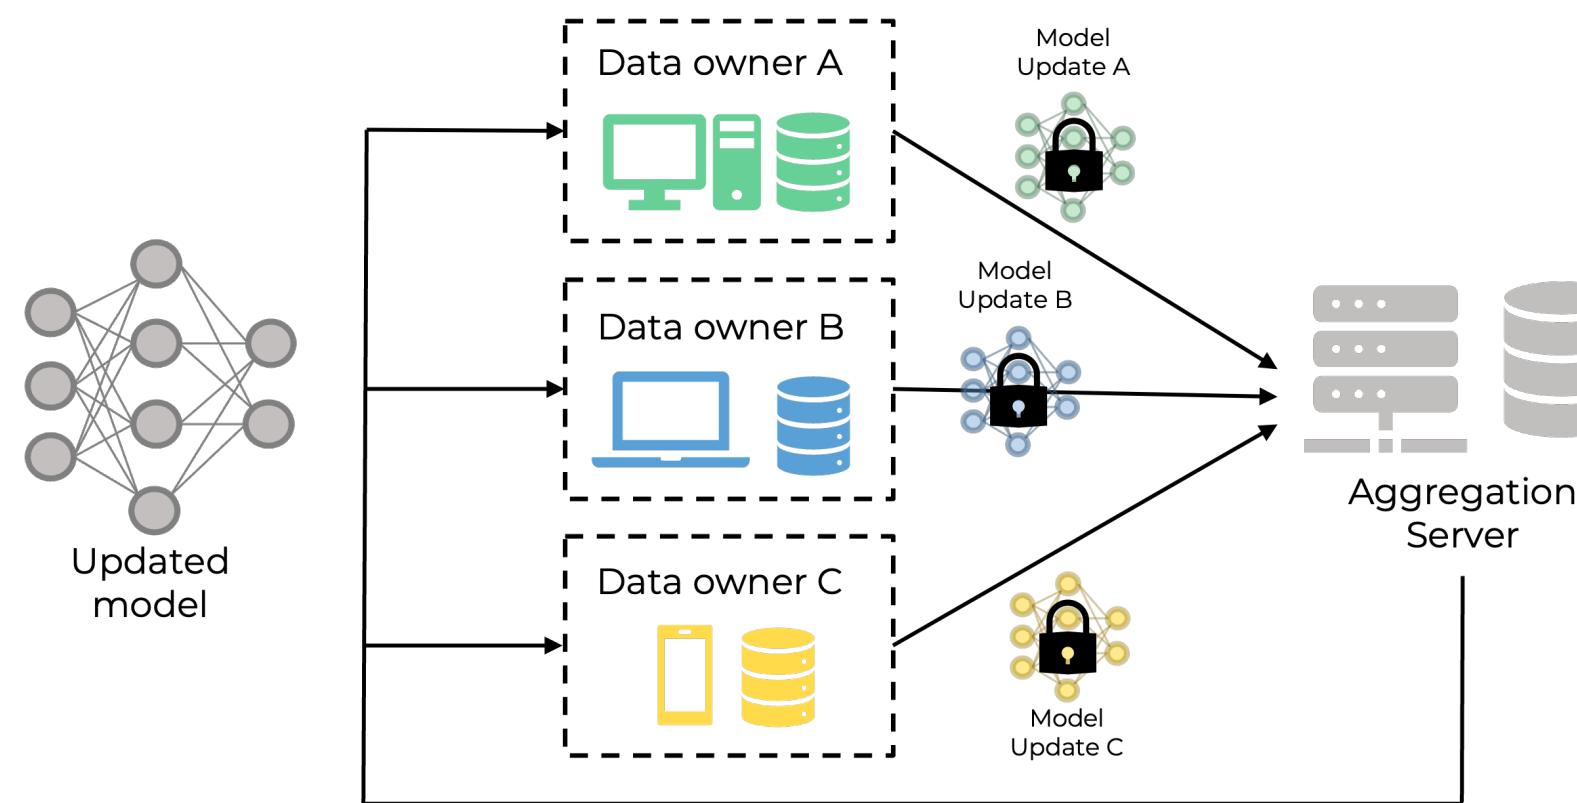


Federated Learning - Privacy

Security & Privacy - Defences

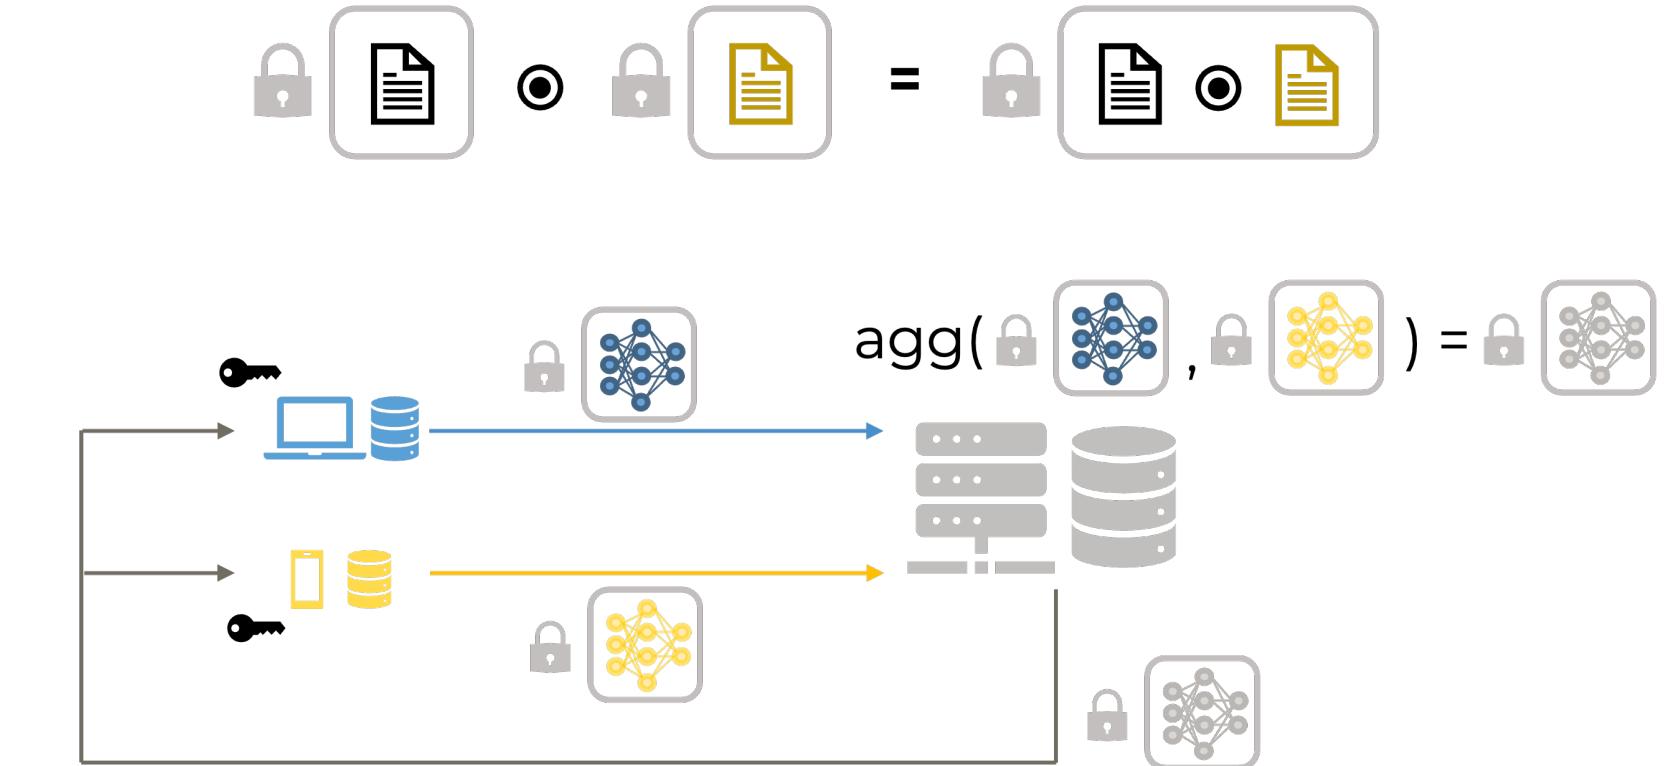
?

Differential Privacy



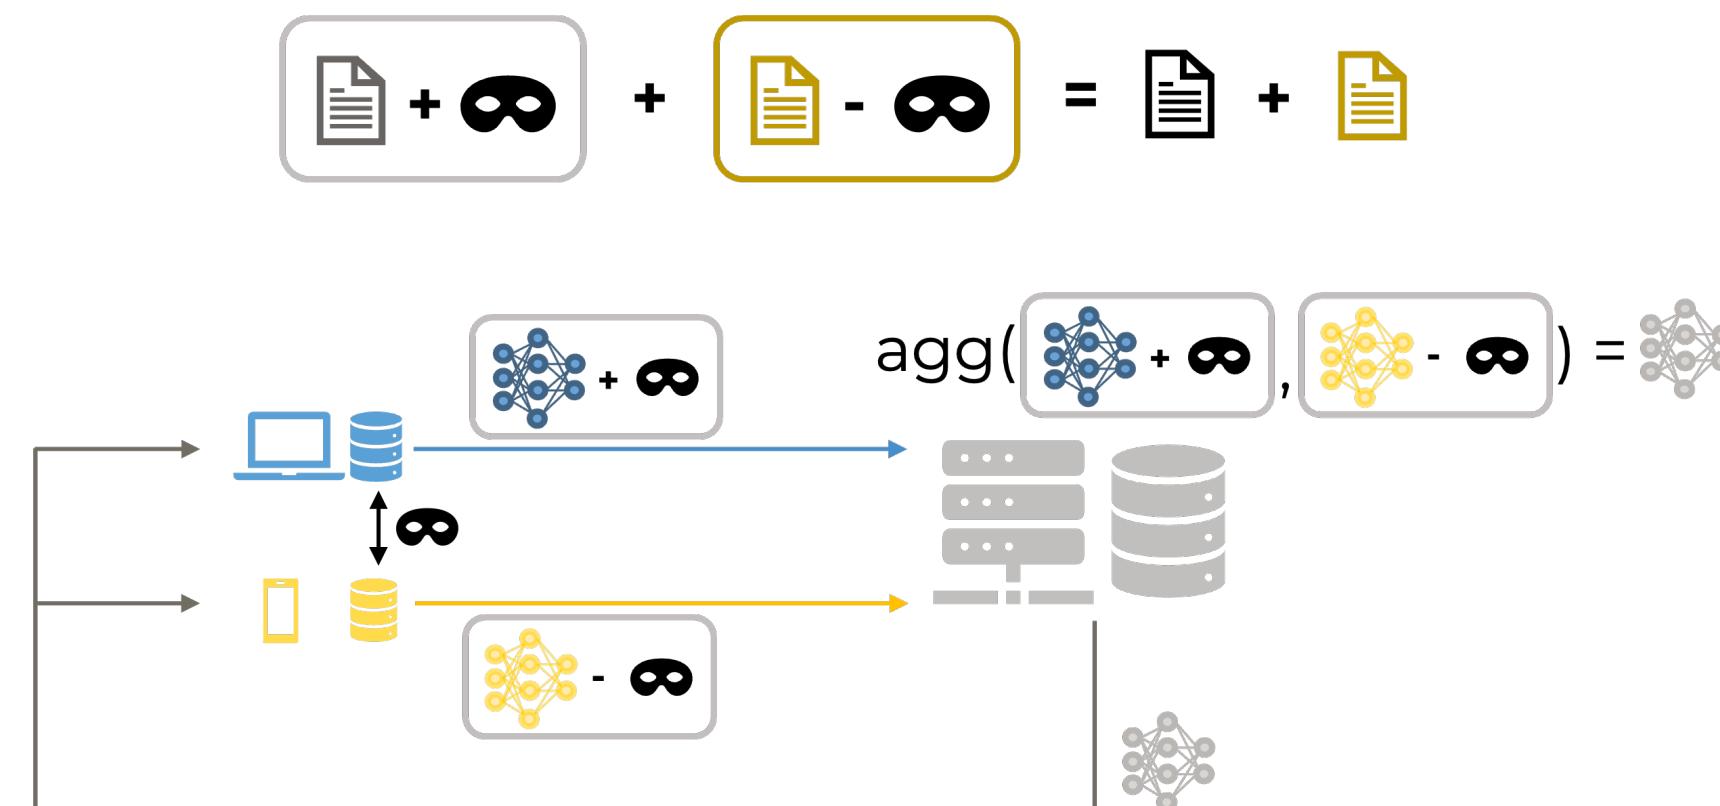
?

Homomorphic encryption



?

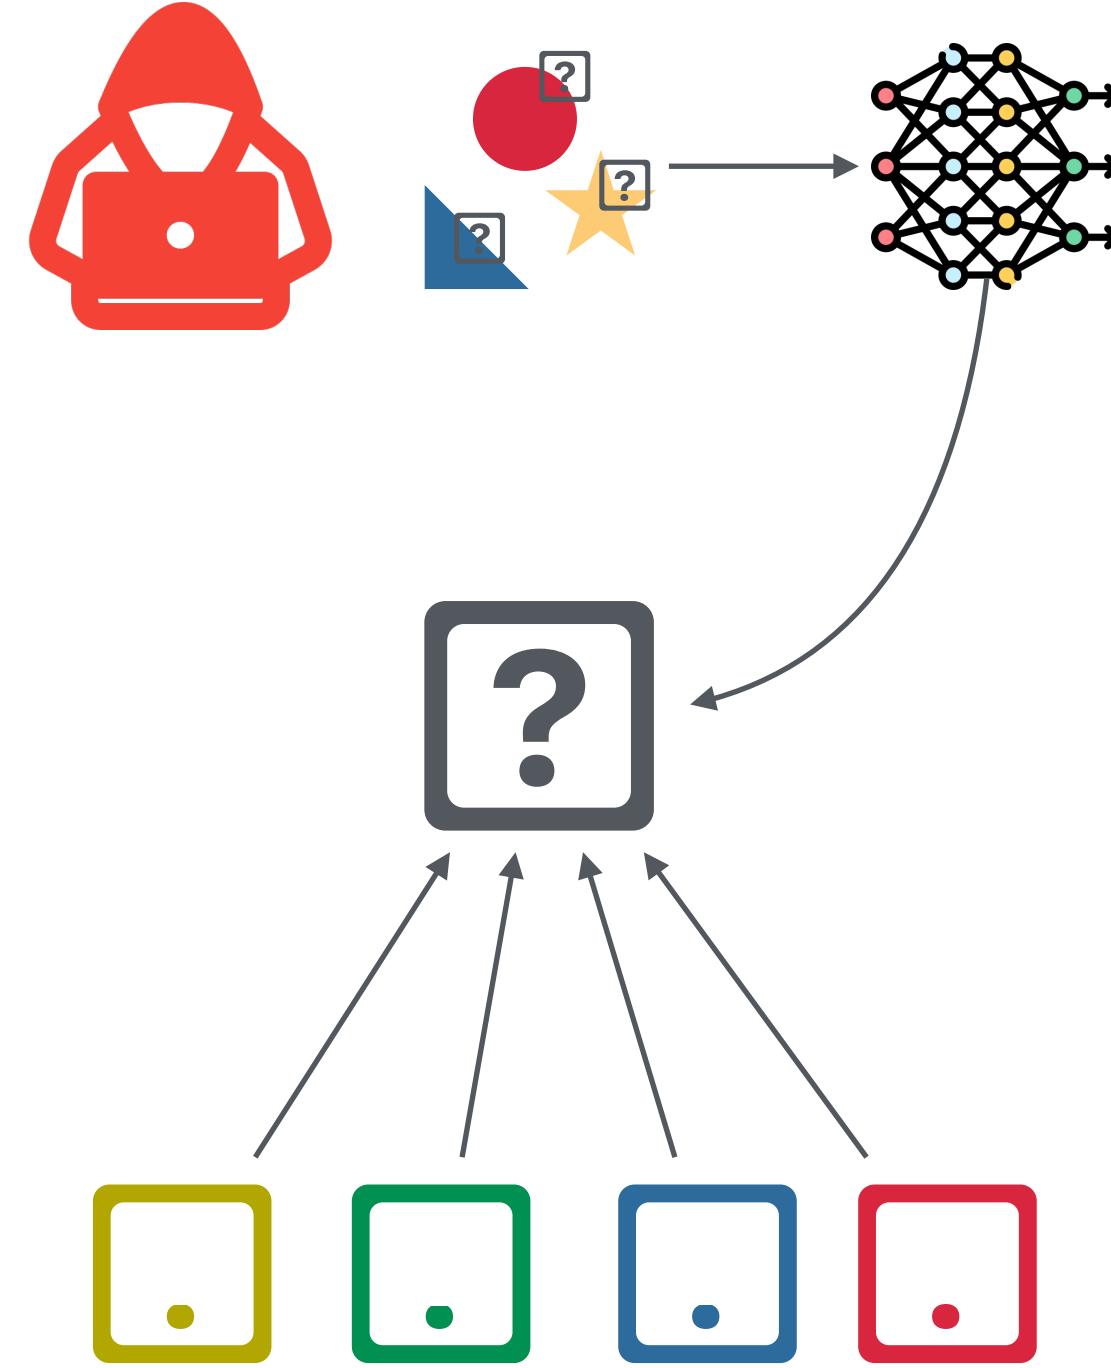
Secure aggregation



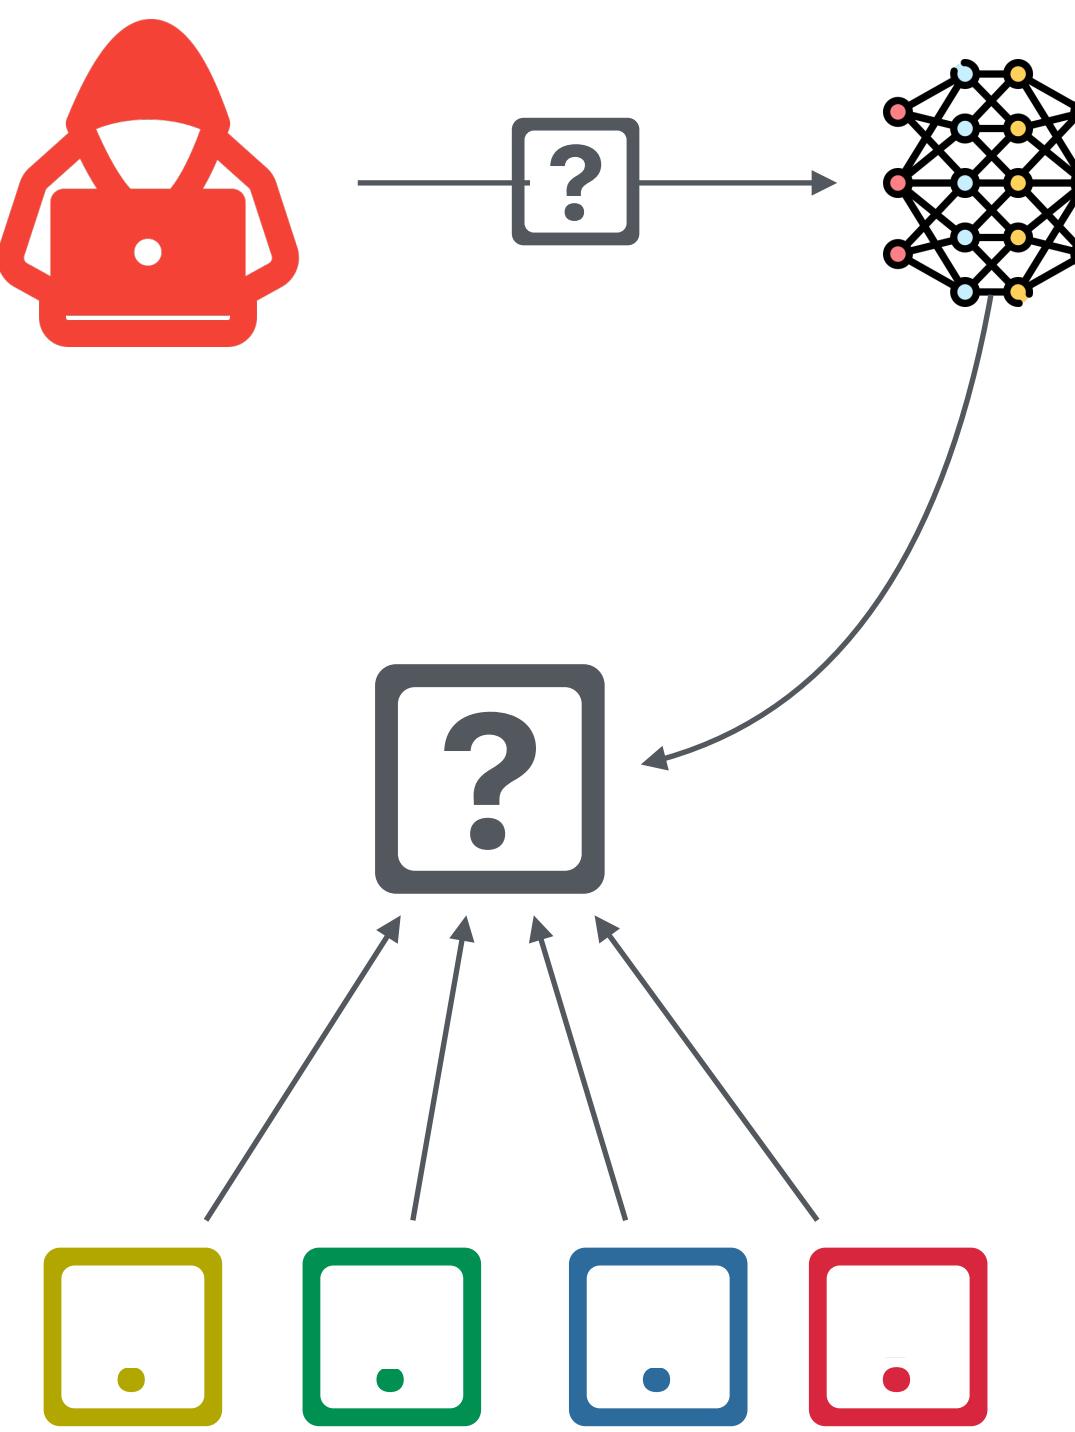
Federated Learning

Robustness

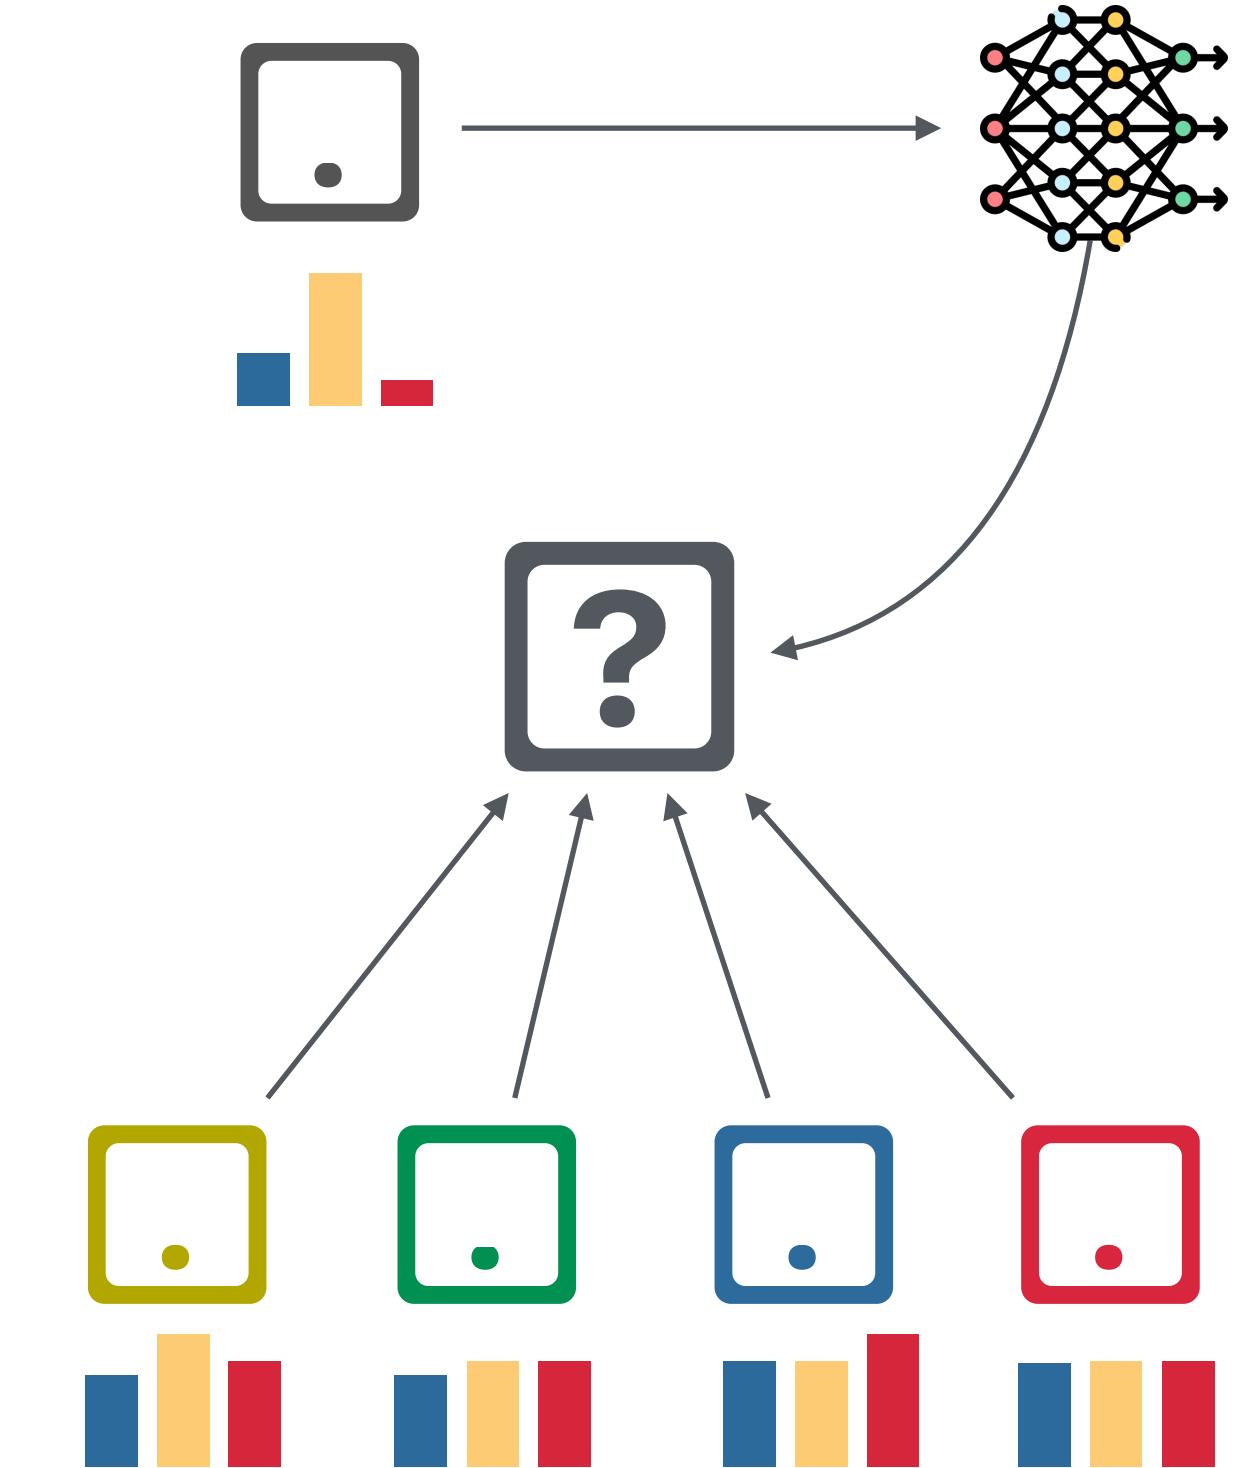
Adversarial Data



Adversarial Model



Noisy updates



Federated Learning

Robustness - possible approaches

Server-side

- Client selection
- Aggregation approach modification
- Aggregation Hyperparameter Tuning

Clients-side

- Objective regularization
- Ad-hoc optimizer
- Data Selection
- Manifold

Federated Learning

Robust aggregation

The **mean** (as in FedAvg) **is not robust** to outliers! A malicious client can deviate the model in a consistent way

Possible alternatives

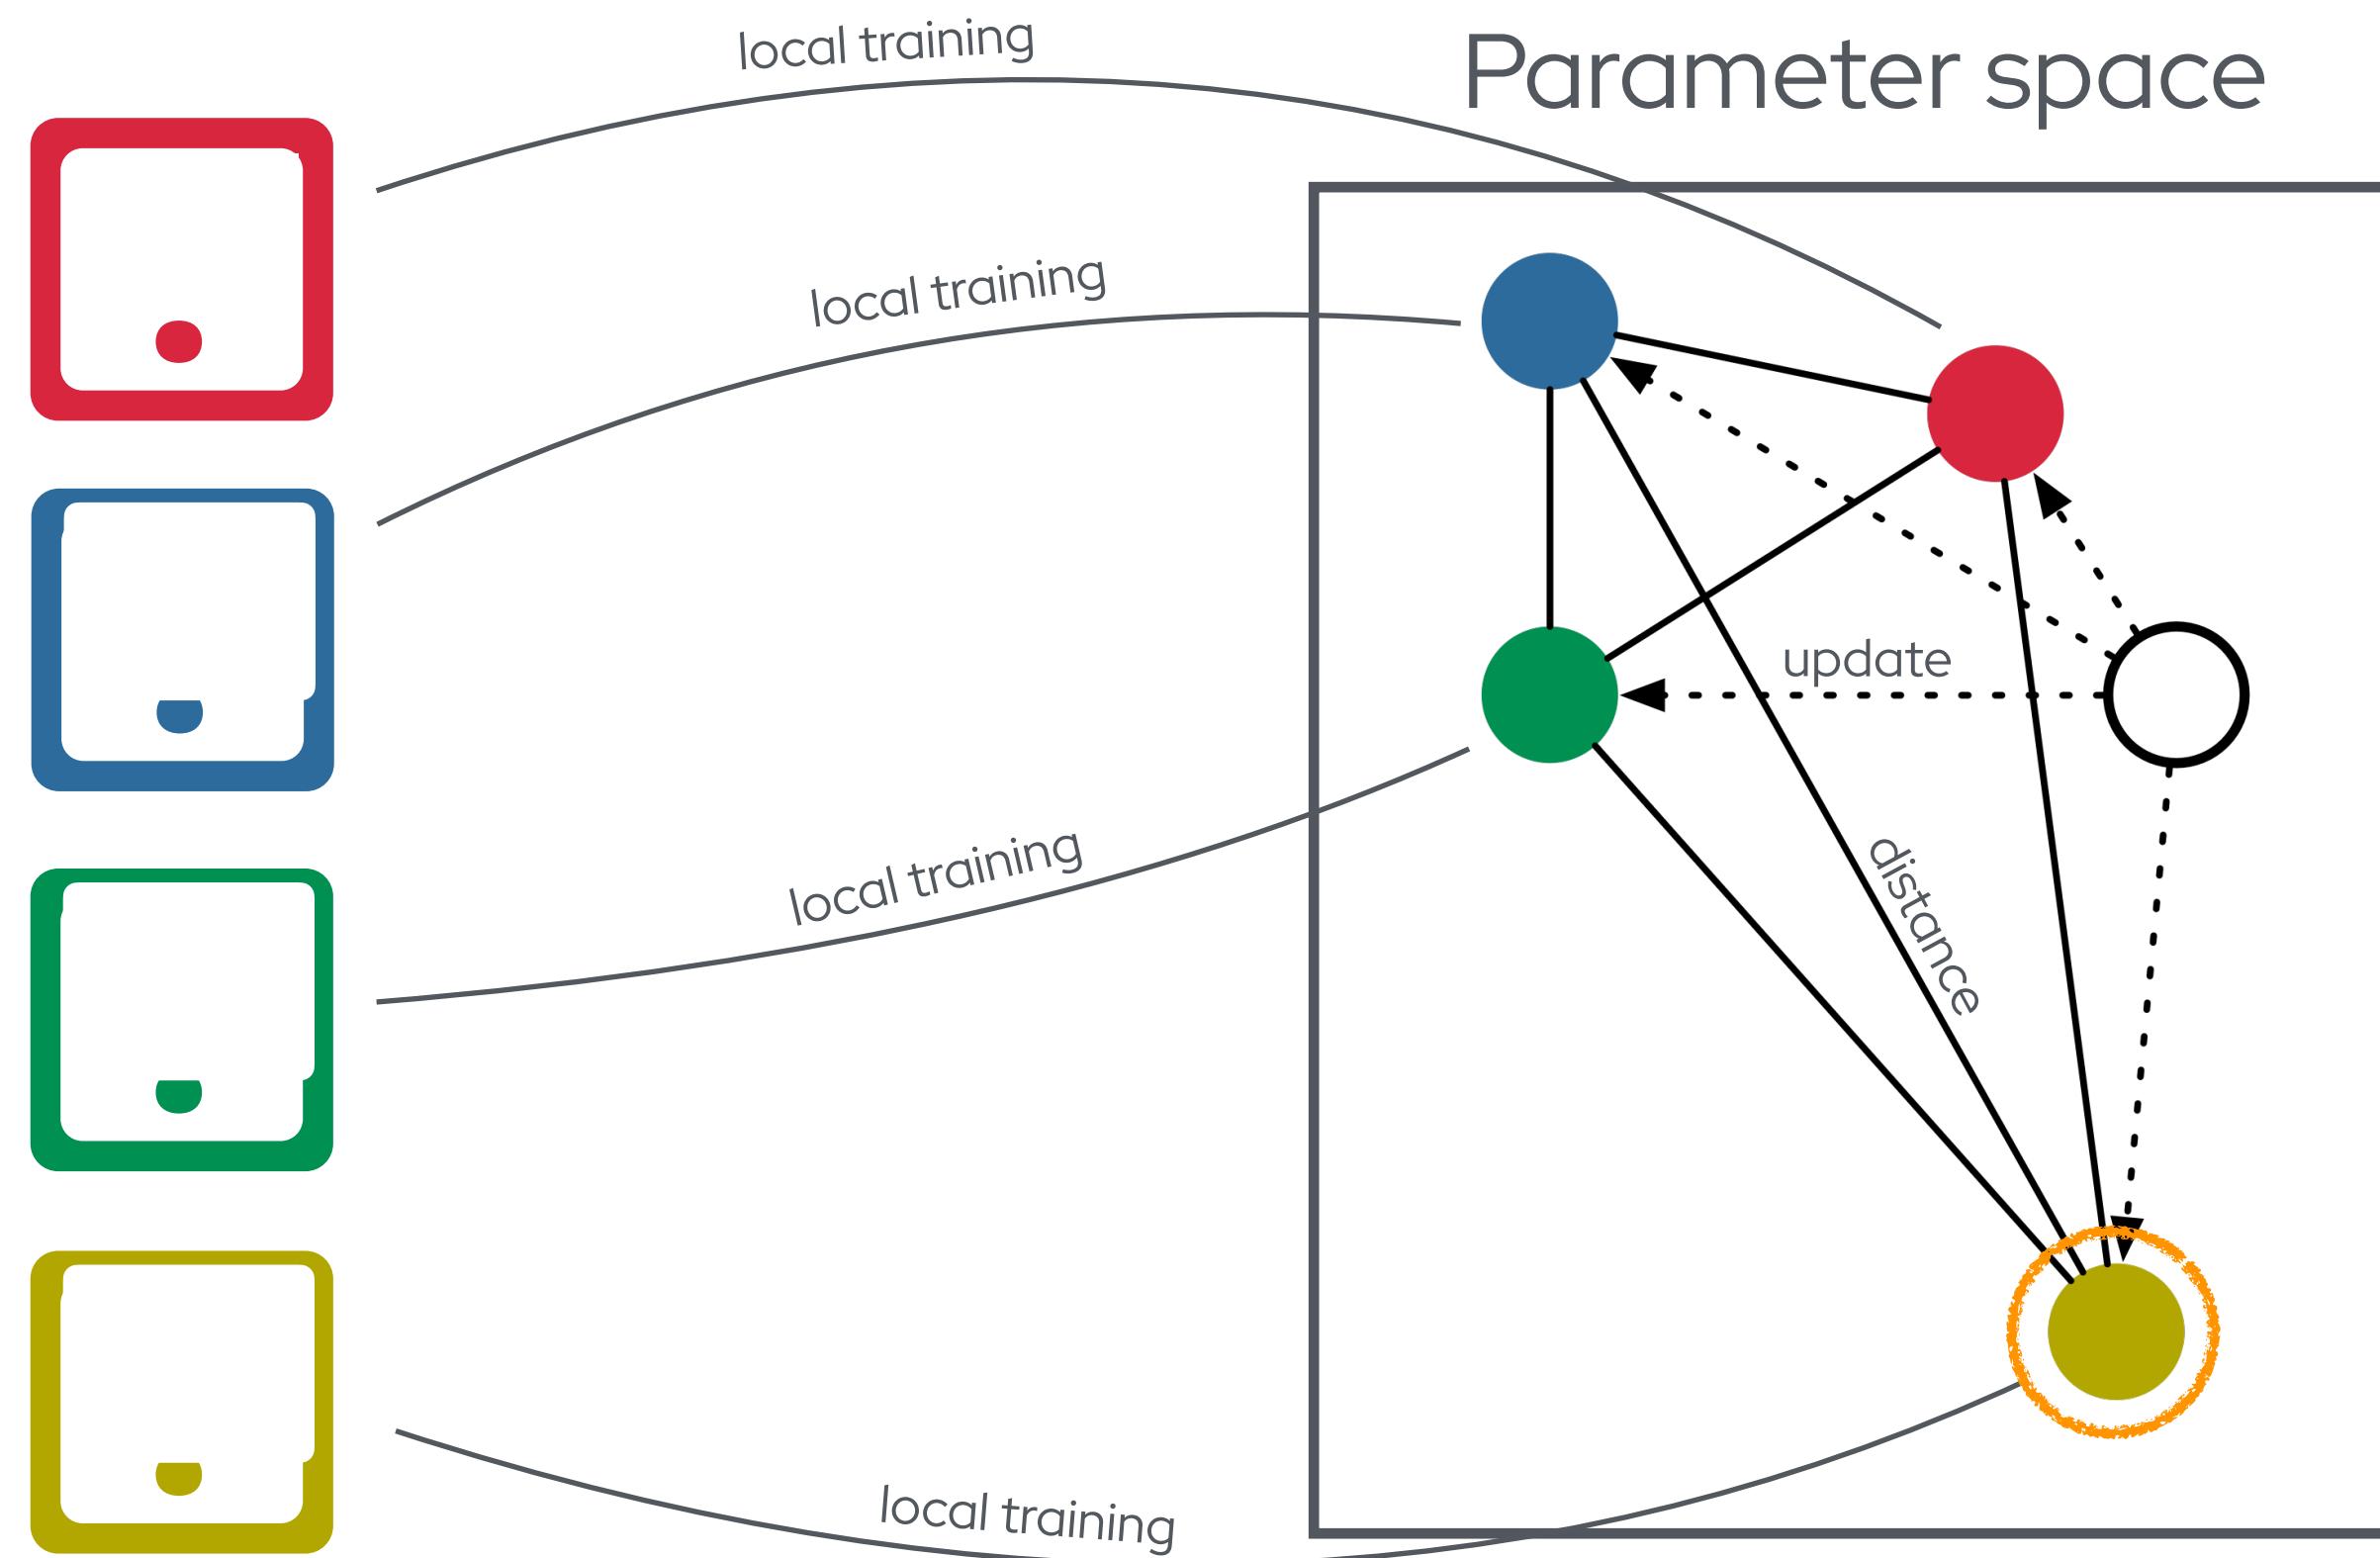
- **Geometric median**¹
- **Coordinate-wise median**²
- **Trimmed-mean**²

¹Pillutla, Krishna, Sham M. Kakade, and Zaid Harchaoui. "Robust aggregation for federated learning." IEEE Transactions on Signal Processing 70 (2022): 1142-1154.

²Yin, Dong, et al. "Byzantine-robust distributed learning: Towards optimal statistical rates." International conference on machine learning. Pmlr, 2018.

Federated Learning

Robust aggregation: **KRUM**

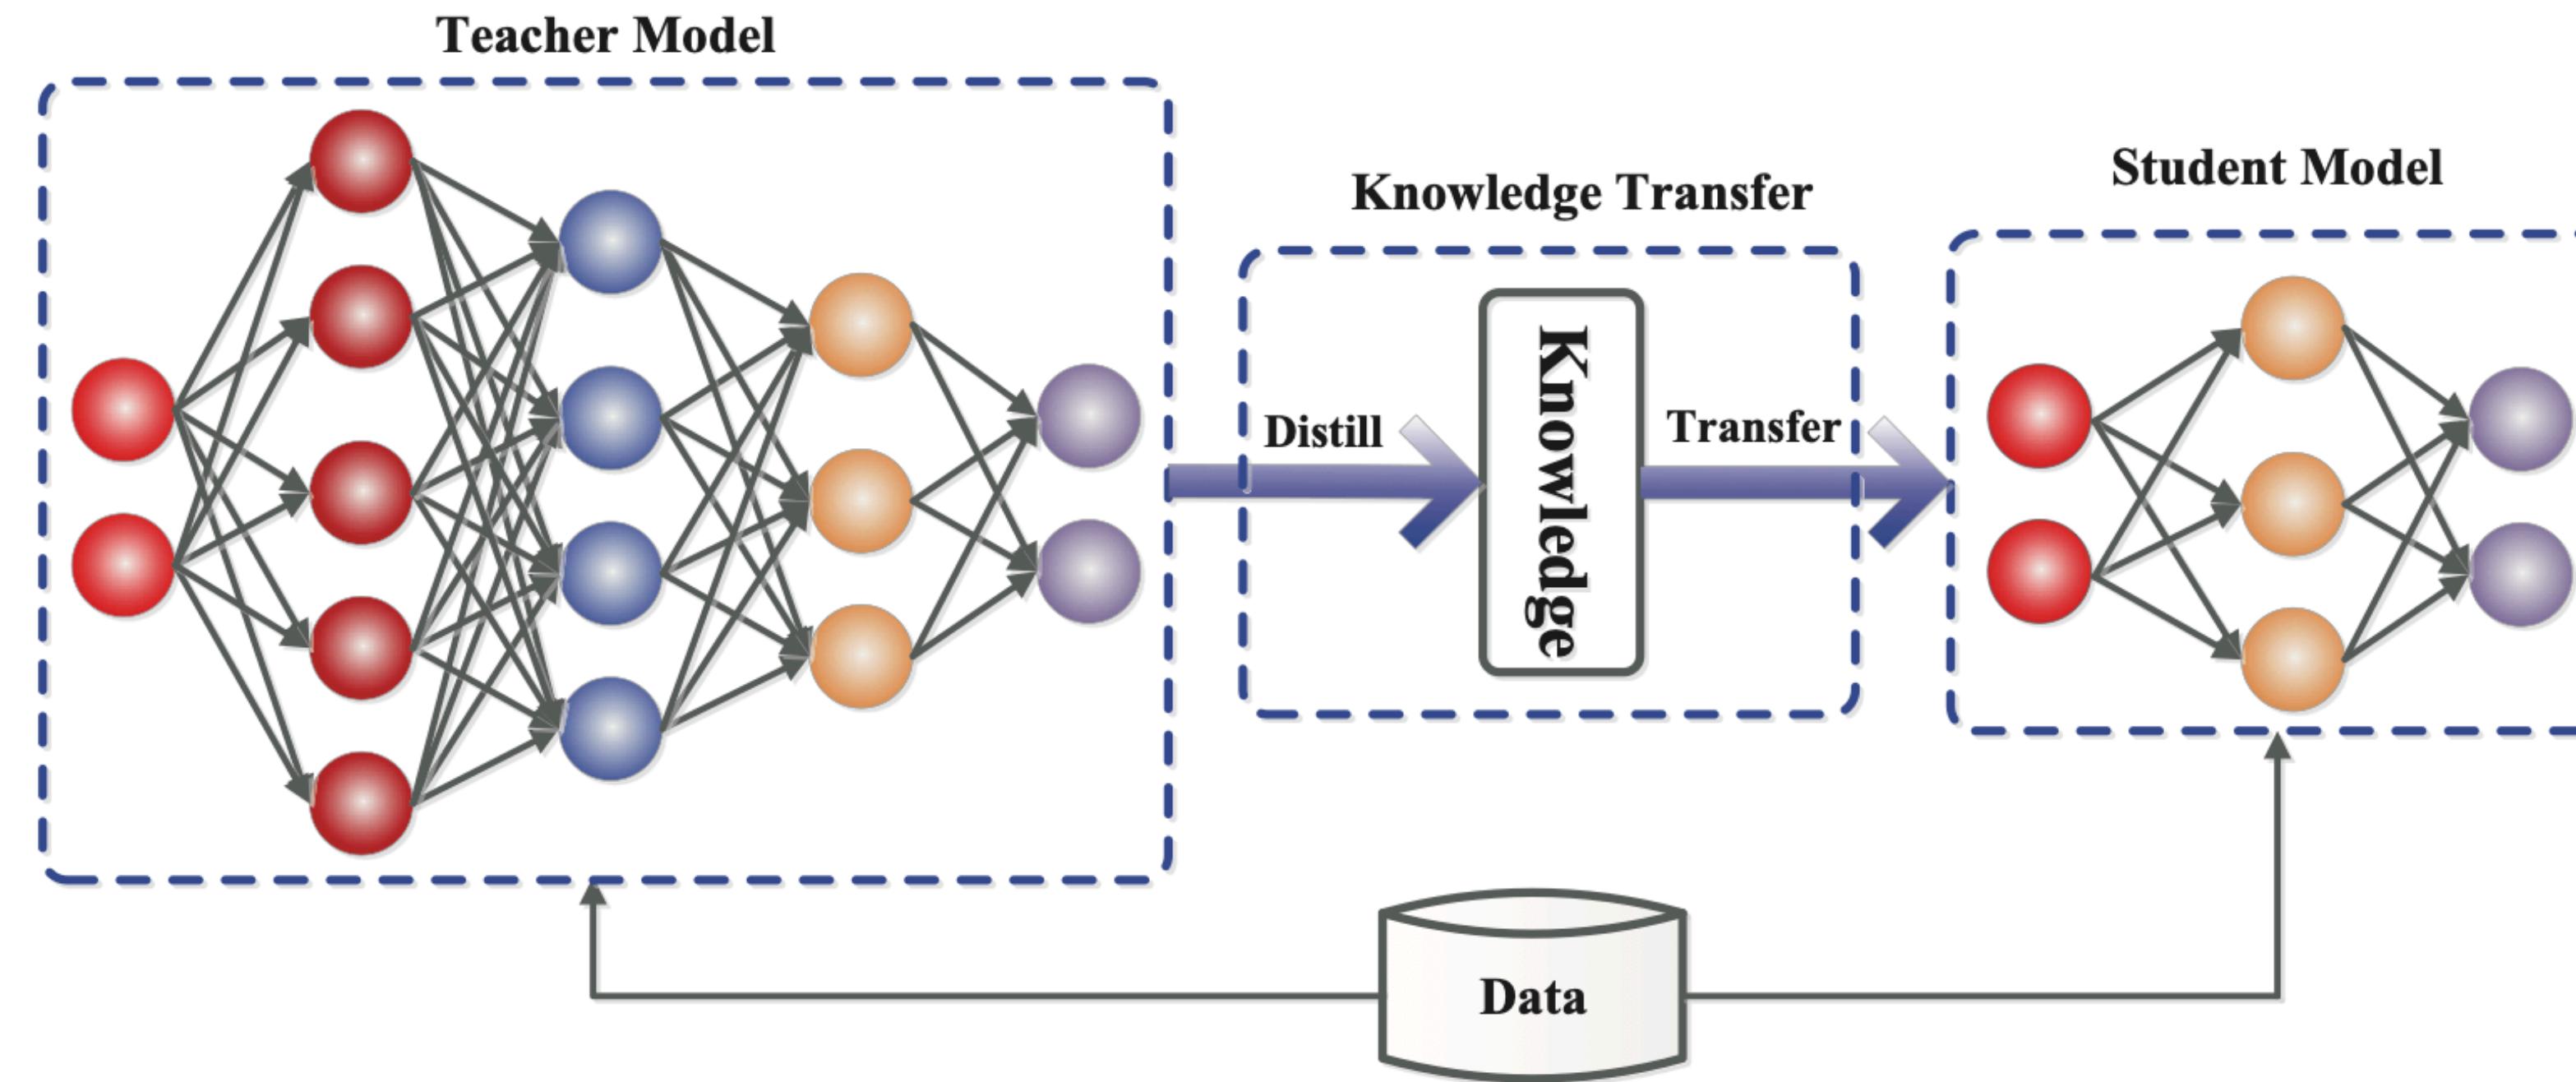


KRUM: Identifies the **outliers** as the models that are one average more distant than the others

Federated Learning

Communication

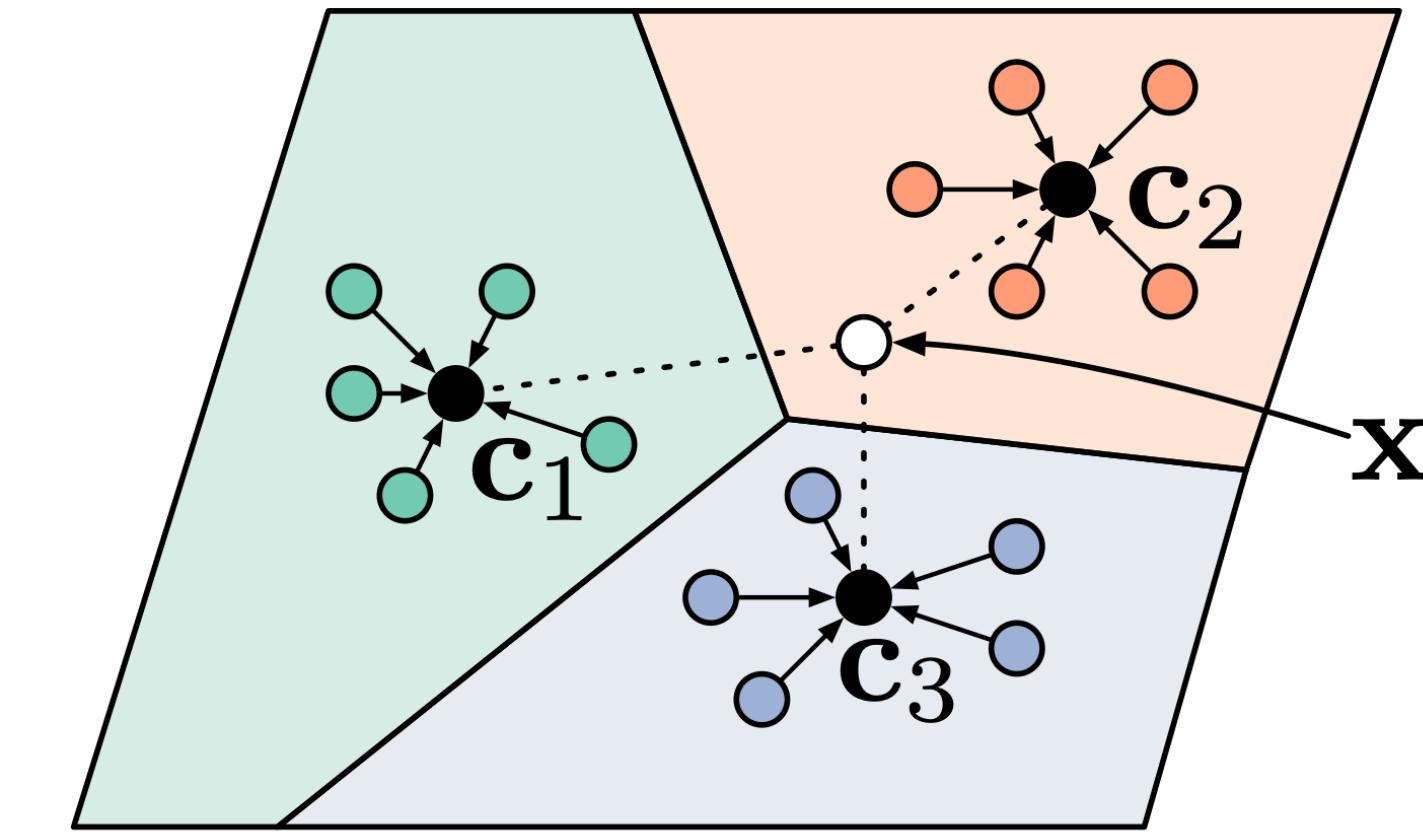
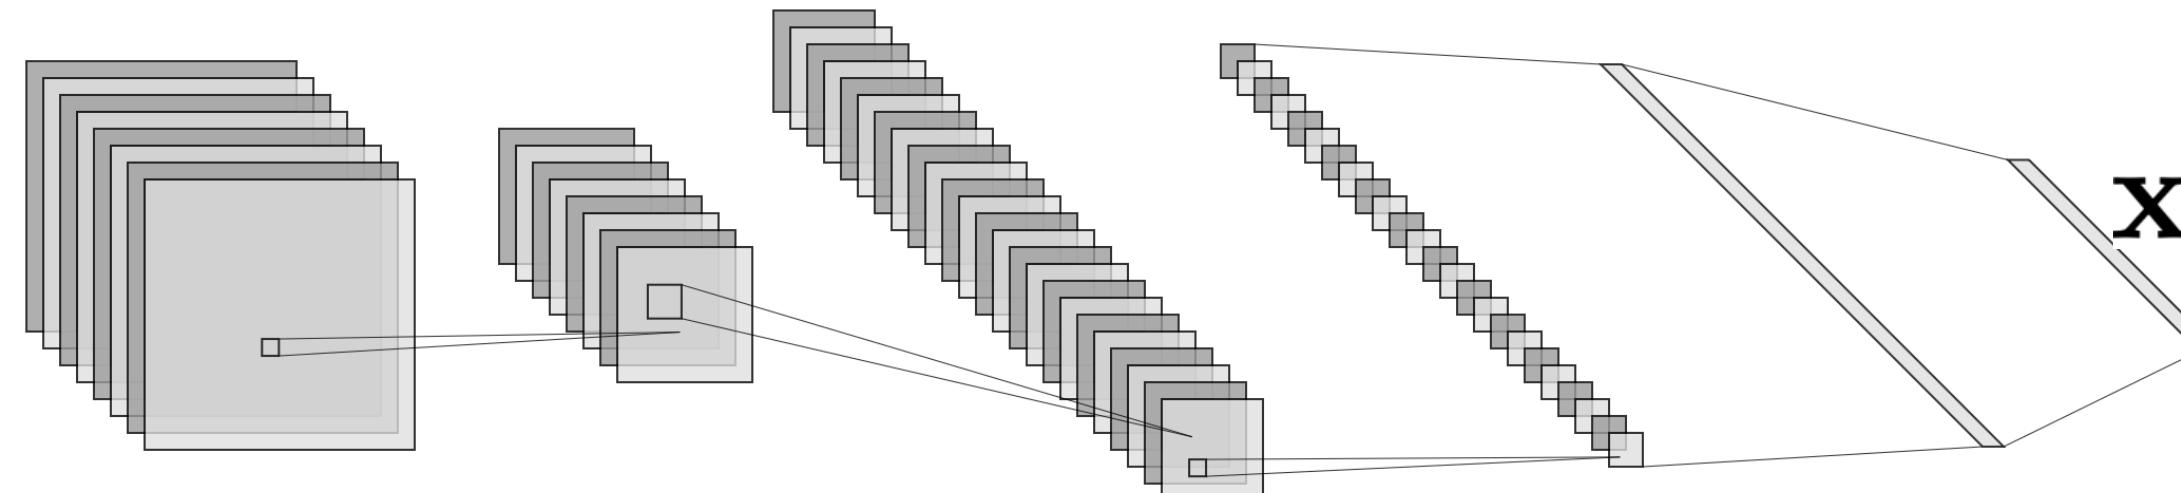
A simple technique to reduce the communication is to leverage **knowledge distillation**



Federated Learning - Communication

Communication - Prototype Learning

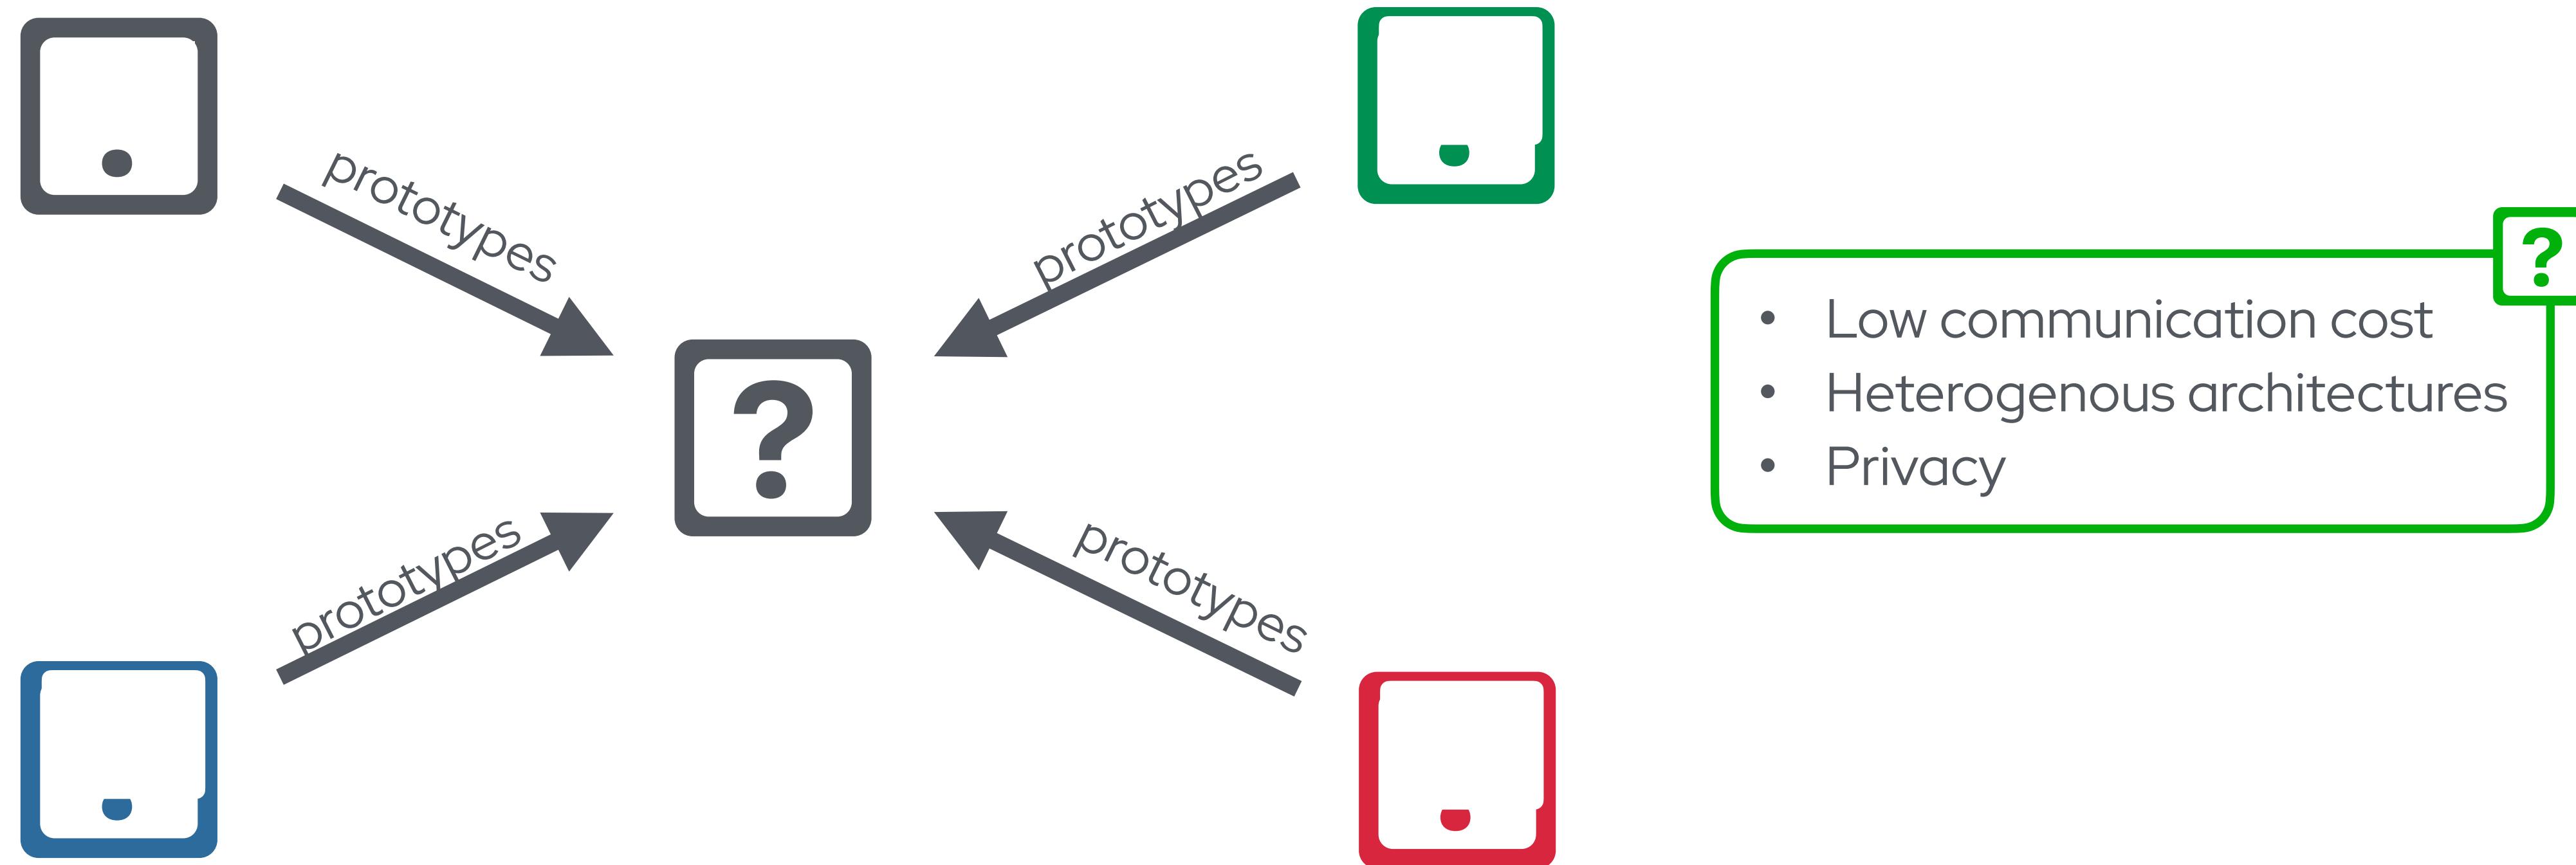
Prototype Learning is a well-known approach that relies heavily on the direct representation of the data



Federated Prototype Learning

Communication - Prototype Learning

The main idea is to **exchange only the prototypes** thus hugely reduce the communication cost



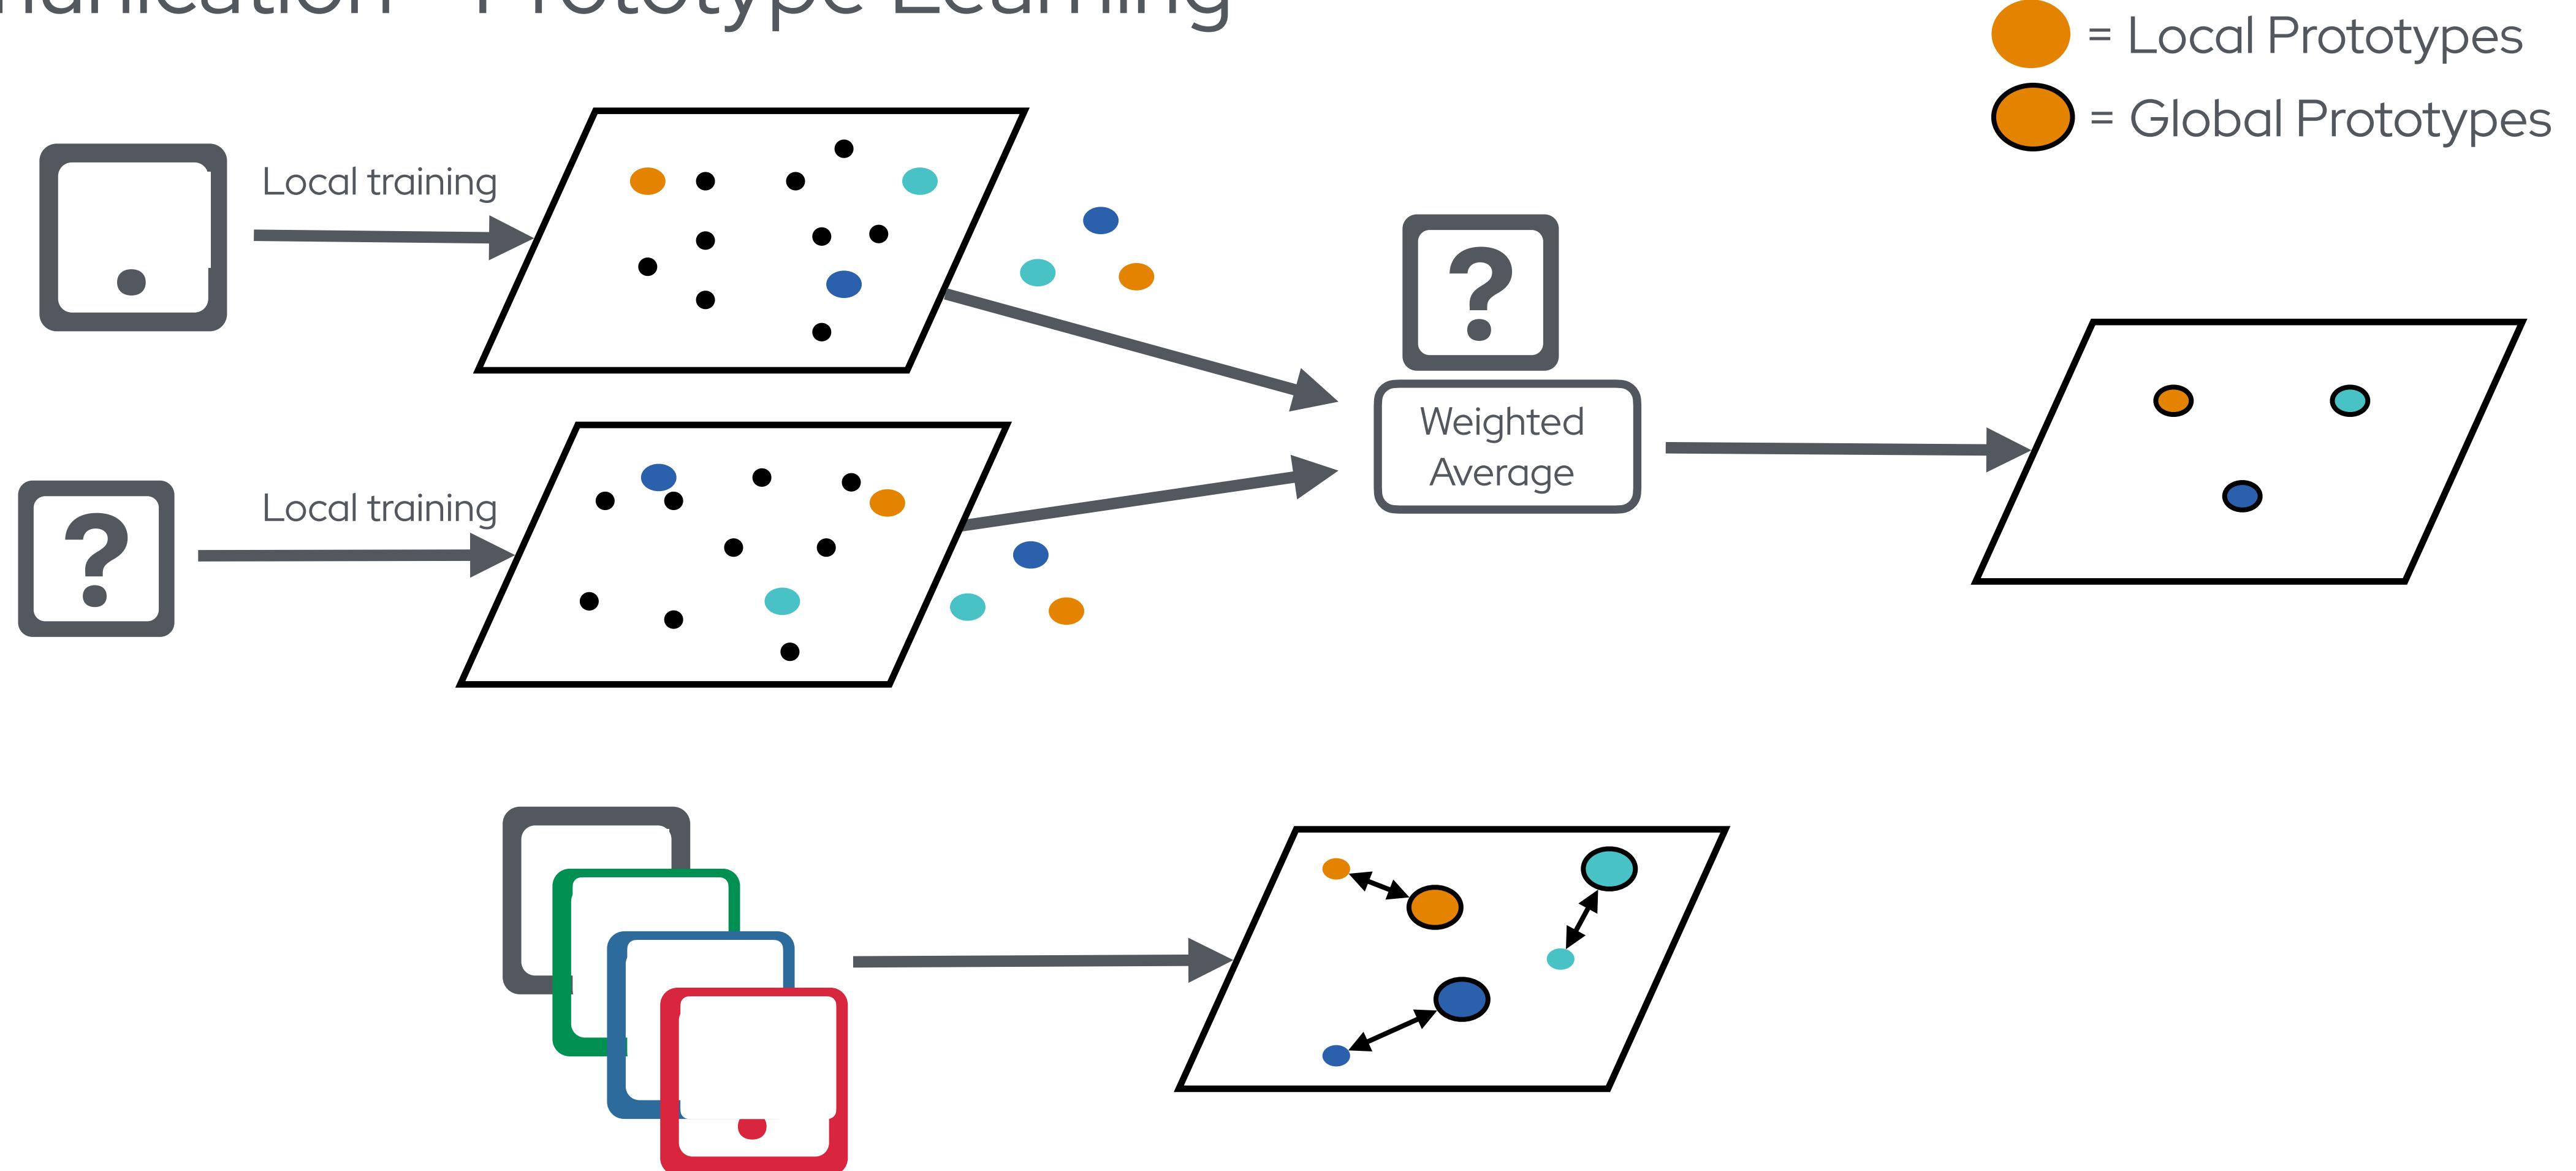
Federated Prototype Learning

The naive way



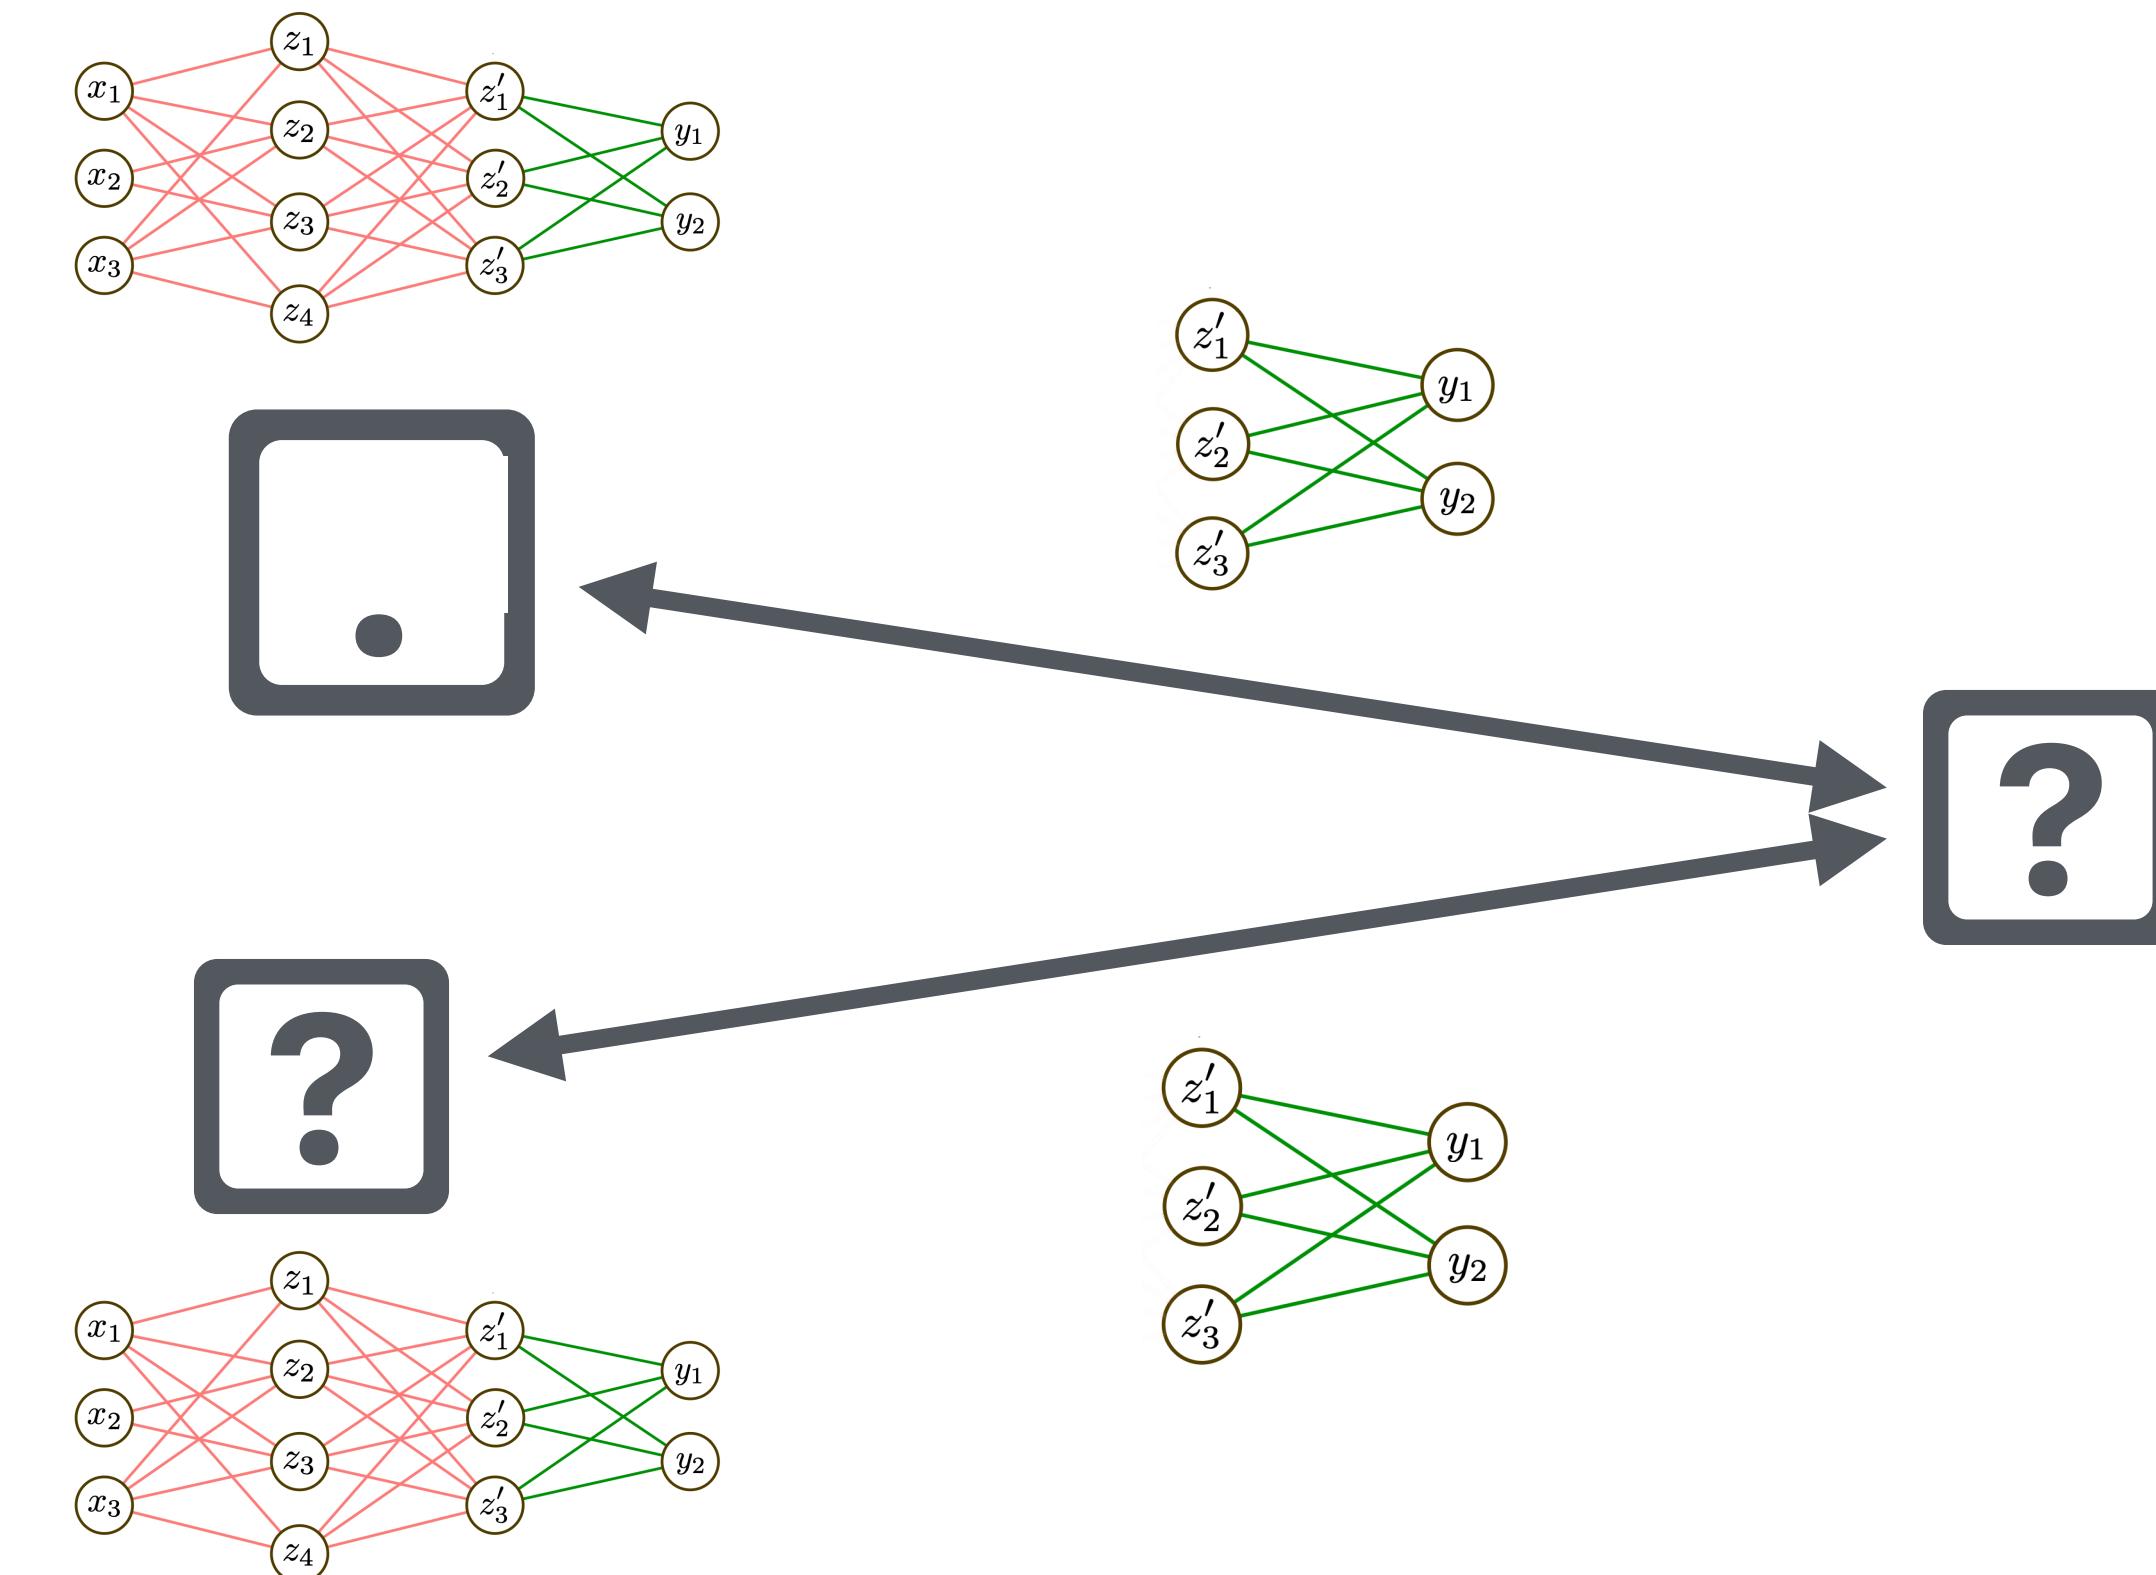
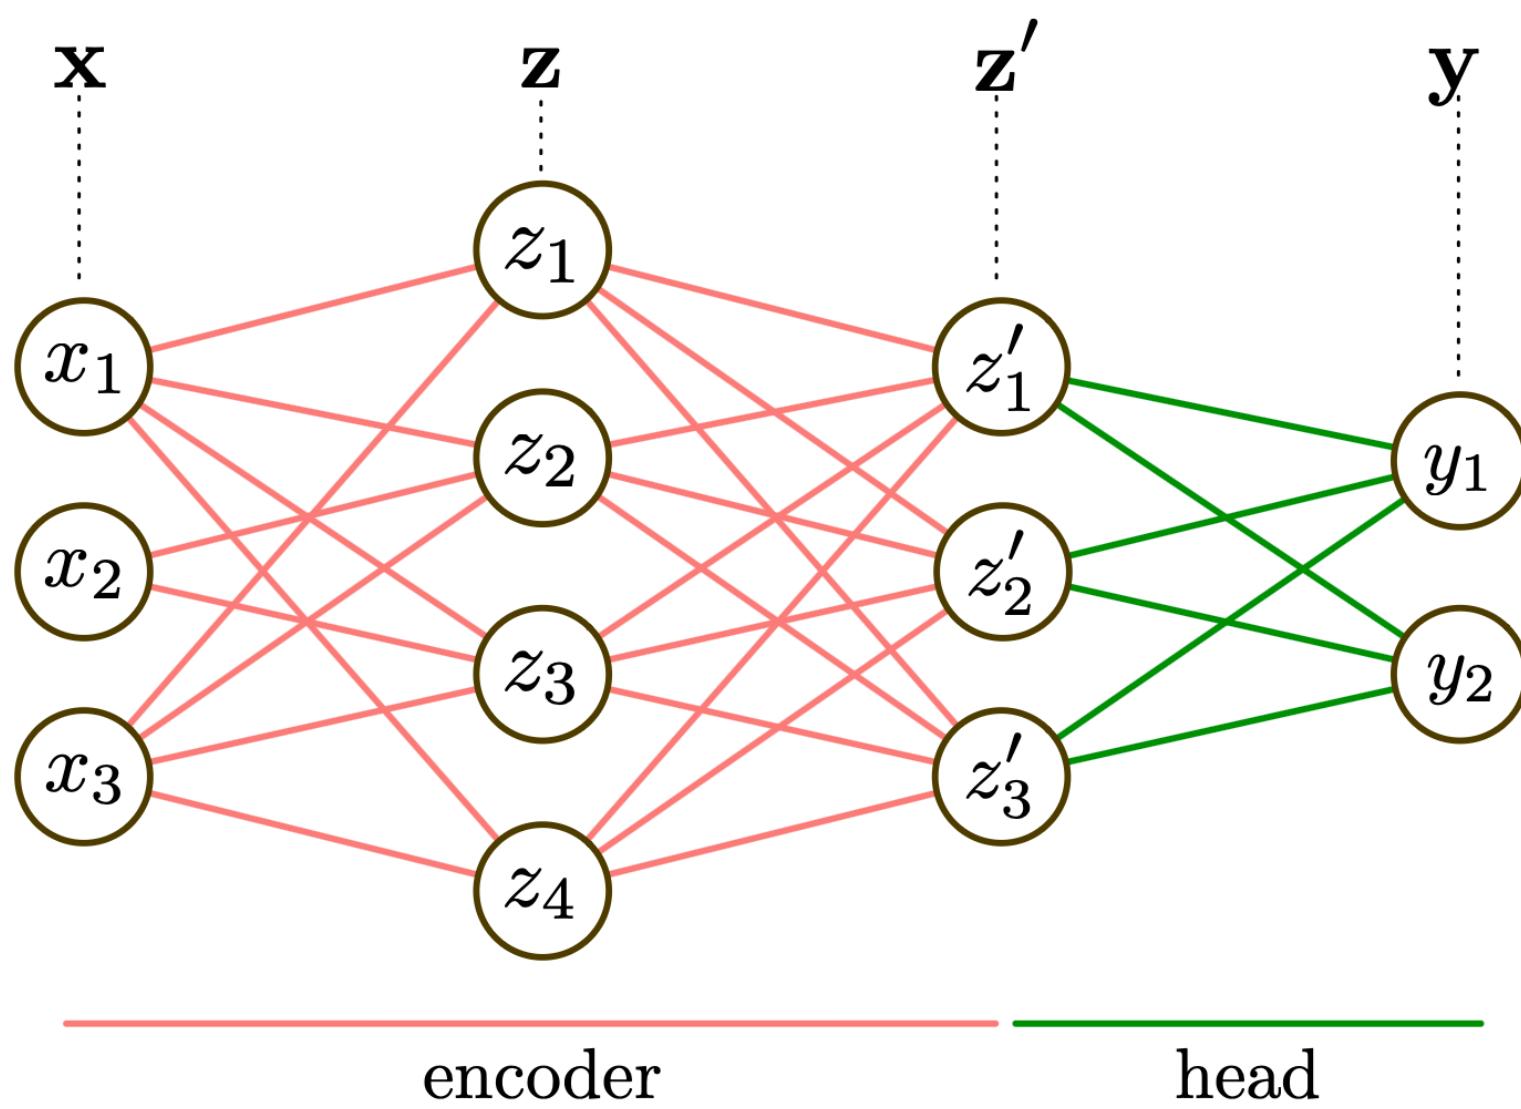
Federated Prototype Learning

Communication - Prototype Learning



LG-FedAVG

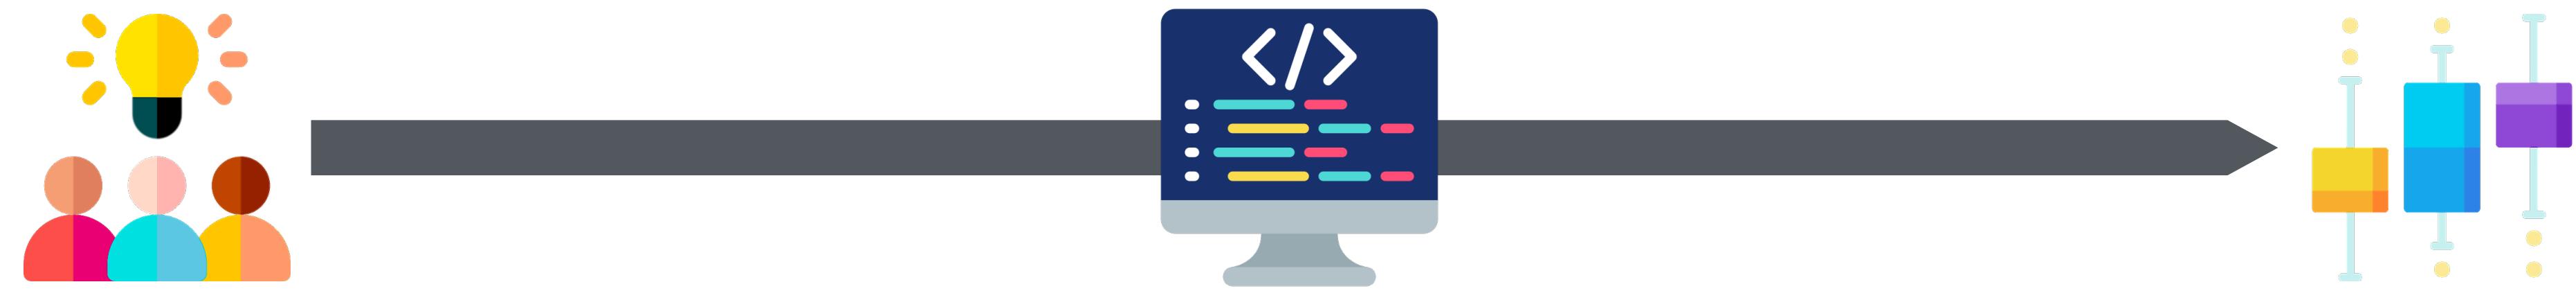
Partial model sharing: shared head and private backbone



Introduction to

Yet another FL framework?

Not really!



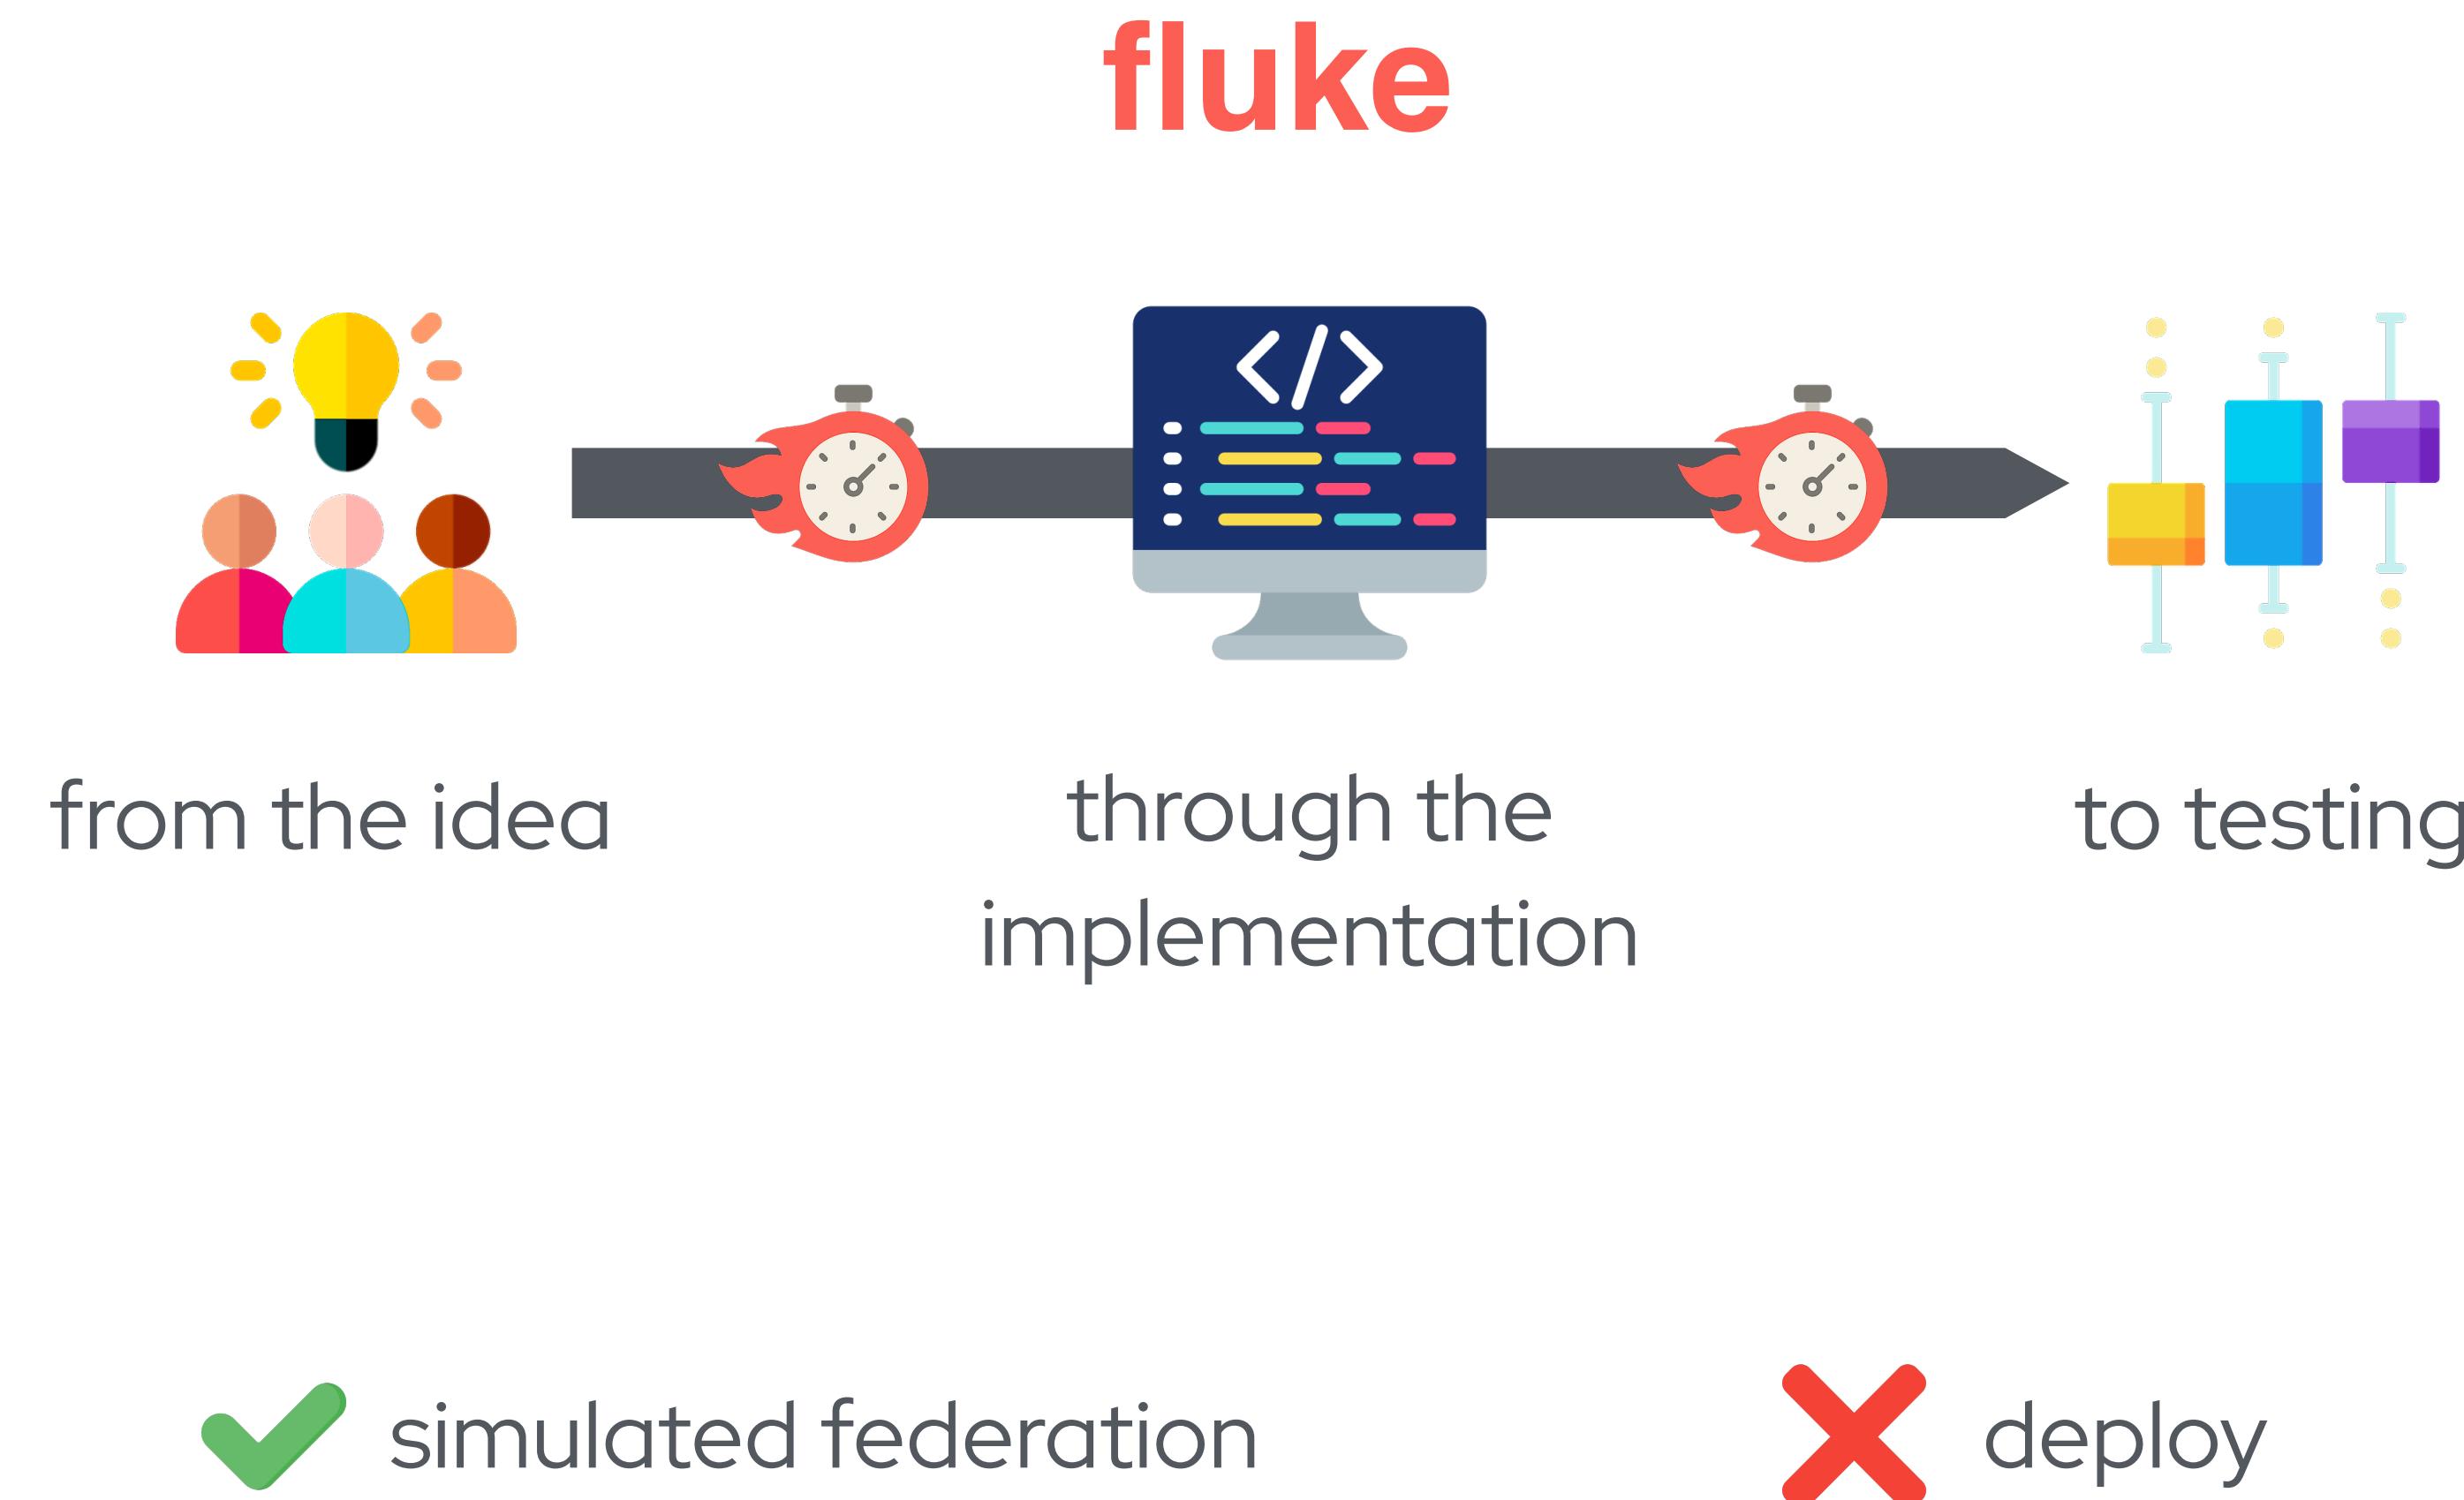
from the idea

through the
implementation

to testing

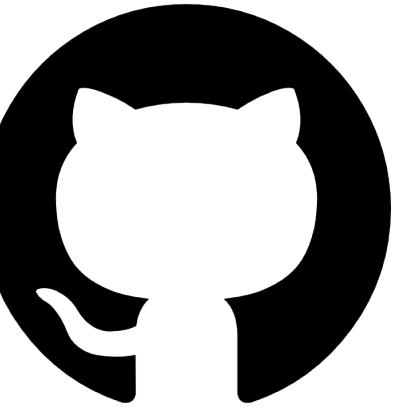
Yet another FL framework?

Not really!



Main features of fluke

Designed for fast prototyping and testing



<https://github.com/makgyver/fluke>

- **Open source:** fluke is an open-source Python package;
- **Easy to use:** fluke is designed to be extremely easy to use out of the box;
- **Easy to extend:** fluke is designed to minimize the overhead of adding new algorithms;
- **Up-to-date:** fluke comes with several (30+) state-of-the-art federated learning algorithms and datasets and it is regularly updated to include the latest affirmed techniques;
- **Easy to read:** the source code of the algorithms is written to mimic as close as possible the description in the reference papers.



fluke is on PyPi!

Install it using a single command



```
$ pip install fluke-fl
```

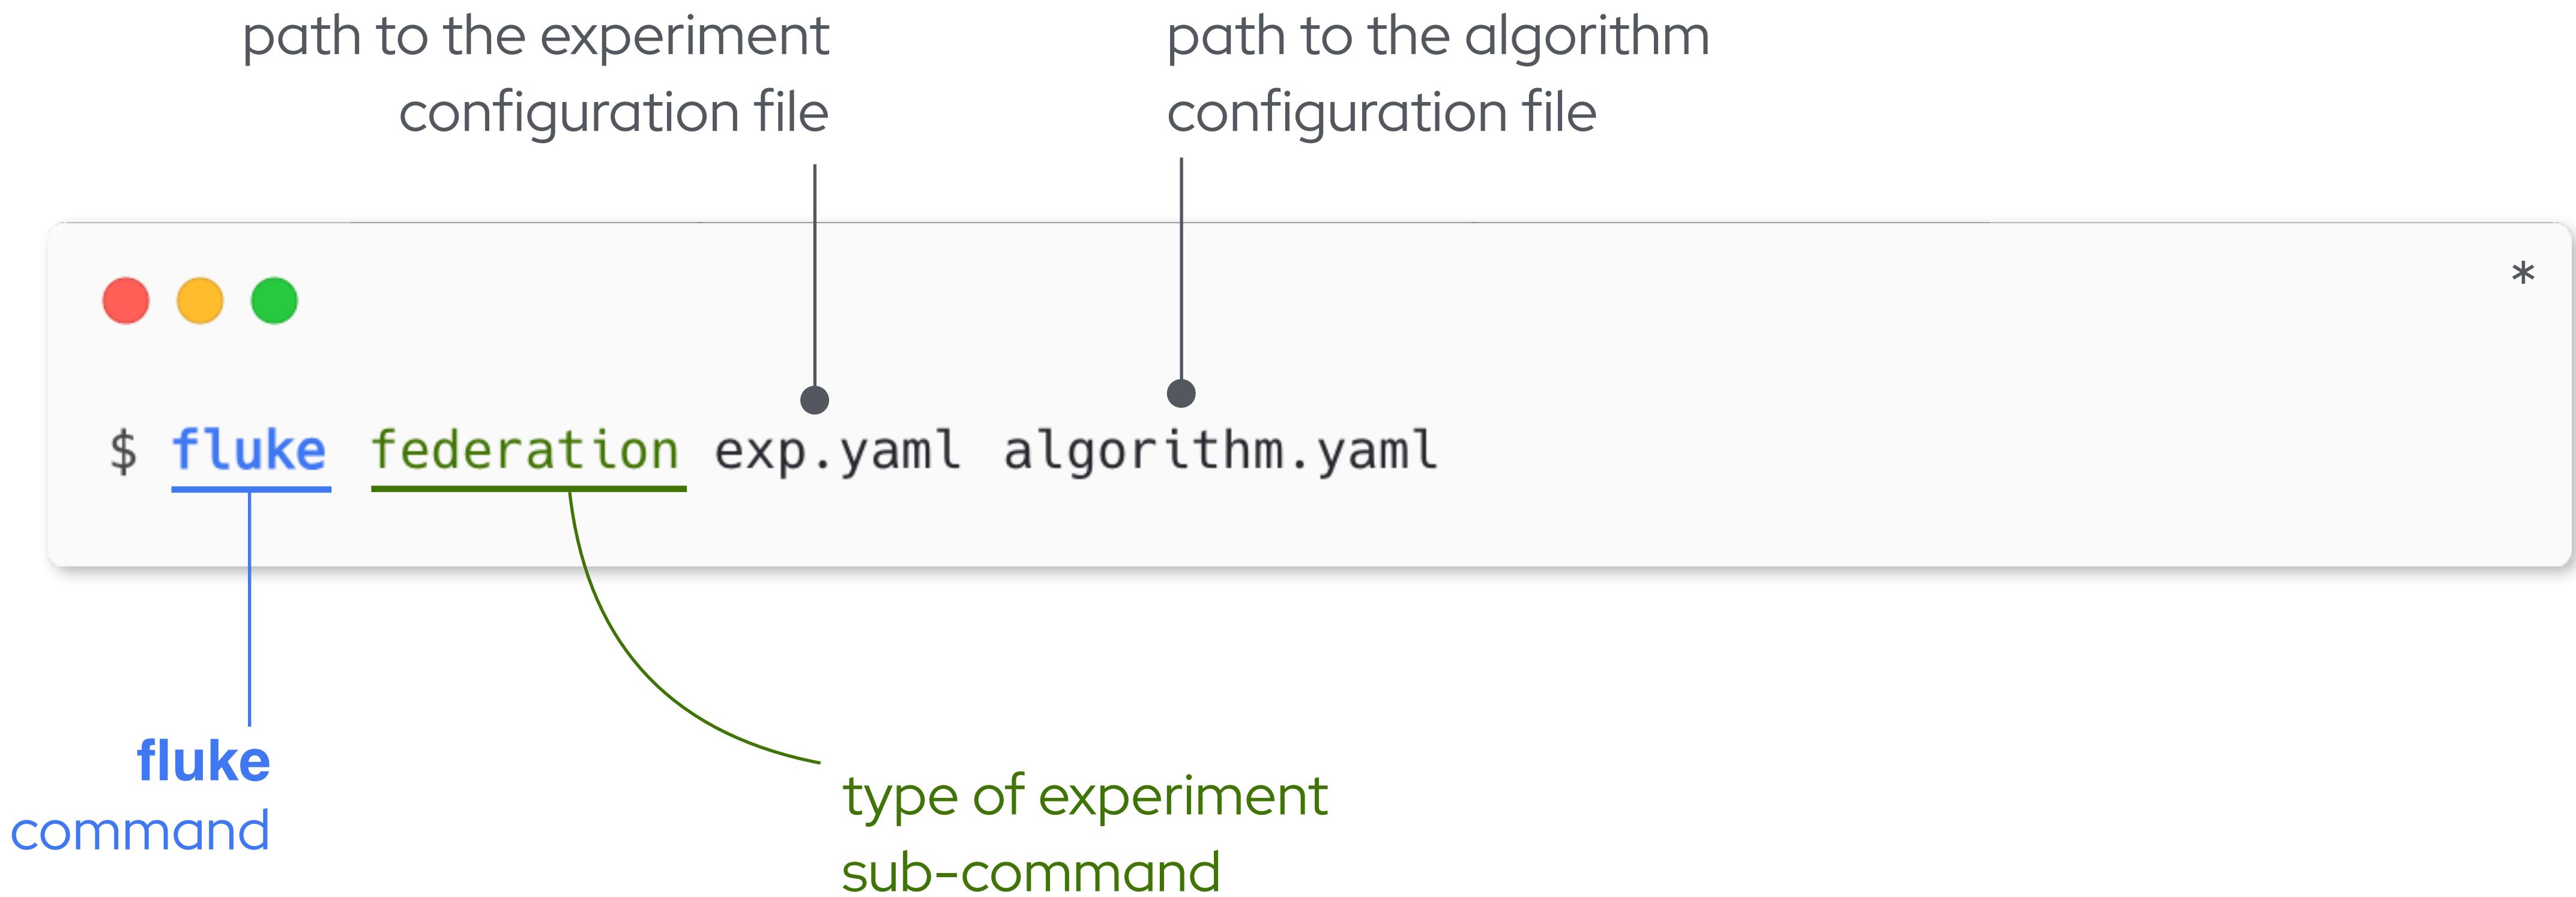
... or by cloning the repo



```
$ git clone https://github.com/makgyver/fluke.git  
$ cd fluke  
$ pip install -r requirements.txt
```

fluke CLI

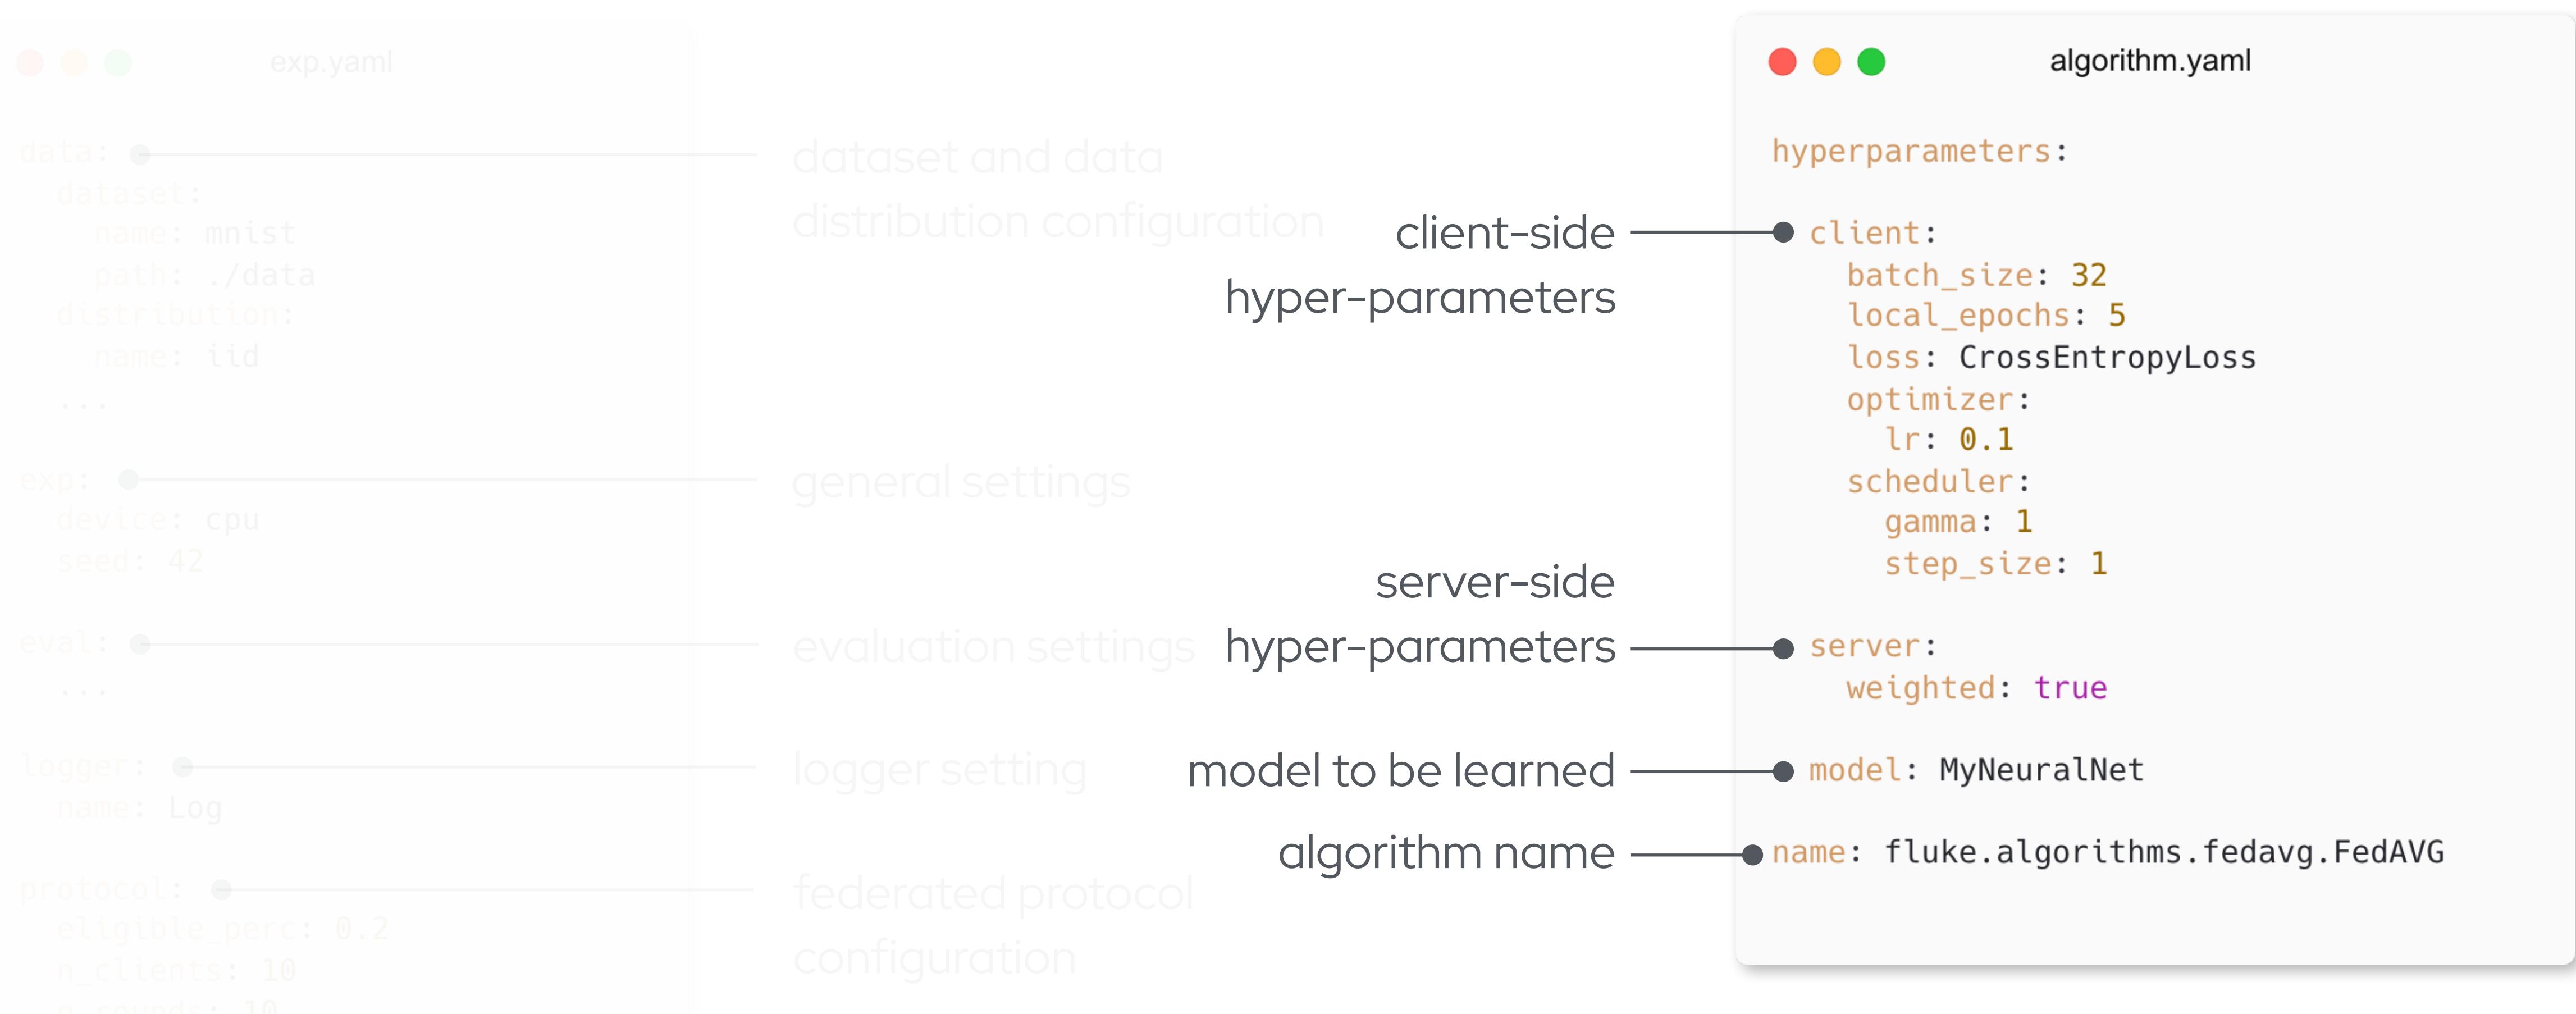
You can run your experiment outright with a single command!



(*) If you cloned the repo, the command (launched from the fluke folder) is `$ python -m fluke.run federation exp.yaml algorithm.yaml`

Configuration files

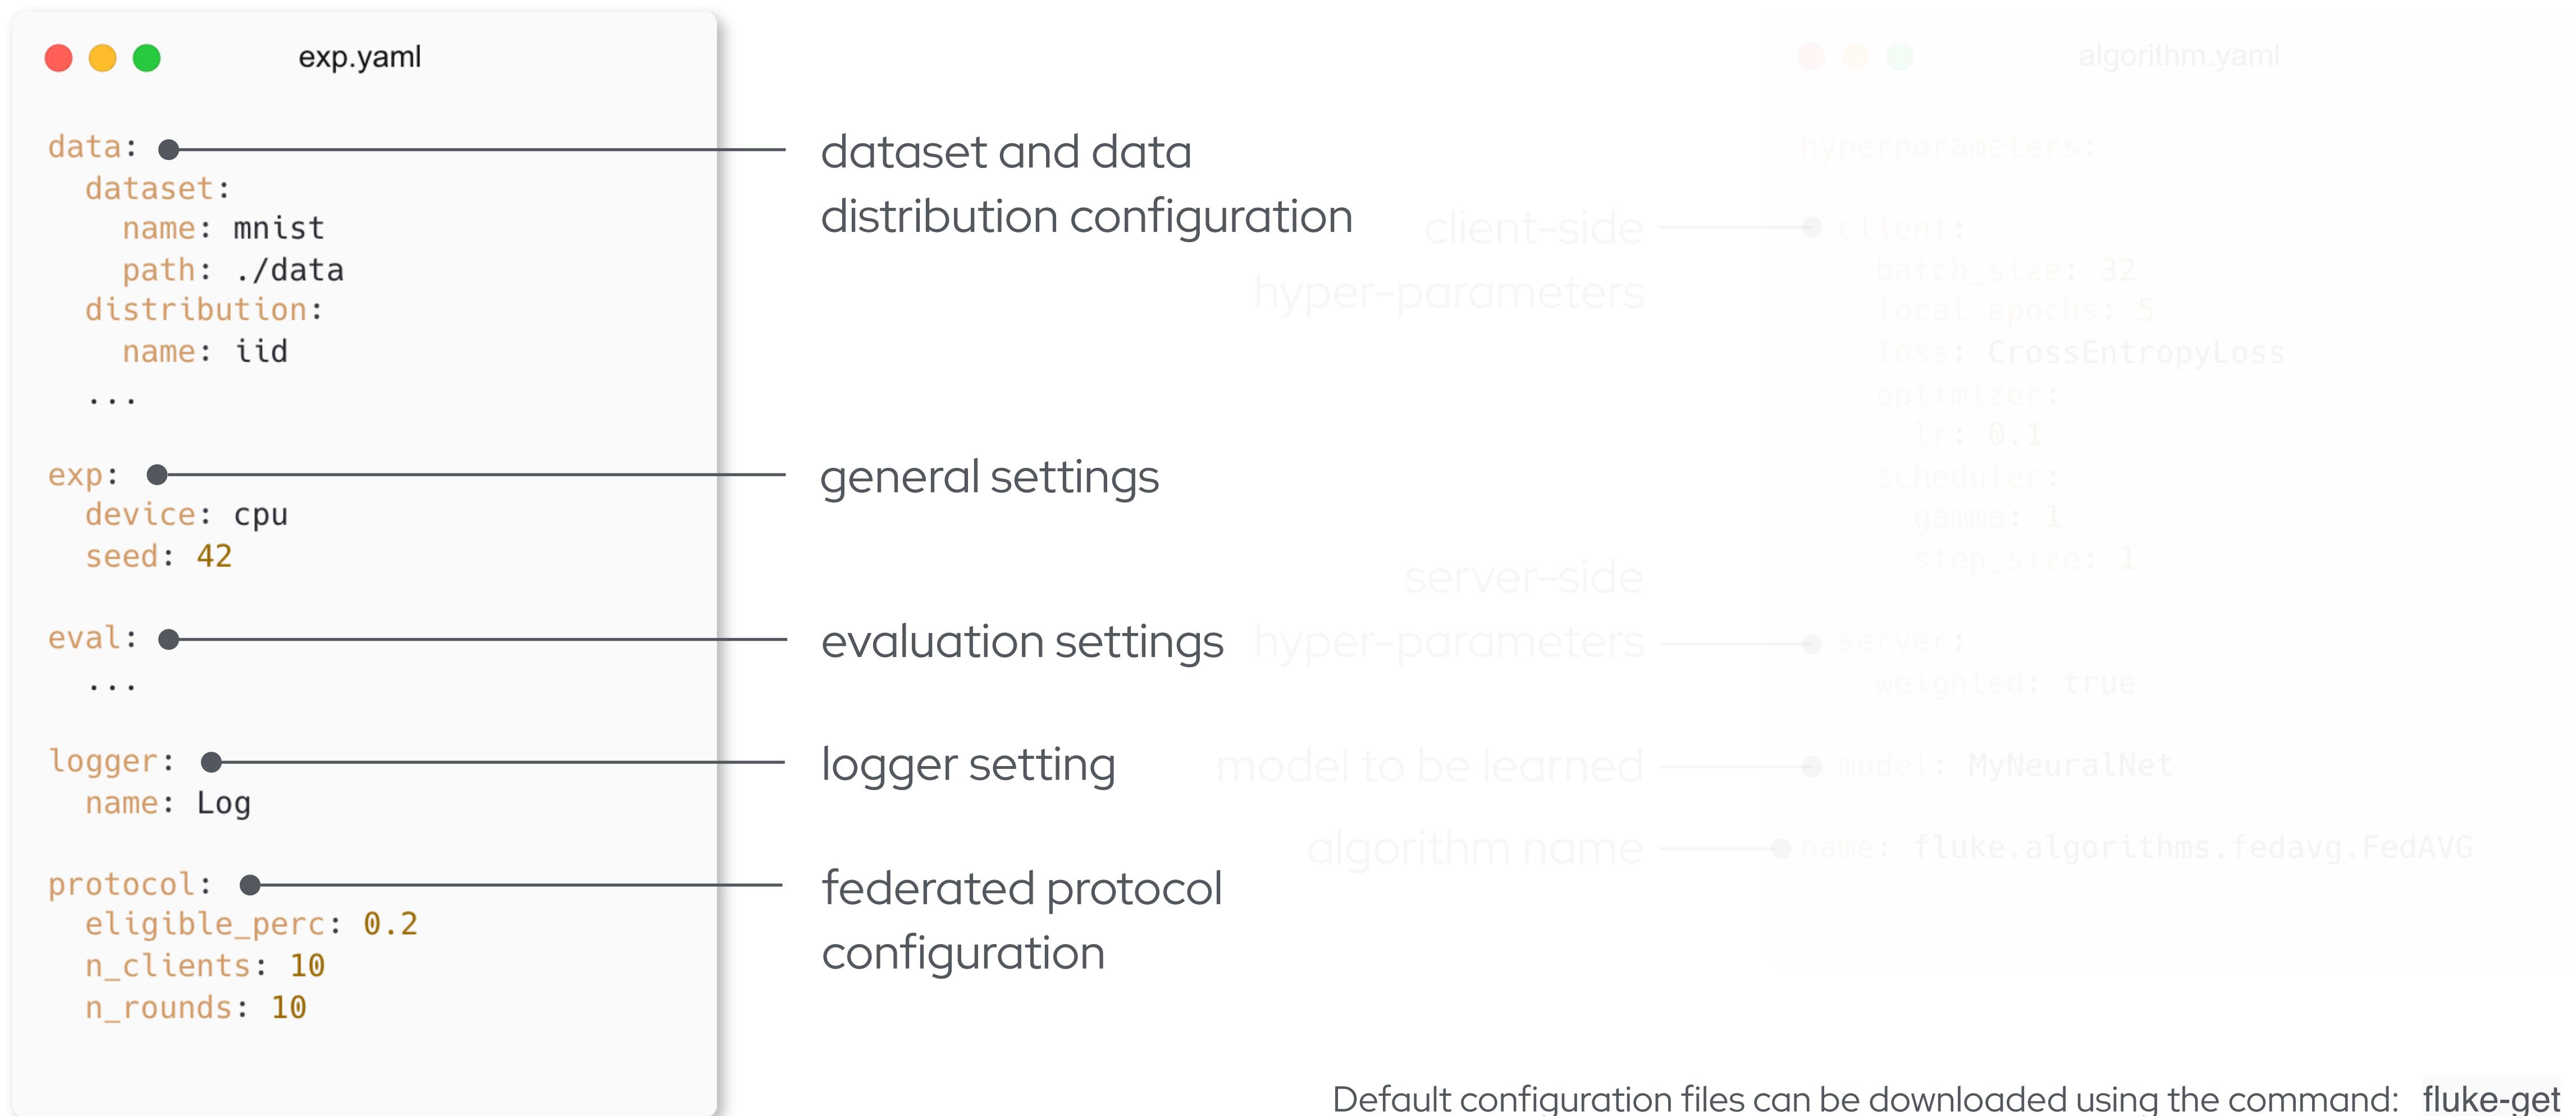
Two YAML files for everything you want to configure



Default configuration files can be downloaded using the command: `fluke-get`

Configuration files

Two YAML files for everything you want to configure



Default configuration files can be downloaded using the command: `fluke-get`

fluke CLI – not only federation

You can run the “same” experiment without the federation for comparison

same experiment but without the federation –
the number of epochs client-side are calculated*
as $n_rounds * eligible_perc * local_epochs$



```
$ fluke clients-only exp.yaml algorithm.yaml
```



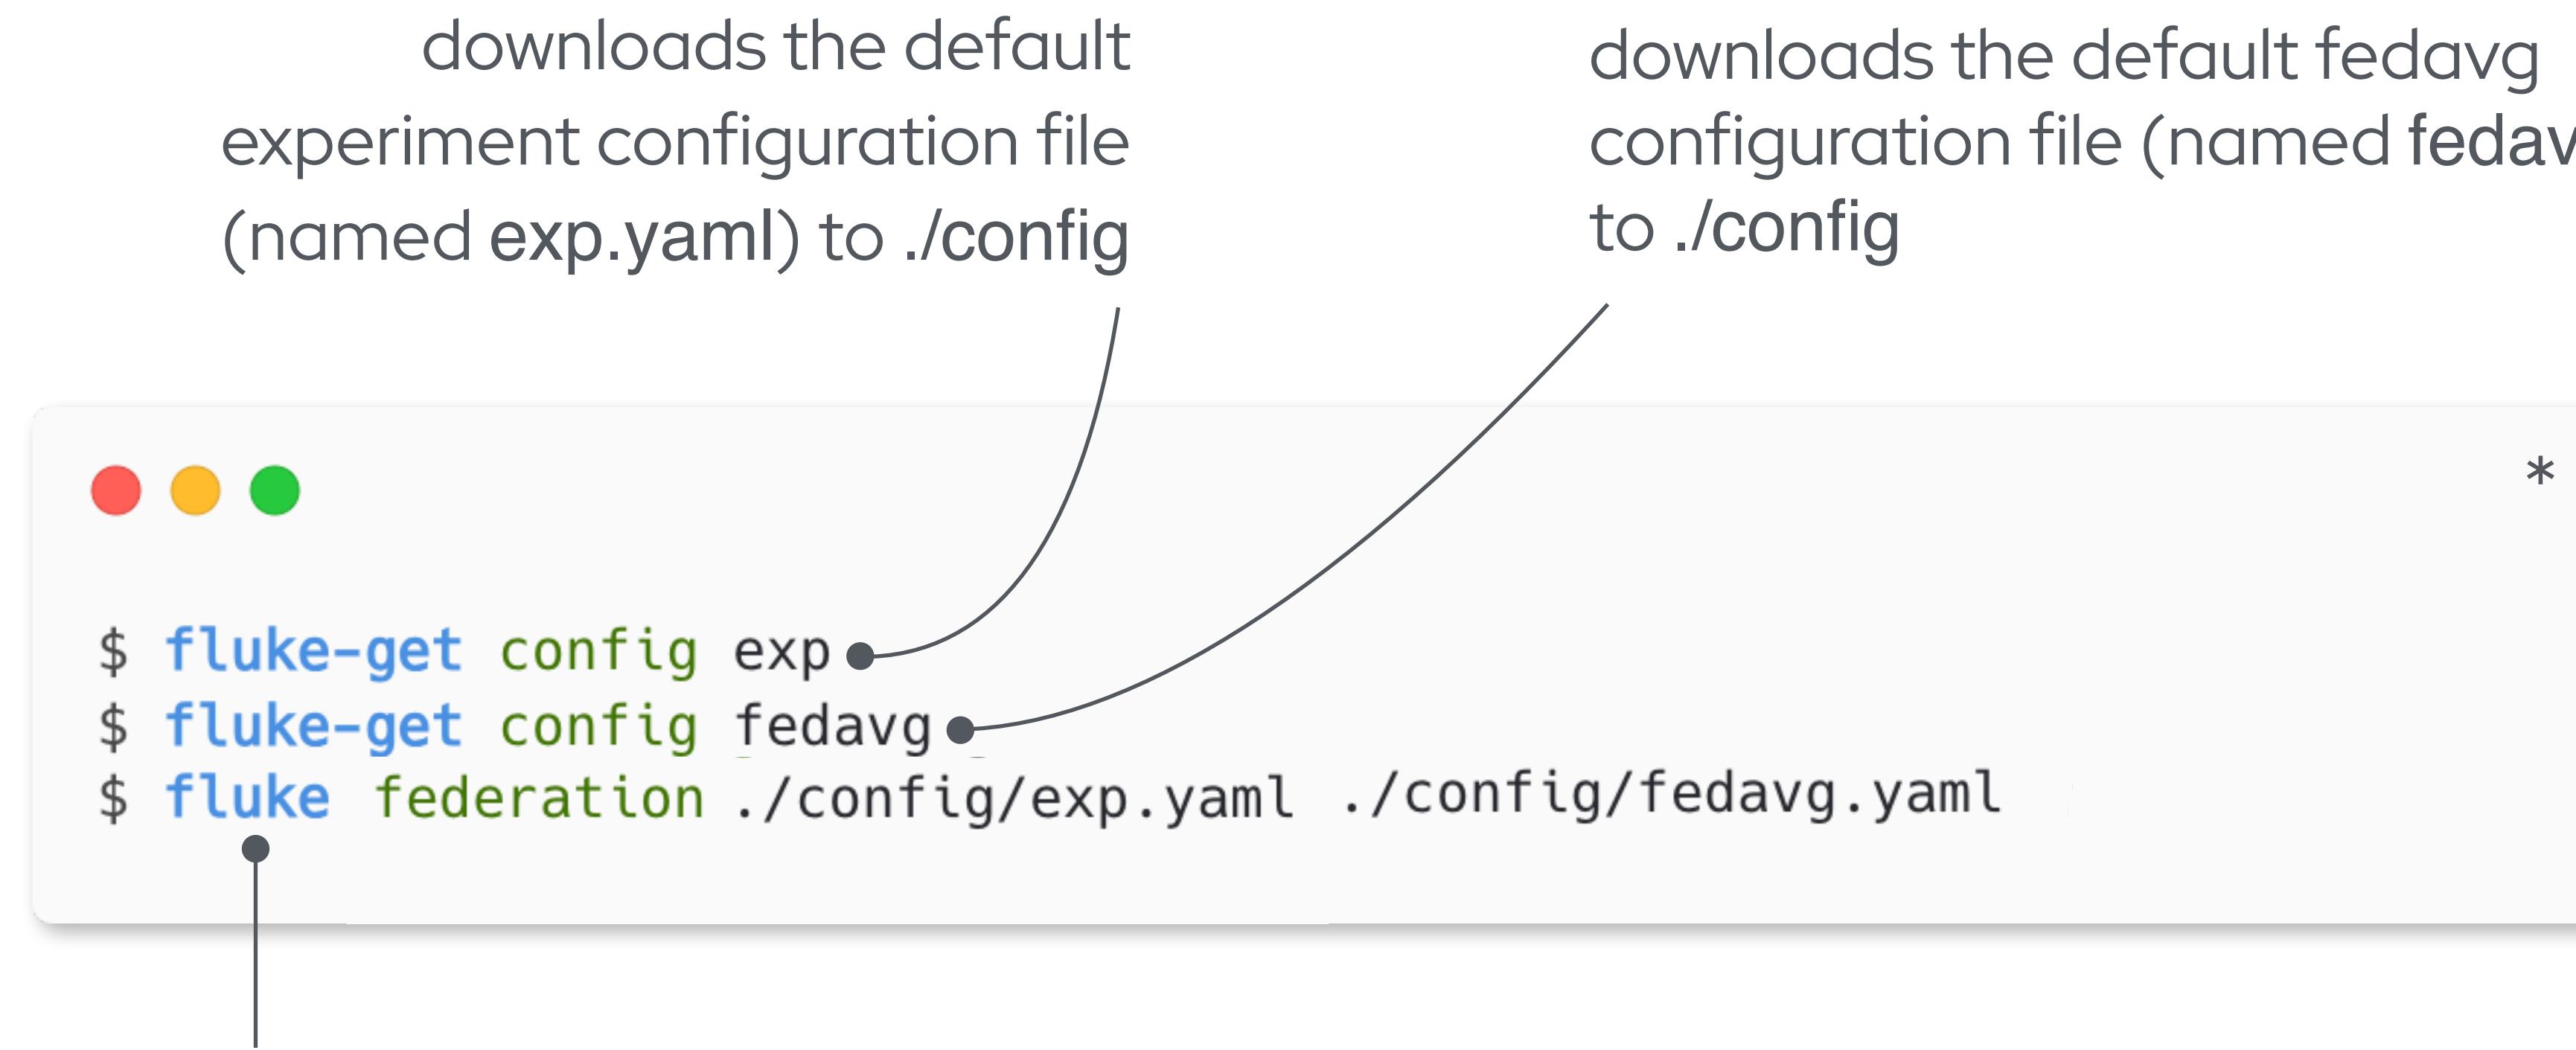
```
$ fluke centralized exp.yaml algorithm.yaml
```

same type of experiment but with all the dataset centralised

(*) The number of epochs can be set by the user via an option of the command, e.g., `--epochs=100`

Example: FedAVG on MNIST

Fluke comes with many downloadable configuration files ready to be used/modified



downloads the default experiment configuration file (named `exp.yaml`) to `./config`

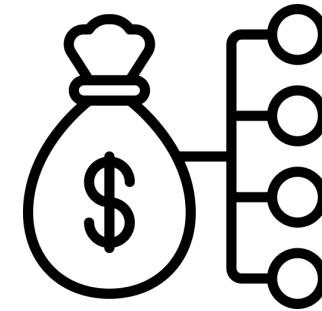
downloads the default `fedavg` configuration file (named `fedavg.yaml`) to `./config`

runs the federated algorithm specified in `fedavg.yaml` on the dataset specified in `exp.yaml`

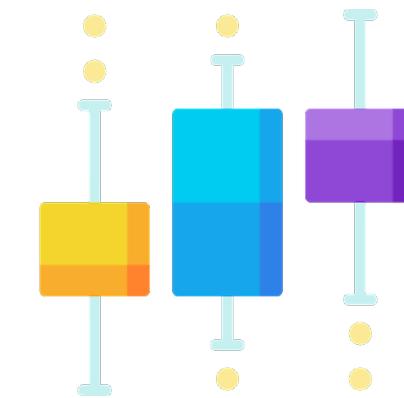
(*) If you want to know all the available default configuration files use the command `$ fluke-get list`

fluke logging

Performance can be logged on your preferred tool



communication cost*



classification
performance**
(e.g., accuracy, precision, recall, F1
- marco and micro)



system performance***



(*) The communication cost is estimated as the number fo floating points numbers exchanged by the entire federation.

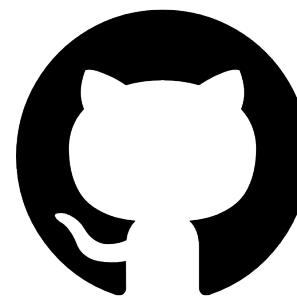
(**) Currently, fluke supports only classification but it is straightforward to extend to other tasks.

(***) System performance are automatically logged by W&B and ClearML.

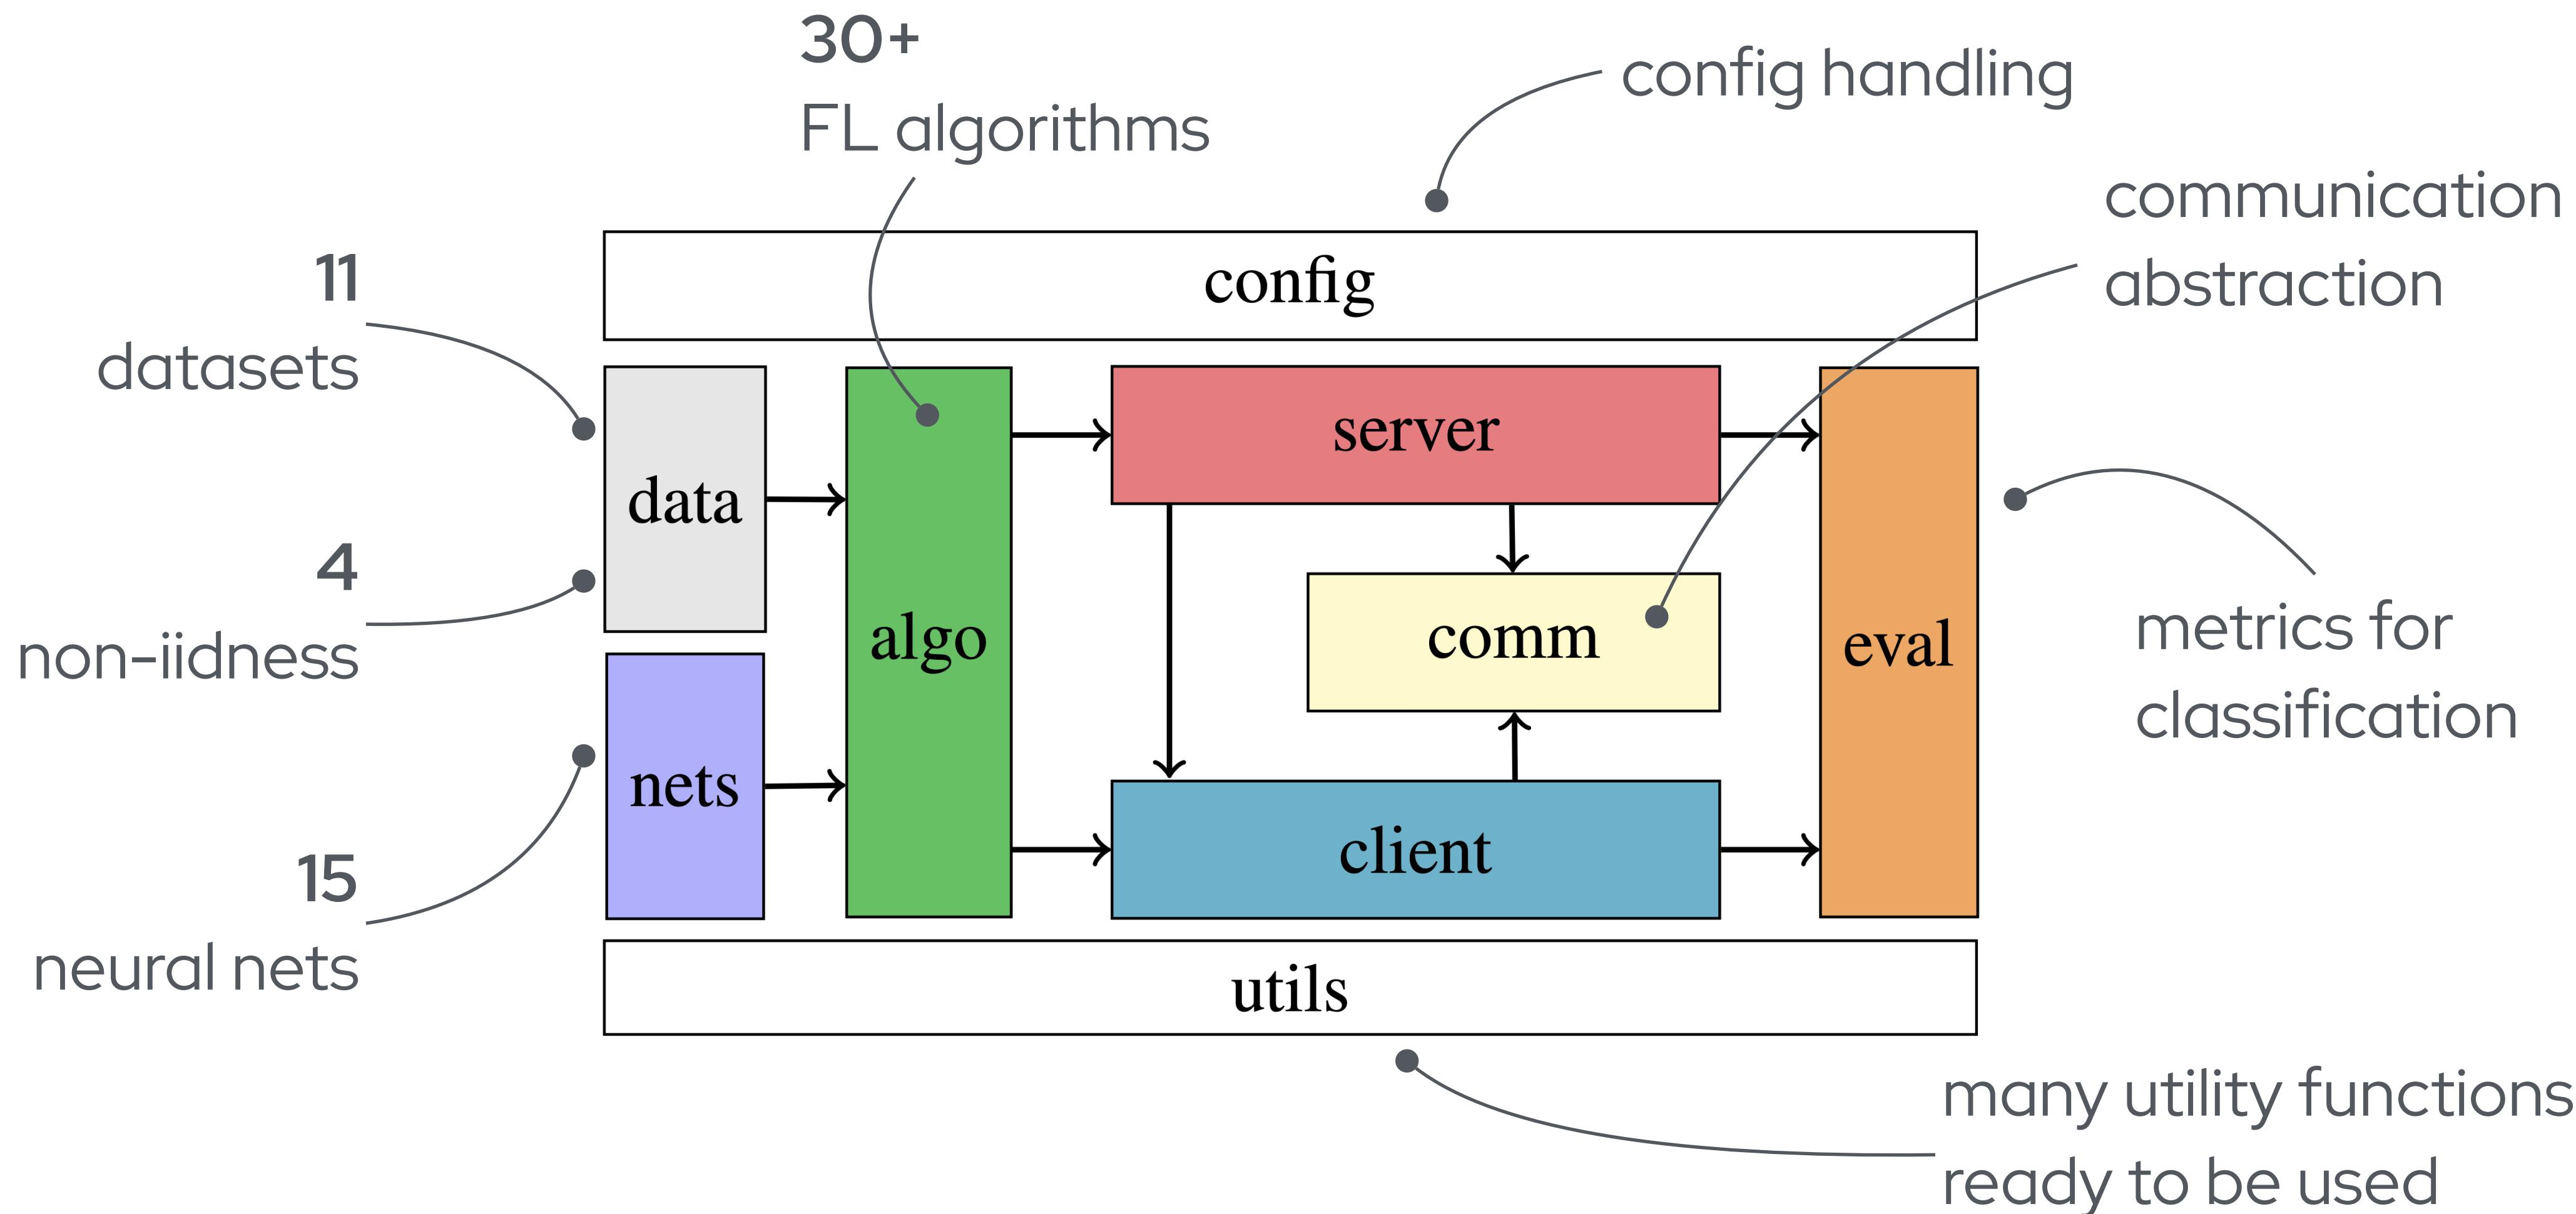
fluke python API

The python API offers all you need to implement and test your FL ideas

<https://github.com/makgyver/fluke>



SCAN ME



fluke server

Code readability is a key feature of fluke

Algorithm 1 FederatedAveraging. The K clients are indexed by k ; B is the local minibatch size, E is the number of local epochs, and η is the learning rate.

Server executes:

```
initialize  $w_0$ 
for each round  $t = 1, 2, \dots$  do
     $m \leftarrow \max(C \cdot K, 1)$ 
     $S_t \leftarrow (\text{random set of } m \text{ clients})$ 
    for each client  $k \in S_t$  in parallel do
         $w_{t+1}^k \leftarrow \text{ClientUpdate}(k, w_t)$ 
     $m_t \leftarrow \sum_{k \in S_t} n_k$ 
     $w_{t+1} \leftarrow \sum_{k \in S_t} \frac{n_k}{m_t} w_{t+1}^k$ 
```

```
1 def fit(self, n_rounds: int, eligible_perc: float) -> None:
2
3     for round in range(n_rounds):
4         eligible = self.get_eligible_clients(eligible_perc)
5         self.broadcast_model(eligible)
6
7         for c, client in enumerate(eligible):
8             client.local_update(round + 1)
9
10        self.aggregate(eligible)
```

fluke client

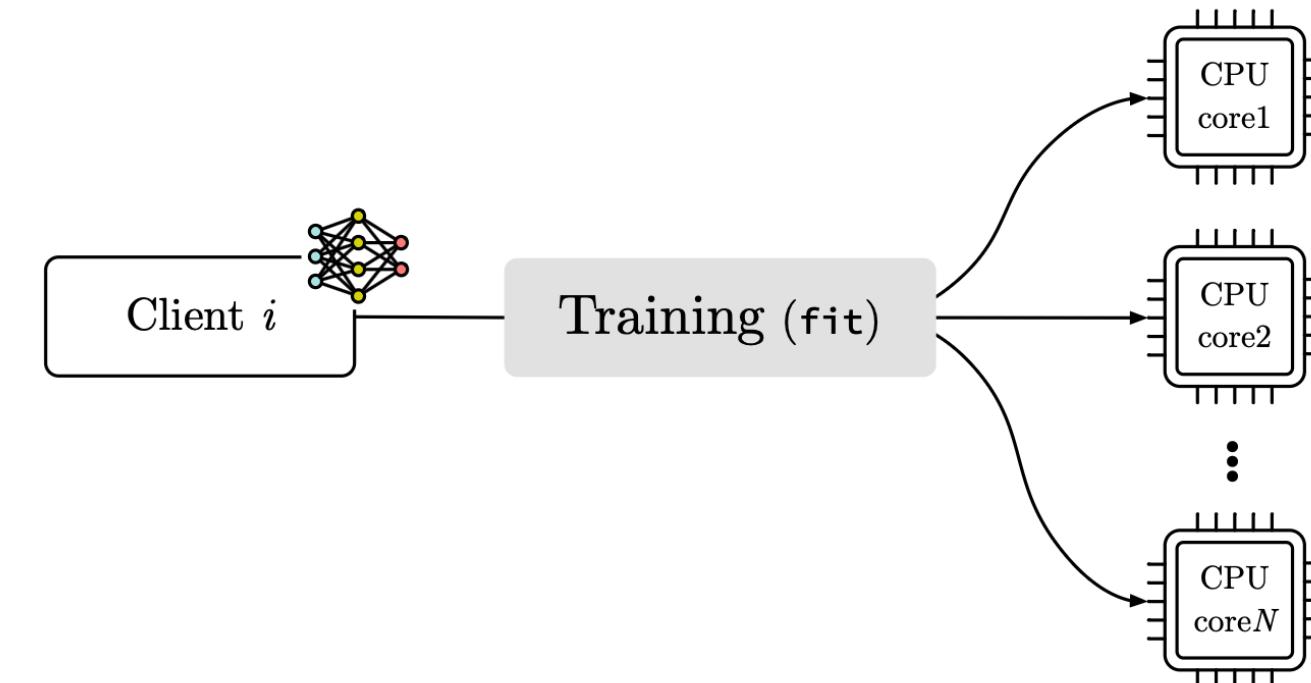
Code readability is a key feature of fluke

```
ClientUpdate( $k, w$ ): // Run on client  $k$ 
 $\mathcal{B} \leftarrow$  (split  $\mathcal{P}_k$  into batches of size  $B$ )
for each local epoch  $i$  from 1 to  $E$  do
    for batch  $b \in \mathcal{B}$  do
         $w \leftarrow w - \eta \nabla \ell(w; b)$ 
    return  $w$  to server
```

```
Client
1 def local_update(self, current_round: int):
2     self.receive_model()
3     self.fit()
4     self.send_model()
5
6 def fit(self, override_local_epochs: int = 0):
7     self.model.train()
8
9     if self.optimizer is None:
10         self.optimizer, self.scheduler = self.optimizer_cfg(self.model)
11
12     for _ in range(epochs):
13         for _, (X, y) in enumerate(self.train_set):
14             X, y = X.to(self.device), y.to(self.device)
15             self.optimizer.zero_grad()
16             y_hat = self.model(X)
17             loss = self.hyper_params.loss_fn(y_hat, y)
18             loss.backward()
19             self.optimizer.step()
20             self.scheduler.step()
21
22 def receive_model(self):
23     msg = self.channel.receive(self, self.server, msg_type="model")
24     self.model.load_state_dict(msg.payload.state_dict())
25
26 def send_model(self):
27     self.channel.send(Message(self.model, "model", self), self.server)
```

Parallelization in fluke

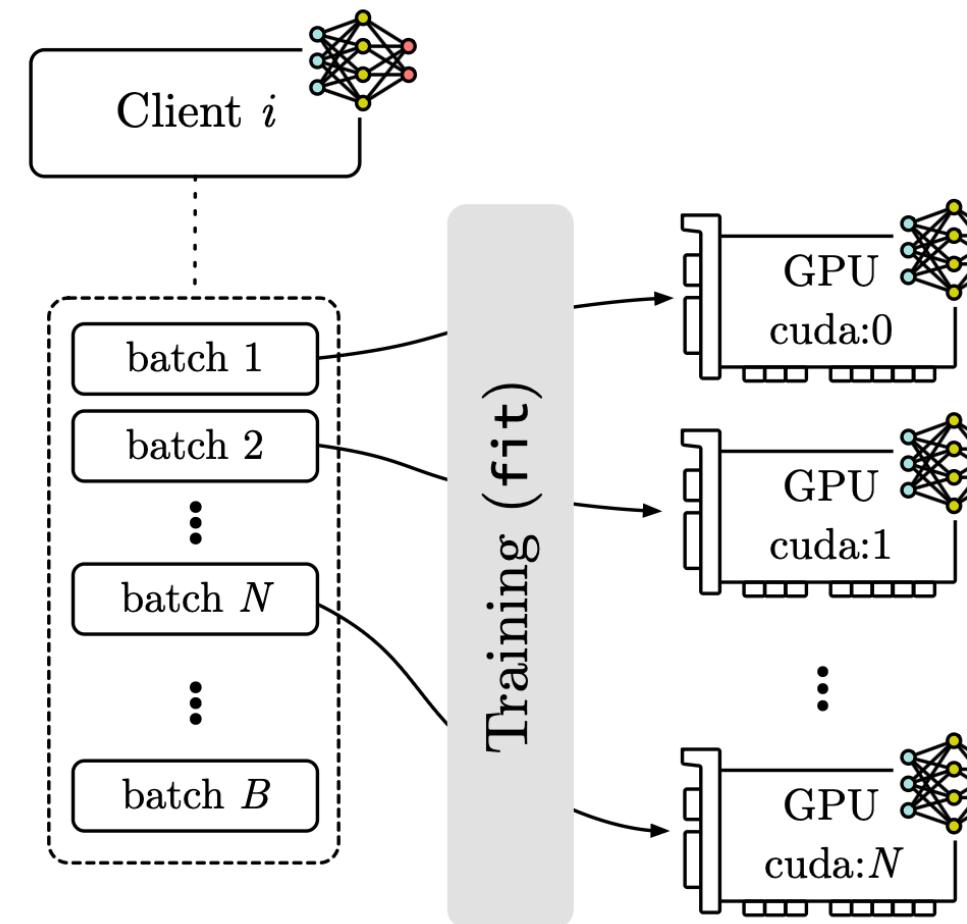
fluke *can easily handle multi-cpus/gpus*



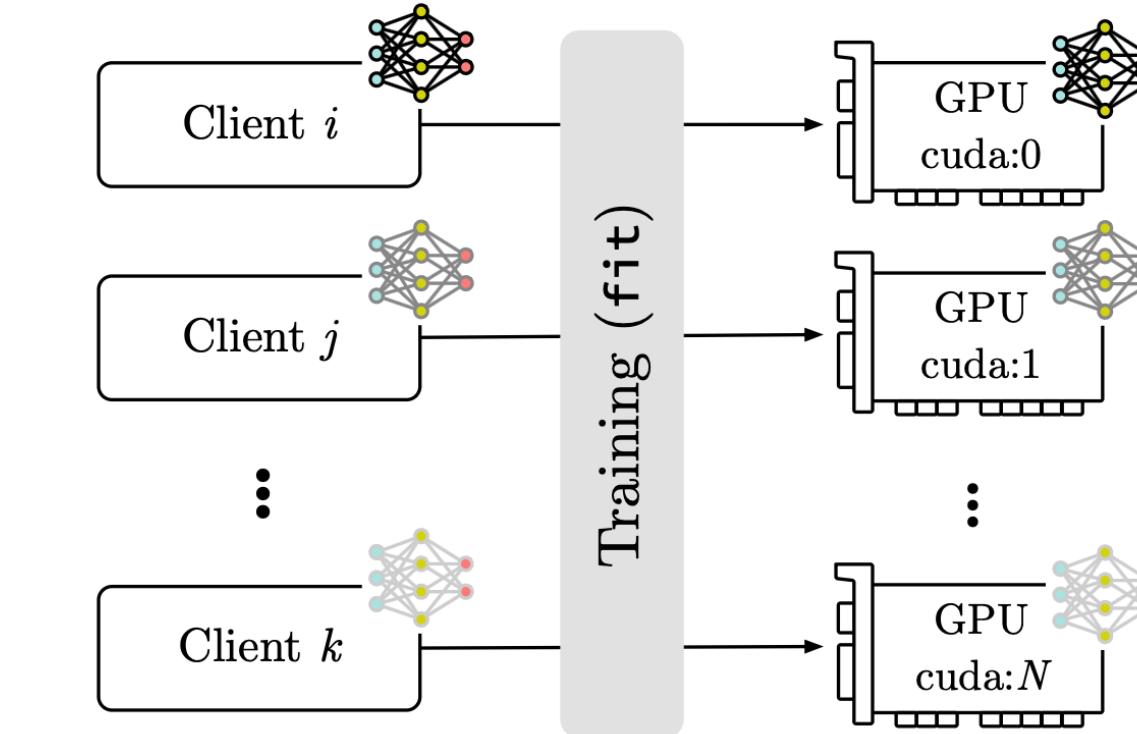
(a) Single client training on multiple CPU cores.



(b) Single client training on a single GPU.



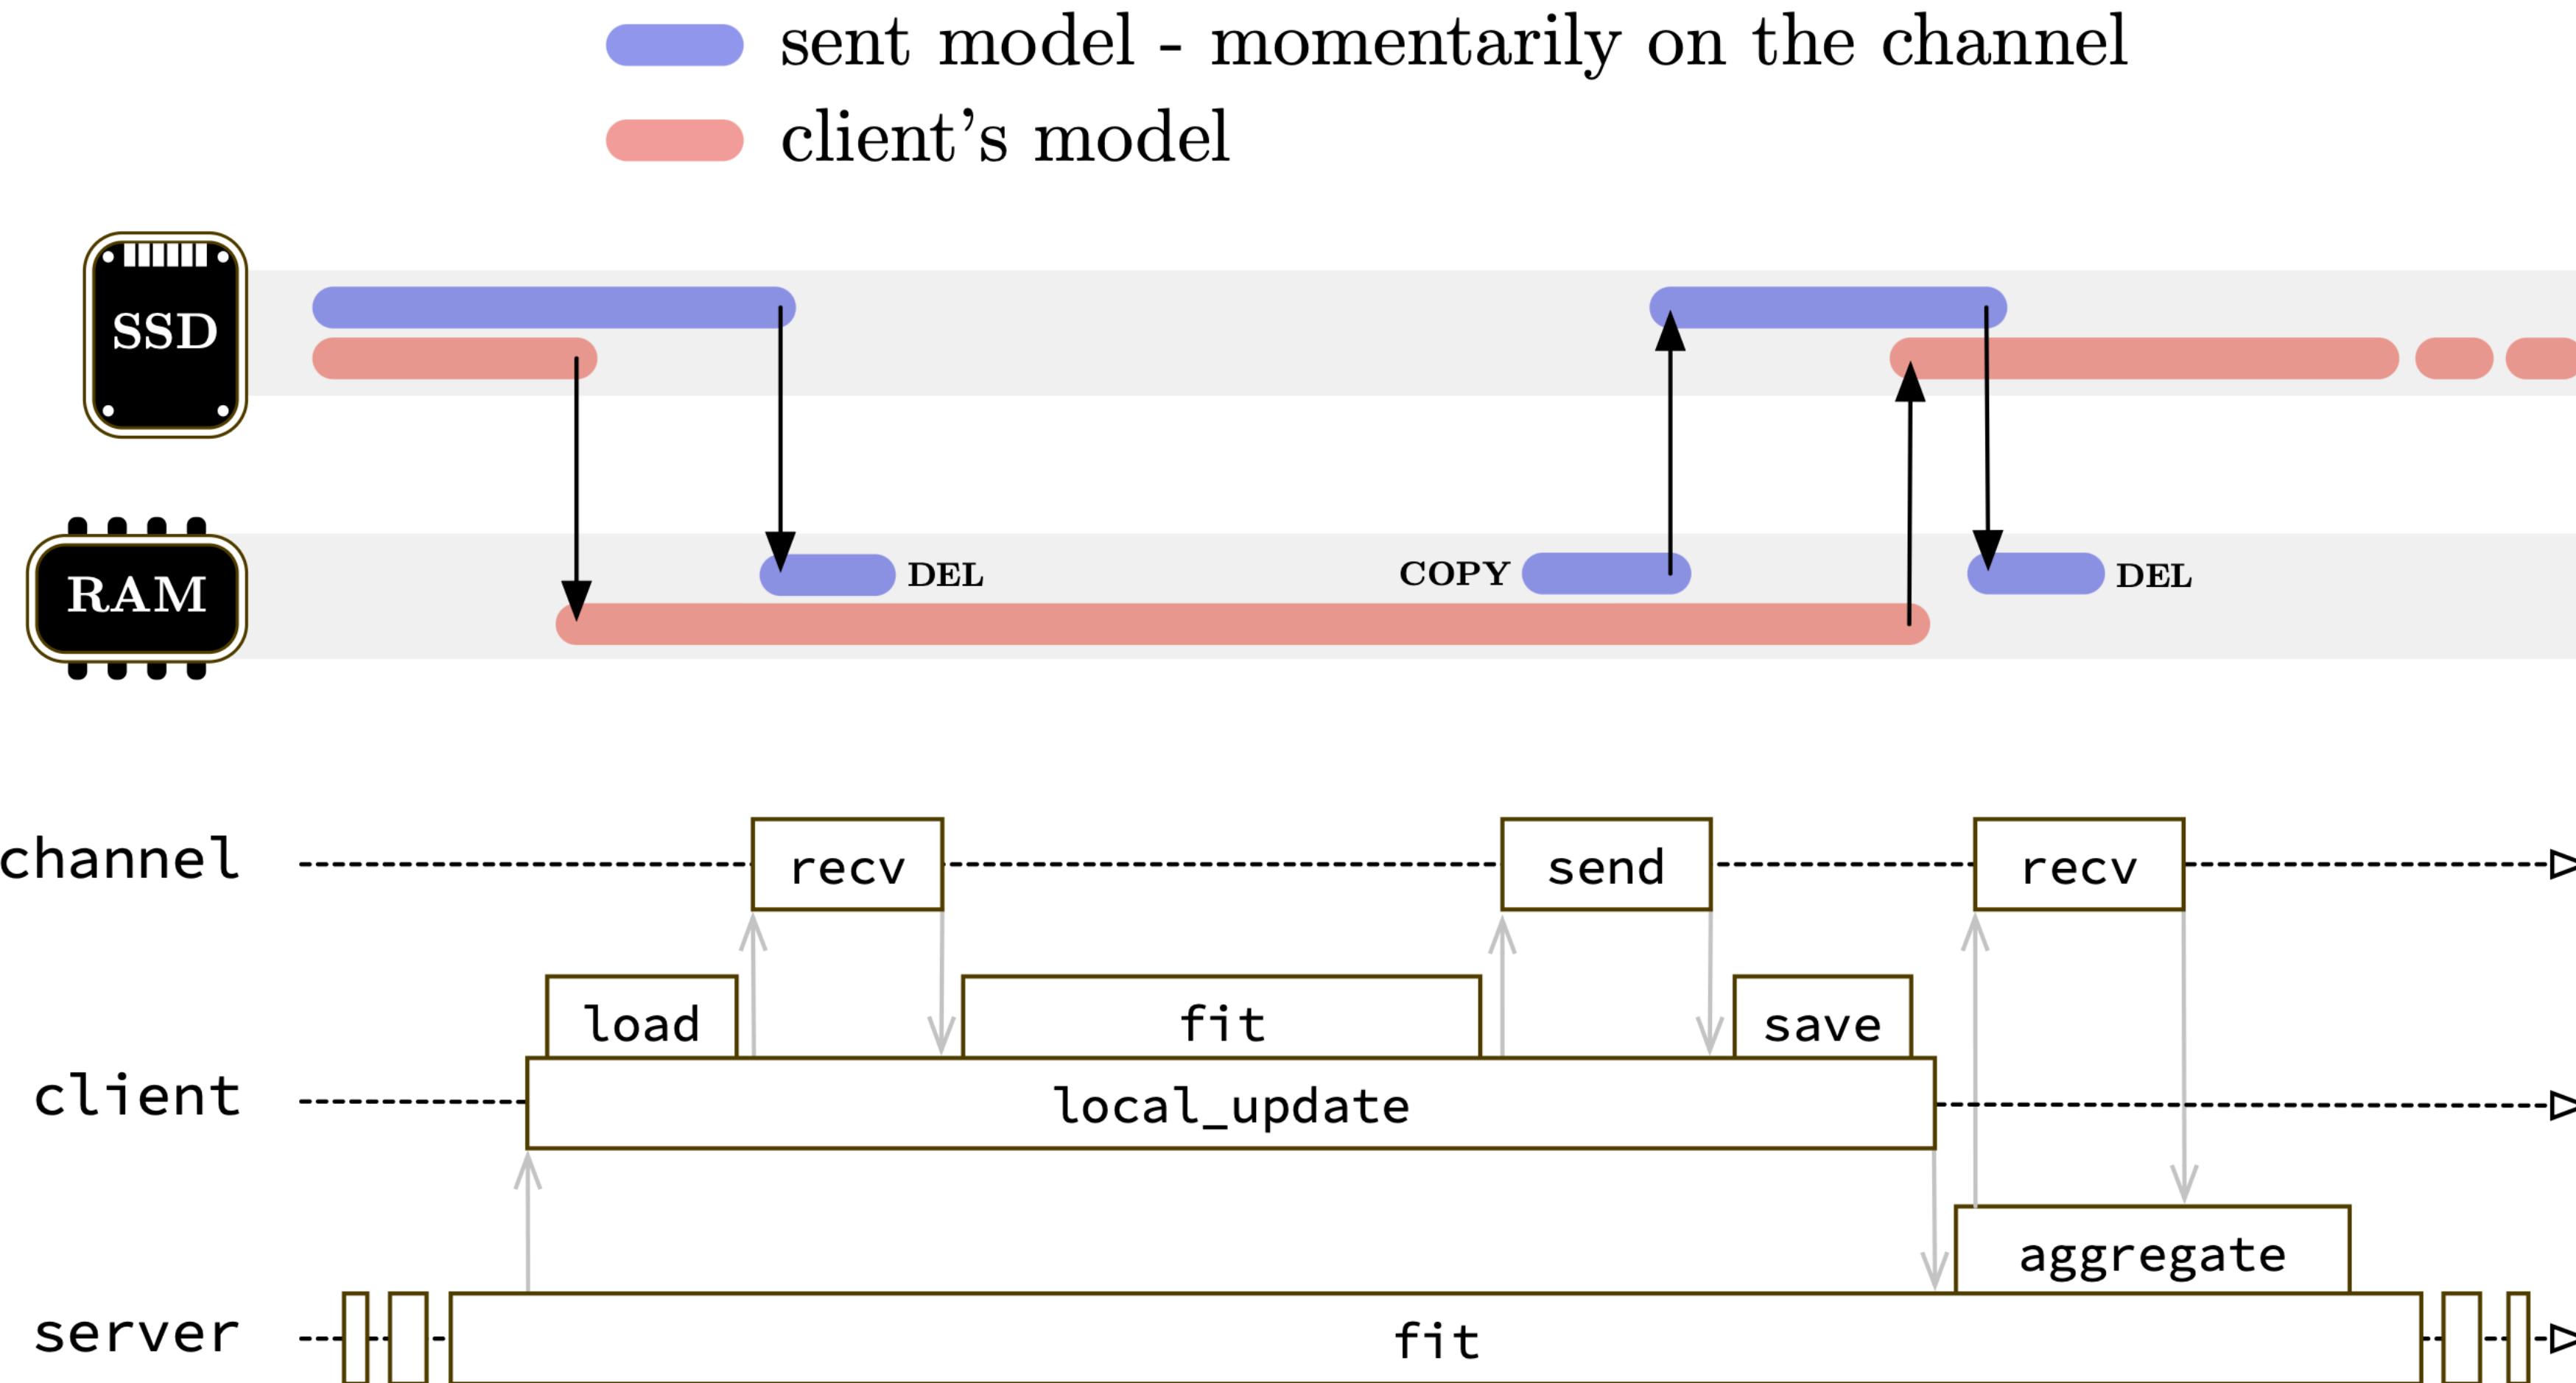
(c) Single client training on many GPUs.



(d) N clients training on N GPUs (one client per each GPU).

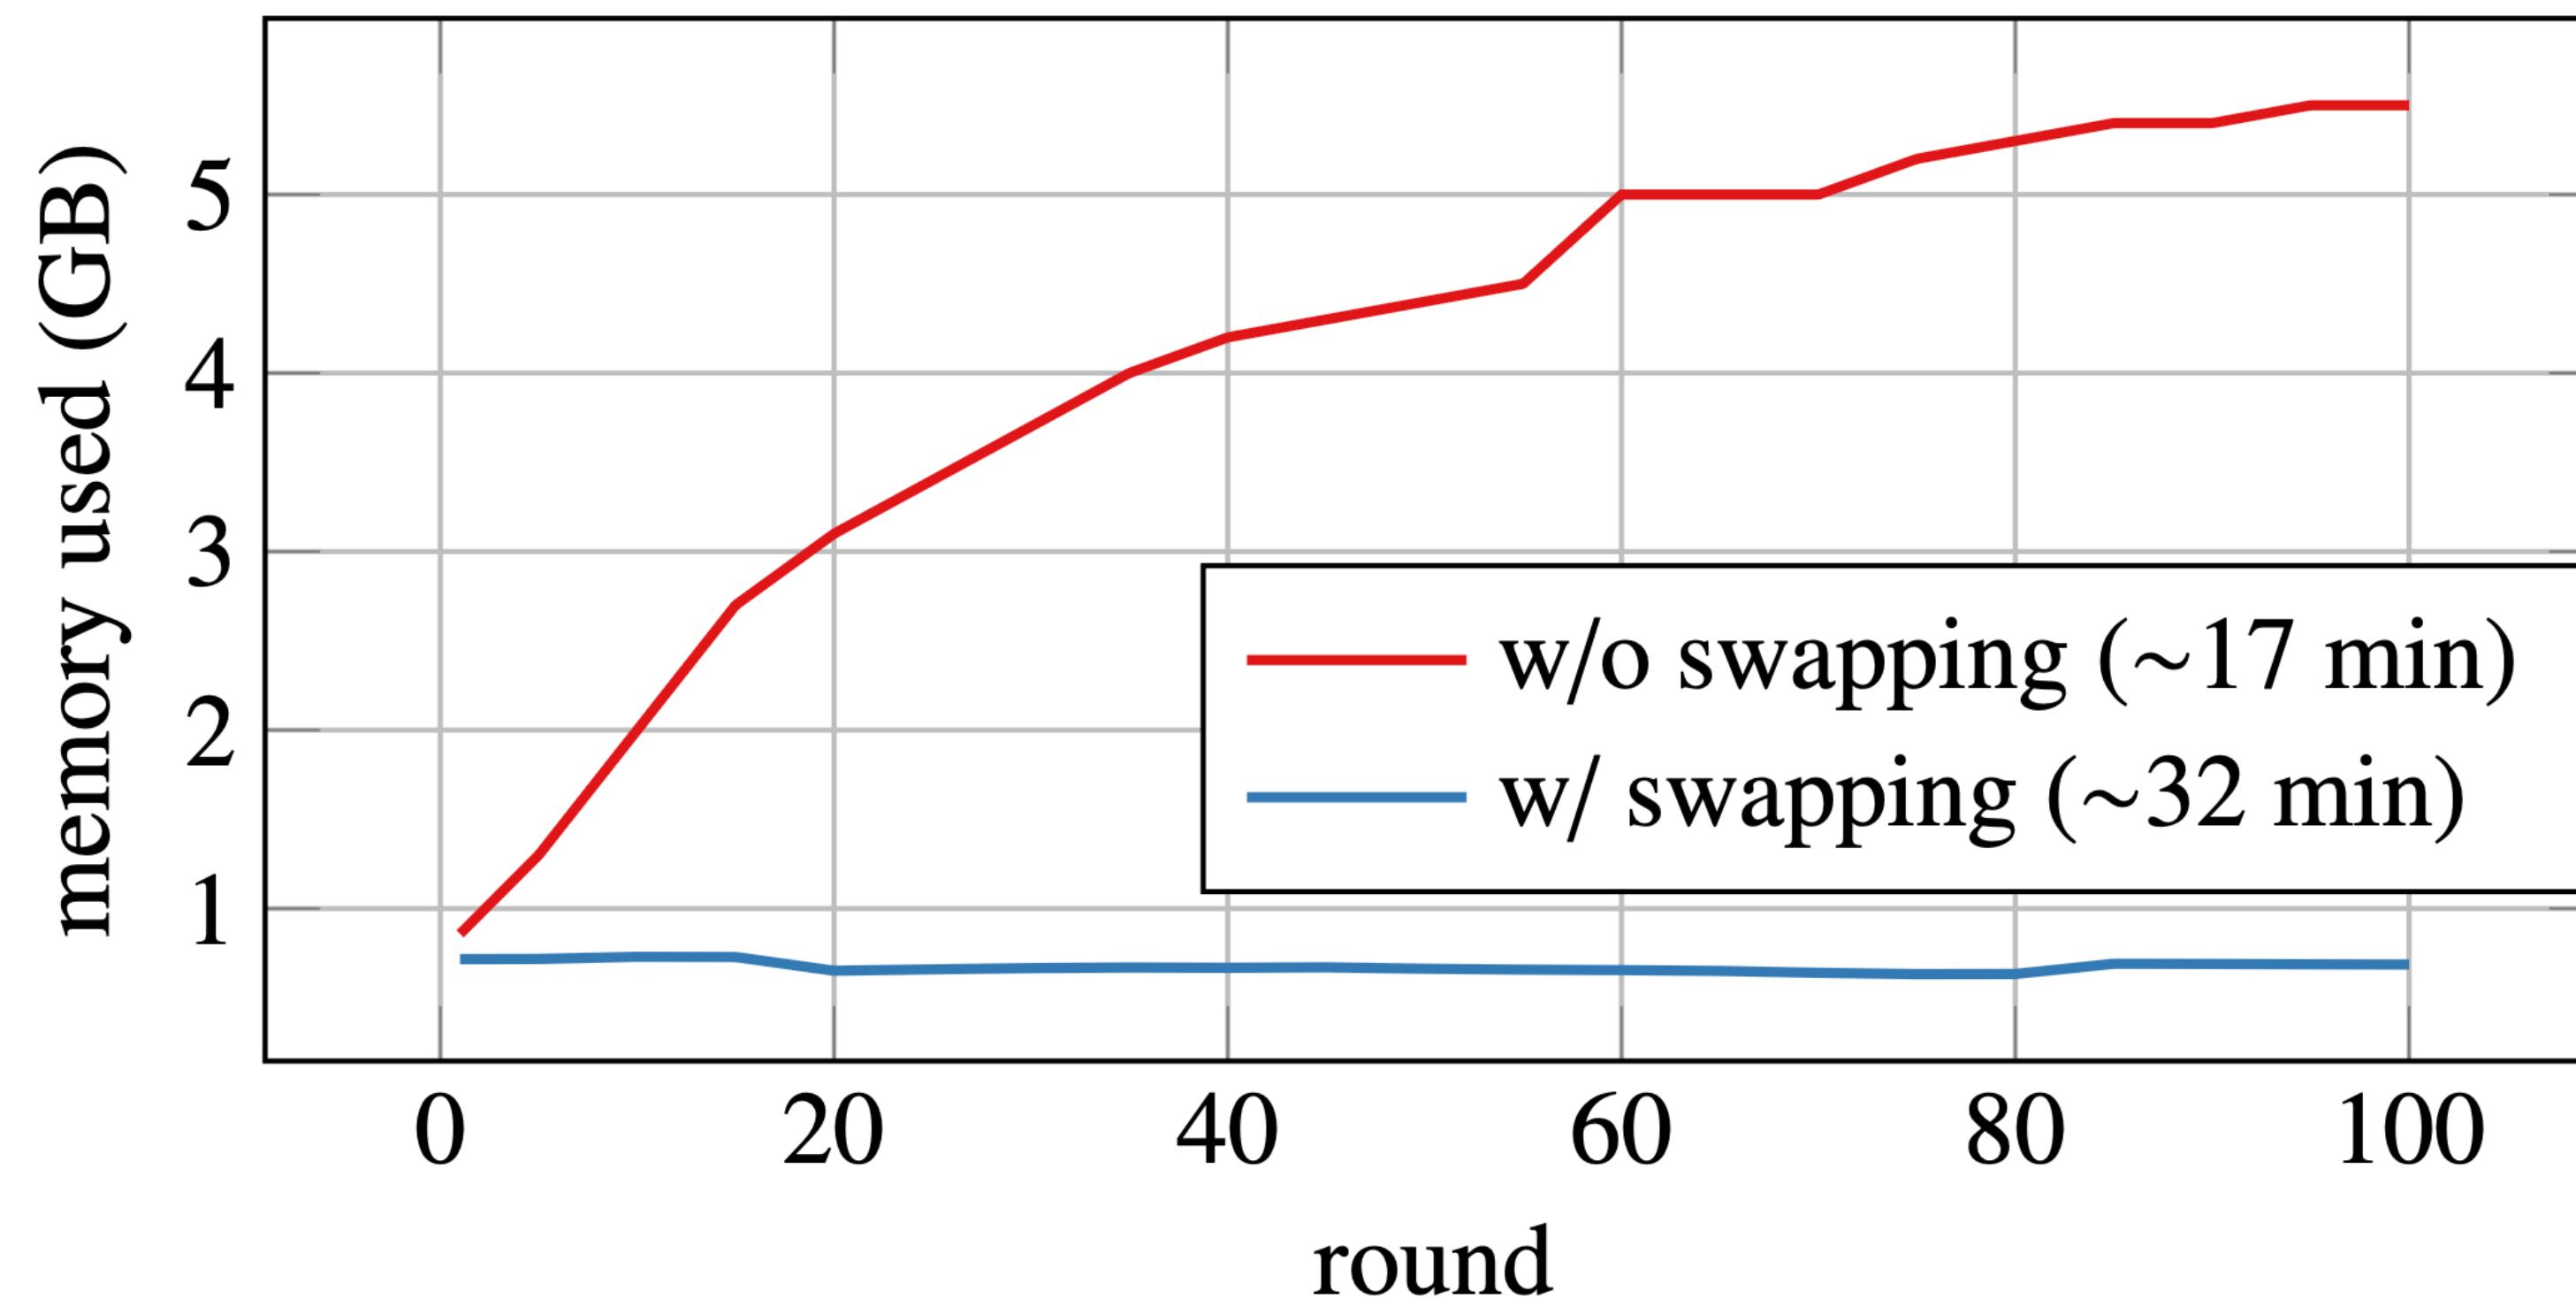
Memory management in fluke

fluke can optimize the memory usage



Memory management in fluke

The swapping mechanism guarantee (almost) constant memory consumption



Memory consumption in fluke with and without swapping on
MNIST with 1000 clients and a CNN model architecture.

Give it a try, you won't regret it!

fluke 0.7.9 (current release) - fluke 0.7.10 available soon!

30+

FL algorithms

11

datasets

15

neural networks



propose your FL algorithm
to be included in fluke!



<https://github.com/makgyver/fluke>

<https://makgyver.github.io/fluke/>

Hands-on session



<https://github.com/CasellaJr/Fluke-Tutorial-ECAI25>